
A Unified Convergence Theory for Non-Stationary Reinforcement Learning

Anonymous Author(s)

Affiliation

Address

email

Abstract

Non-stationary reinforcement learning lacks a unified convergence theory: variation-budget regret bounds, sliding-window methods, tempo-and-forgetting analyses, two-term decompositions, and causal-RL frameworks each address fragments without composing. We present a unified theory through the composition of four structural elements: (i) a two-gap diagnostic separating goal-feasibility from policy-quality; (ii) a point-mass reverse-KL/TV identity under deterministic optimum, $\text{TV}(\delta_{a^*}, Q) = 1 - e^{-D_{\text{KL}}}$, an *exact* algebraic identity that lies strictly below the Pinsker and Bretagnolle–Huber envelopes at this deterministic corner (collapsing the BH inequality’s envelope to a single value rather than introducing a new bounding technique); (iii) a multi-factor strategic tempo $\mathcal{T}_\Sigma = \min_{(i,j)} \nu_{ij} \cdot \iota_{ij} \cdot (1 - \lambda_{ij})$ with a structural forgetting prerequisite $\overline{\mathcal{T}}_\Sigma > \rho_\Sigma / R_\Sigma$ as a survival inequality; (iv) closed-loop interventional access making regret bounds learnable from on-policy data via the Pearl causal hierarchy. Composing these gives a cumulative dynamic regret theorem with rate $O(V_{\max} \sqrt{(B_T + 1)T})$ where B_T counts optimum-change events. We further establish a structural failure class: every gain-decay update (effective per-element gain $\alpha_{ij}^{(t)} \rightarrow 0$, including count-accumulating Bayesian schemes and Robbins–Monro gradient methods) eventually violates the forgetting prerequisite for any positive disturbance rate, while finite-gain mechanisms admit bidirectional thresholds (constant-step stochastic approximation, sliding windows, bounded memory, block restart). The point-mass identity is coordinate-optimal among bounds depending only on total variation, chain-rule compositional to multi-step trajectory regret as behavior-cloning loss against the optimal trajectory, and load-bearing for the cumulative theorem. The deterministic-optimum scope extends perturbatively to ϵ -stochastic and softmax-regularized policies with $O(\epsilon \log(1/\epsilon))$ and $O(\tau^{-1} \exp(-\Delta_{\min}/\tau))$ corrections respectively (exponentially small with polynomial prefactor). Theory-only; Lemma 5.2 makes the ProST reduction rigorous at the sector-parameter level, while their published experiments motivate the strategic-tempo component.

1 Introduction

Non-stationary reinforcement learning has matured along four largely separate research tracks. Each has produced rigorous results; none has been composed with the others into a single convergence theory.

The four tracks are visible in the recent literature as parallel lineages. **Variation-budget dynamic regret under drift**^[1–4] gives sublinear dynamic regret under bounded total variation in reward and transition dynamics. **Two-term regret decompositions**^[5–7] split the dynamic regret along an *exploration-vs-adaptation* axis, isolating the cost of confidence-set construction from the cost of

tracking a moving optimum. **Tempo and forgetting analyses**^[8–12] make policy-update timing or evidence-discount rate an explicit convergence variable. **Causal and interventional access**^[13–18] uses causal-graph structure to sharpen sample complexity, but operates in stationary settings only.

The information-theoretic regret literature^[19–24] traverses these tracks at a different layer, using entropy, mutual information, information ratio, Pinsker, or Hellinger to bound regret. Notably absent from that literature, as we document below (Appendix F), is an *exact* reverse-KL / total-variation identity at the deterministic-optimum corner: a sharp action-distribution regret coordinate that strictly improves both Pinsker and the Bretagnolle–Huber inequality^[25] at this point. The Bretagnolle–Huber inequality is itself absent from the retrieved RL/non-stationary corpus *as an upper-bound regret coordinate* (its use in lower-bound proofs via Le Cam’s method^[26,27] is well established and not contested here); the *exact identity at its point-mass corner* — which we establish — is *a fortiori* absent.

These four tracks share an evident common ancestor — the dynamic-regret analysis of online MDPs — but no published framework composes them. In particular, no framework combines (a) a regret decomposition along the *goal-feasibility-vs-policy-quality* axis (distinct from the exploration-vs-adaptation decompositions); (b) an exact point-mass reverse-KL/TV regret identity under deterministic optimum, strictly tighter than the Bretagnolle–Huber inequality at this corner; (c) a structural survival inequality threading rate-of-policy-revision against environment-side disturbance; and (d) a closed-loop causal-access argument that makes the regret bound *learnable* on-policy rather than merely provable analytically. The fundamental question remains open:

Can we design a unified convergence theory for non-stationary reinforcement learning that handles non-stationarity, has explicit metric structure on policy space, and is learnable from on-policy data?

In this paper we present the first composition that delivers all three properties jointly. Composing four structural elements — a two-gap diagnostic, a point-mass reverse-KL/TV identity, a multi-factor strategic tempo with forgetting prerequisite, and closed-loop interventional access — yields a cumulative dynamic regret theorem with rate $\text{DynReg}(T) \leq \tilde{O}(N_h \sqrt{(B_T + 1)T})$ in the piecewise-stationary regime, where B_T counts optimum-change events along the realized trajectory and N_h is the planning horizon. The per-round coordinate of this rate is sharper than both Pinsker and Bretagnolle–Huber at the deterministic- π^* corner. The technical anchor is an exact algebraic identity:

$\boxed{\text{TV}(\delta_{a^*}, Q) = 1 - e^{-D_{\text{KL}}(\delta_{a^*} \| Q)}}$, which holds for any policy Q with $Q(a^*) > 0$ and lies *strictly below* both the Pinsker envelope $\sqrt{D_{\text{KL}}/2}$ and the Bretagnolle–Huber envelope $\sqrt{1 - e^{-D_{\text{KL}}}}$ at this corner.

Roadmap Section 2 situates the four-track landscape. Section 3 sets minimal preliminaries. Section 4 presents the four-component composition theorem with narrative unpacking and a defense of why each component is necessary. Section 5 sketches the proof via four key lemmas chained through a block decomposition. Section 6 gives concluding observations and practitioner takeaways. Full proofs and auxiliary mathematical material are in the appendices.

1.1 Contribution

The contribution has the shape — recognized by the NeurIPS Theory Track guidelines as a valid form of originality — of “*a novel combination of existing techniques [where] the reasoning behind this combination is well-articulated.*” Three specific moves carry the paper.

(i) Point-mass reverse-KL/TV regret identity (Section 4, full development Section 5.1). Under deterministic optimum policy $\pi^* = \delta_{a^*}$, an elementary calculation gives the *exact identity* $D_{\text{KL}}(\pi^* \| Q) = -\log Q(a^*) = -\log(1 - \text{TV}(\pi^*, Q))$, equivalently $\text{TV}(\pi^*, Q) = 1 - e^{-D_{\text{KL}}(\pi^* \| Q)}$. Substituting this into the classical Bretagnolle–Huber inequality $\text{TV}(P, Q) \leq \sqrt{1 - e^{-D_{\text{KL}}(P \| Q)}}$ yields $1 - e^{-D_{\text{KL}}} \leq \sqrt{1 - e^{-D_{\text{KL}}}}$, which holds *strictly* on $D_{\text{KL}} > 0$ since $x < \sqrt{x}$ on $(0, 1)$. At the deterministic- π^* corner the identity supplies the *exact* TV value strictly below the BH envelope, positioning BH as the lineage the identity sits within rather than as the source it instantiates. Composed with the textbook total-variation regret bound (full development Section 5.1), the identity yields a tight two-sided characterization $\Delta_{\min}(1 - e^{-D_{\text{KL}}(\pi^* \| Q)}) \leq R(Q) \leq V_{\max}(1 - e^{-D_{\text{KL}}(\pi^* \| Q)})$, *strictly below* the Pinsker form $V_{\max} \sqrt{D_{\text{KL}}/2}$ — at $D_{\text{KL}} = 0.01$ the identity bound is $7\times$ tighter; at

88 $D_{\text{KL}} > 2$ Pinsker becomes vacuous against the trivial V_{max} envelope while the identity continues to
 89 saturate at V_{max} . The deterministic- π^* scope is not a hard wall: the identity extends *perturbatively* to
 90 ϵ -greedy and softmax-regularized optima with quantified $O(\epsilon \log(1/\epsilon))$ and $O(\tau^{-1} \exp(-\Delta_{\text{min}}/\tau))$
 91 corrections respectively (full development in Section 5.1 and Appendix B.2), recovering exact equality
 92 at $\epsilon = 0 / \tau \rightarrow 0$.

93 **(ii) Strategic tempo with forgetting prerequisite (Section 4, full development Section 5.2).** A
 94 multi-factor strategic tempo $\mathcal{T}_{\Sigma}$ — composed per revisable policy element of an observation rate
 95 ν , an identifiability coefficient ι , and an update gain η — yields a structural survival inequality
 96 $\mathcal{T}_{\Sigma}^{\text{bn,ss}} := \min_{(i,j) \in E} \nu_{ij} \cdot \iota_{ij} \cdot (1 - \lambda_{ij}) > \rho_{\Sigma}/R_{\Sigma}$ in which all three factors are multiplicatively
 97 load-bearing per element, with the *bottleneck* (per-element minimum) as aggregator. The familiar
 98 single-factor form $(1 - \lambda) > \rho_{\Sigma}/R_{\Sigma}$ is recovered as a corollary under uniform-Regime-A Beta-
 99 Bernoulli normalization. A structural-class theorem identifies the *gain-decay* class $\mathcal{A}_{\text{decay}}$ (count-
 100 accumulating Bayesian updates, observation-aggregating schemes without restart, gradient-based
 101 Robbins–Monro decaying-step methods) as one in which every agent eventually violates the threshold
 102 for any positive disturbance rate, while non-accumulating mechanisms (constant-step SA, sliding-
 103 window, bounded-memory, block-restart) face *bidirectional* ceilings. This *lifts* the tempo result of^[8]
 104 from a single-factor hyperparameter-optimization claim to a multi-factor structural-threshold claim.
 105 The forgetting prerequisite is the strategic analog of the persistence condition in adaptive control^[28],
 106 with $(1 - \lambda)$ playing the role of adaptive tempo per element.

107 **(iii) Closed-loop interventional access (Section 4, full development Section 5.3).** With explicit
 108 grounding in Bareinboim’s causal-hierarchy theorem^[29], the agent in the feedback loop *generates*
 109 Pearl Level-2 (interventional) data by construction: the action a_t causally precedes o_{t+1} , so under
 110 positivity, sequential ignorability, and a known action mechanism, the loop’s data-generating process
 111 is interventional rather than associational. This makes the regret bound of (i) *learnable* from on-policy
 112 interaction in regimes with sufficient identifiability, not merely provable in principle. The argument is
 113 implicit in the action-perception-loop framing of active inference^[30,31] and the causal-RL line^[13–18];
 114 we make the Bareinboim-hierarchy connection explicit, partition usable strength of identification
 115 into three regimes (A: full; B: partial; C: observation-only), and flag honestly where the loop yields
 116 interventional *data* without yielding identified *do*-estimates.

117 **Connective tissue: the two-gap diagnostic (Section 4).** A separate decomposition of the regret
 118 signal — the *satisfaction gap* δ_{sat} (the goal is unmet under the best available one-step improvement)
 119 versus the *control regret* δ_{regret} (the gap between the best available improvement and the agent’s current
 120 policy) — provides the corrective-action routing that ties (i)–(iii) together. The 2×2 disambiguation
 121 routes four regimes (success, strategy problem, capability limit, both) to four distinct corrective
 122 actions, along an axis orthogonal to the exploration-vs-adaptation decompositions of^[5–7].

123 **Composition** The four together yield (Section 4) a non-stationarity-aware convergence theory with
 124 three properties no existing strand exhibits jointly: the bound *handles non-stationarity* (via tempo and
 125 forgetting); has *explicit metric structure on policy space* (via the point-mass reverse-KL identity); is
 126 *learnable from on-policy data* (via closed-loop interventional access). Composition with a variation-
 127 budget block argument under a restart-on-change base learner yields the cumulative dynamic regret
 128 bound stated above — sharper per-round coordinate than Pinsker/BH, with the cumulative rate sitting
 129 in the piecewise-stationary B_T family of^[1] (*not* directly comparable to the continuous-variation
 130 rate $\tilde{O}(SA V_T^{1/3} H T^{2/3})$ of^[3]; the two regimes coincide up to constants under abrupt-magnitude
 131 piecewise-stationarity). Each of the four components is itself a theorem (cited or proved); the paper’s
 132 contribution is the recognition that the four together close a story none of the strands closes individually,
 133 plus the cumulative-regret rate that follows from composing the per-round identity with a base-learner
 134 stationary rate.

135 1.2 Scope and limitations

136 **What “convergence” means here** We do not claim pointwise policy convergence — under genuine
 137 non-stationarity the optimum is itself moving. By *convergence* we mean: per-round regret-coordinate
 138 convergence within stationary blocks; cumulative dynamic regret with sublinear / vanishing-average
 139 rate when $B_T = o(T)$ (Cesàro tracking $\frac{1}{T} \sum_t (V_t^* - V(\pi_t)) \rightarrow 0$); and ultimate boundedness of
 140 the strategic mismatch $\|\delta_{\Sigma}\|$ under the tempo threshold. Pointwise $V(\pi_t) \rightarrow V^*$ is structurally
 141 unavailable under genuine non-stationarity; the Cesàro tracking statement is the right replacement.

Theory only No experiments. The point-mass identity is mathematically airtight (a two-line direct calculation that specializes the BH family at its deterministic- π^* corner). The empirical grounding for the strategic-tempo claim is provided indirectly through reduction to^[8] ProST as a special case (Section 5); at the steady-state sector-parameter level the reduction is rigorous.

Canonical scope The composition theorem holds under deterministic optimum policy, bounded value range, isolated optimum (so $\Delta_{\min} > 0$), and a singular causal trajectory in the sense made precise in Section 3. The deterministic- π^* scope extends *perturbatively* to ϵ -stochastic and softmax-regularized optima (Section 5.1); tied-optimum is a degraded one-sided form (Appendix C.5). Each axis of the canonical scope degrades cleanly outside it (we discuss the failure modes in Section 6).

Comparator regime The cumulative rate is in the piecewise-stationary B_T family, not the continuous-variation V_T family of^[3]. Generalizing the per-round-identity route to continuous variation is open (Section 6).

2 Related Work

The four tracks named in Section 1 organize the literature.

Variation-budget dynamic regret ^[1–4,32] establish dynamic regret bounds under a *variation budget* V_T on cumulative reward and transition variation.^[3] achieves the near-optimal continuous-variation rate $\tilde{O}(SA V_T^{1/3} H T^{2/3})$ via RestartQ-UCB;^[1] treats the piecewise-stationary case with restart-on-change. Our cumulative rate $\tilde{O}(N_h \sqrt{(B_T + 1)T})$ sits in the piecewise-stationary family — same block-shape as Cheung-style restart-on-change — recovered as a corollary of composition with the per-round identity. The novelty is the *threshold form* of the strategic-tempo inequality (environment regimes where no $\lambda \in (0, 1)$ stabilizes the sector-Lyapunov mismatch), which dynamic-regret-optimization analyses do not surface.

Two-term regret decompositions along the exploration-vs-adaptation axis ^[5–7] decompose the dynamic regret into a *confidence-set construction* term and an *adaptation-under-non-stationarity* term.^[33] treats a structurally adjacent question — constraint-region feasibility — using a two-term decomposition along that axis. Same shape, different axis: ours is goal-feasibility-vs-policy-quality (the satisfaction gap distinguishes “the goal is unattainable from M_t ” from “the current policy is suboptimal”); theirs is uncertainty-source. Neither subsumes the other: a Long-Fei Li-style exploration term can be present in any of our 2×2 cells, and an adaptation term likewise. The two axes are orthogonal.

Tempo and forgetting ^[8] (ProST) defines an agent tempo as the schedule of policy update times $\{t_1, \dots, t_K\}$ and computes the schedule that minimizes a dynamic regret upper bound.^[9] (Pausing Policy Learning) refines this with non-zero hold durations.^[10–12] establish forgetting / sliding-window analyses for non-stationary linear MDPs and bandits. We *lift* the ProST move along two axes simultaneously: single-factor $\rightarrow$ multi-factor (per-element $\nu \cdot \iota \cdot (1 - \lambda)$ rather than scalar update frequency); and hyperparameter-optimization $\rightarrow$ structural-survival inequality (the threshold below which no schedule persists, irrespective of which schedule is chosen). ProST recovers as the Beta-Bernoulli special case $|E| = 1, \nu = \iota = 1$.

Causal and interventional access ^[13–18,34] use causal-graph structure to sharpen sample complexity in RL. Their setting is *stationary*. We compose the closed-loop interventional-access argument with non-stationarity. The connection between Pearl’s causal hierarchy^[35] and the data-generating process of an agent in a feedback loop^[29] is the technical bridge.

Information-theoretic regret (cross-cutting) ^[19–24,36] traverse the above tracks at a different layer, using entropy / mutual information / information ratio / Pinsker / Hellinger as the divergence backbone. The Russo–Van Roy information ratio uses Pinsker as its Cauchy–Schwarz follow-up; the IDS line and its descendants use mutual-information bounds. The Bretagnolle–Huber inequality^[25] — well-known in lower-bound proofs via Le Cam’s method^[26,27] — does not appear *as an upper-bound regret coordinate* in this corpus, and its exact point-mass corner (our Lemma 5.1) is *a fortiori* absent. We document the retrieval scope and methodology in Appendix F.

191 **Adjacent (satisficing / feasibility)** ^[37,38] use *satisficing* to mean “any policy above acceptance
 192 level β is acceptable.” Our δ_{sat} is structurally distinct: $V_O^{\min}$ is a property of the objective (set by
 193 the application domain), and $\delta_{\text{sat}} > 0$ signals goal-relative *infeasibility* (no policy in Π meets the
 194 threshold), not policy-relative satisficing.

195 **Contemporaneous (post-March 2026)** ^[38,39]; cited and distinguished without empirical-
 196 comparison demand per the contemporaneous-work cutoff convention.

197 3 Preliminaries

198 We adopt strict minimalism: only the notation and assumptions used directly in the main result and
 199 the mechanism statements appear here. Convention-hierarchy details (one-step vs. receding-horizon
 200 vs. Bellman continuation), the action-gap matching argument, tied-optimum extensions, and the
 201 deterministic-trajectory deployment modes are deferred to Appendix C.

202 **Episodic non-stationary MDP** A finite-horizon non-stationary MDP $(\mathcal{S}, \mathcal{A}, P_t, r_t, N_h)$ with state
 203 space $\mathcal{S}$, finite action space $\mathcal{A}$, time-indexed transition kernels $P_t(\cdot \mid s, a)$ and reward functions
 204 $r_t(s, a)$ allowed to vary across rounds t , and finite planning horizon N_h . Standard non-stationary-RL
 205 boundedness assumptions^[1]: per-step reward in $[0, 1]$.

206 **Variation regimes** A variation budget V_T measures continuous variation in transitions and rewards.
 207 We work in the *piecewise-stationary* specialization: time partitions into $B_T + 1$ stationary intervals
 208 separated by optimum-change events, with $B_T := |\{t : a_t^* \neq a_{t-1}^*\}|$. The two regimes are distinct in
 209 general; under abrupt-magnitude- δ piecewise-stationarity they coincide up to constants ($V_T \asymp \delta B_T$).
 210 Our results are stated for B_T ; the continuous- V_T extension is open (Section 6).

211 **Policy distribution** The agent’s policy at round t is a distribution $Q_t(\cdot \mid s)$ over actions; we write Q
 212 when state and round are understood. We adopt the canonical scope: the optimum policy at (M_t, s)
 213 is *deterministic*, $\pi^*(\cdot \mid M_t, s) = \delta_{a^*(s)}$ where $a^*(s) := \arg \max_a Q^\pi(s, a)$ is the optimal action
 214 under the agent’s current model M_t . Deterministic π^* is canonical for finite-MDP RL with isolated
 215 optima^[27]; the perturbative extension to ϵ -stochastic and softmax-regularized optima (Section 5.1,
 216 full derivation Appendix B.2) covers smooth deviations.

217 **Value object** Given an objective O , model M_t , policy π , and horizon N_h : $Q_O(M_t, a; \pi_{\text{cont}}, N_h) =$
 218 $\mathbb{E}[V_O(\tau) \mid M_t, \text{do}(a_t = a), a_{t+1} : \sim \pi_{\text{cont}}]$. The $\text{do}(\cdot)$ ^[35] flags intervention on a_t , not conditioning.
 219 We adopt the **one-step improvement** convention $\pi_{\text{cont}} = \pi_{\text{current}}$ as default: most conservative, compa-
 220 rable across rounds, fixed-point-free. Receding-horizon and Bellman alternatives are in Appendix C.1.

221 **Bounded value range** $V_{\max}(M_t) := \max_a Q_O(M_t, a) - \min_a Q_O(M_t, a)$, finite under bounded
 222 reward and finite horizon.

223 **Action gap (isolated optimum)** $\Delta(a) := Q_O(a^*) - Q_O(a) \in [0, V_{\max}]$, $\Delta_{\min} := \min_{a \neq a^*} \Delta(a) >$
 224 0 whenever the optimum is isolated.

225 **Strategic regret** For policy distribution Q : $R(Q) := Q_O(a^*) - \mathbb{E}_{a \sim Q}[Q_O(a)] = \sum_{a \neq a^*} Q(a) \cdot$
 226 $\Delta(a)$. Three classical forms — $\mathbb{E}_{\pi^*}[V] - \mathbb{E}_Q[V]$, $V(a^*) - \mathbb{E}_Q[V]$, $\mathbb{E}_{\pi^*}[V - V_Q]$ — coincide under
 227 deterministic π^* .

228 **Total variation and reverse-KL** For finite action measures P, Q on $\mathcal{A}$, $\text{TV}(P, Q) :=$
 229 $\frac{1}{2} \sum_a |P(a) - Q(a)|$, $D_{\text{KL}}(P \parallel Q) := \sum_a P(a) \log \frac{P(a)}{Q(a)}$ with the convention $0 \log 0 = 0$.
 230 Our bounds use the *reverse* KL direction (with $P = \pi^*$); the forward direction $D_{\text{KL}}(Q \parallel \pi^*)$ is $+\infty$
 231 whenever Q has off-optimum mass and is therefore vacuous as a regret coordinate (full direction-
 232 forcing argument in Appendix C.3).

233 **Singular trajectory** The agent operates on a single, non-forkable causal trajectory: a_t causally
 234 precedes o_{t+1} , and the model state M_t ’s sufficiency is defined relative to *this* trajectory rather than to

235 a model-state equivalence class. This is the operational ground for the do-operator above and for the
 236 closed-loop interventional-access argument (Section 5.3).

237 4 Main Result

238 We present the four components first as motivation, then state the formal composition theorem, then
 239 unpack what its conclusions mean.

240 4.1 The four components

241 Each component supplies a *distinct epistemic property* that the others cannot deliver. The composition
 242 is necessary for the joint properties (rate + persistence + learnability + corrective-action routing), not
 243 for the rate alone. The component-by-component failure-mode analysis is in Section 4.4 below.

244 **Component 1: Two-gap diagnostic — corrective-action routing** Let Π denote the policy class
 245 available to the agent and let $V_O^{\min}$ denote the threshold trajectory value above which the objective is
 246 “met.” Define *objective attainability* $A_O(M_t; \Pi, N_h) := \sup_{\pi \in \Pi} V_O(M_t, \pi; N_h)$, the *satisfaction*
 247 *gap* $\delta_{\text{sat}} := V_O^{\min} - A_O$ (goal-feasibility diagnostic; positive means *unmet* under the best available
 248 policy), and the *control regret* $\delta_{\text{regret}} := A_O - V_O(M_t, \pi_{\text{current}}; N_h) \geq 0$ (policy-quality diagnostic).
 249 The 2×2 cross of these two gaps routes corrective action: *Success* (no action), *Strategy problem*
 250 (revise the policy), *Capability limit* (revise the model, policy class, or horizon — in that order), *Both*
 251 (revise the policy first, then re-evaluate). Without this routing, an agent receiving a regret signal
 252 cannot distinguish “the goal is unattainable from M_t ” from “the current policy is suboptimal.” A
 253 single $\delta_{\text{objective}} = V_O^{\min} - V_O(M_t, \pi_{\text{current}}; N_h)$ collapses both regimes into a single number and loses
 254 the diagnostic. The two-gap split runs along an axis *orthogonal* to the exploration-vs-adaptation
 255 decompositions of [5–7]: theirs isolates uncertainty *sources*; ours isolates *what is wrong*.

256 **Component 2: Point-mass reverse-KL/TV identity — metric layer on policy space** Under
 257 deterministic $\pi^* = \delta_{a^*}$, the reverse-KL and total variation between π^* and any policy Q are related by
 258 the *exact identity* $D_{\text{KL}}(\pi^* \parallel Q) = -\log Q(a^*) = -\log(1 - \text{TV}(\pi^*, Q)) \iff \text{TV}(\pi^*, Q) =$
 259 $1 - e^{-D_{\text{KL}}(\pi^* \parallel Q)}$. This is a two-line direct calculation under the point-mass assumption (full statement
 260 and proof in Section 5.1). It is *not* the Bretagnolle–Huber inequality at equality: substituting it into
 261 the BH form $\text{TV} \leq \sqrt{1 - e^{-D_{\text{KL}}}}$ gives $1 - e^{-D_{\text{KL}}} \leq \sqrt{1 - e^{-D_{\text{KL}}}}$, *strictly* on $D_{\text{KL}} > 0$. The
 262 identity sits *strictly below* the BH envelope at this corner. Composed with the textbook TV-regret
 263 bound $R(Q) \leq V_{\max} \text{TV}(\pi^*, Q)$ and the matching action-gap lower bound, it yields the two-sided
 264 characterization $\Delta_{\min}(1 - e^{-D_{\text{KL}}(\pi^* \parallel Q)}) \leq R(Q) \leq V_{\max}(1 - e^{-D_{\text{KL}}(\pi^* \parallel Q)})$, *Lipschitz-*
 265 *equivalent* with constants $\Delta_{\min}/V_{\max}$ and 1. The identity coordinate is *coordinate-optimal among*
 266 *bounds depending only on TV*: any tighter bound on $R(Q)$ requires information beyond TV (the
 267 value-landscape spread). Strict improvement over Pinsker is uniform: at $D_{\text{KL}} = 0.01$ the identity
 268 bound is $7 \times$ tighter; for $D_{\text{KL}} > 2$ Pinsker is vacuous against the trivial $V_{\max}$ envelope while the identity
 269 saturates at $V_{\max}$. The deterministic- π^* scope is not a hard wall: the identity extends *perturbatively* to
 270 ϵ -greedy and softmax-regularized optima with corrections $O(\epsilon \log(1/\epsilon))$ and $O(\tau^{-1} \exp(-\Delta_{\min}/\tau))$
 271 respectively — exponentially small with polynomial prefactor in the softmax case (Section 5.1).

272 **Component 3: Multi-factor strategic tempo with forgetting prerequisite — non-stationarity**
 273 **persistence** Consider a structured policy with $|E|$ revisable elements (in our analysis, edges of
 274 an internal causal model the agent maintains over policy components; in Lee et al. 2023 [8]’s ProST,
 275 sub-policies indexed by a tempo schedule). For each element (i, j) , let ν_{ij} denote the per-element
 276 observation rate, $\iota_{ij} \in [0, 1]$ the identifiability coefficient (fraction of evidence that genuinely identifies
 277 the causal effect rather than reflecting confounded association), and $\eta_{\text{edge}, ij}$ the per-element update
 278 gain. Within a sector-Lyapunov reduction with diagonal correction (full Model (S) in Section 5.2),
 279 per-element exponential forgetting at rate λ_{ij} gives steady-state gain $\eta_{\text{edge}, ij}^{\text{ss}} \approx 1 - \lambda_{ij}$, and ultimate
 280 boundedness of the strategic mismatch $\|\delta_{\Sigma}\|$ within the modeled dynamics is *sufficient* under the
 281 bottleneck threshold $\mathcal{T}_{\Sigma}^{\text{bn,ss}} := \min_{(i,j) \in E} \nu_{ij} \cdot \iota_{ij} \cdot (1 - \lambda_{ij}) > \rho_{\Sigma}/R_{\Sigma}$. All three factors are
 282 multiplicatively load-bearing per element (zeroing any collapses the bottleneck; low identifiability
 283 cannot be compensated by faster forgetting), and the *bottleneck* (per-element minimum, not aggregate
 284 sum) is the right aggregator: a Lyapunov derivation under adversarial disturbance saturates at min,

not $\sum$, when the disturbance concentrates on the weakest element. The structural-class theorem identifies the *gain-decay* class $\mathcal{A}_{\text{decay}}$ — count-accumulating Bayesian updates without forgetting (e.g., Beta-Bernoulli with $\eta = 1/(n+1)$), bounded-memory schemes with growing memory, observation-aggregating schemes without restart, and gradient-based methods with vanishing step sizes (the Robbins–Monro regime, $\eta_t = c/t^\beta$ for $\beta \in (0, 1]$) — as one in which *every* agent eventually violates the threshold for any positive disturbance rate (formal statement in Appendix C.6). Non-accumulating mechanisms (constant-step SA, sliding-window, bounded-memory, block-restart) face *bidirectional* ceilings: the threshold is tight in both directions and the schedule has to be in the right regime, not arbitrarily fast or slow. The familiar single-factor form $(1 - \lambda) > \rho_\Sigma / R_\Sigma$ is recovered as a corollary under uniform-Regime-A Beta-Bernoulli normalization.

Component 4: Closed-loop interventional access — on-policy learnability Pearl’s causal hierarchy^[35] and the causal-hierarchy theorem of^[29] establish that Level-2 (interventional) queries $P(Y \mid \text{do}(X))$ are not in general identifiable from Level-1 (associational) data $P(Y \mid X)$. By temporal ordering and the singular-trajectory commitment of Section 3, a_t causally precedes o_{t+1} . Under three conditions on the agent’s action mechanism — (C1) *positivity* ($\pi_t(a \mid H_t) \geq p_{\min} > 0$ on the support of the information state H_t), (C2) *sequential ignorability* (the information state blocks unobserved confounding of action selection), and (C3) *known action mechanism* (the behavior policy is known) — conditioning on H_t blocks the action-selection confounding pathway, so the conditional distribution $P(o_{t+1} \mid a_t, H_t)$ coincides with the interventional distribution $P(o_{t+1} \mid \text{do}(a_t), H_t)$. The loop’s data-generating process is *interventional* by construction (formal statement in Section 5.3). This makes the regret bound of Component 2 *learnable on-policy* in regimes with sufficient identifiability: the agent’s identified optimum $a_{\text{ag},t}^*$ matches the true optimum $\tilde{a}_t^*$ with probability p_{id} , and the per-state, per-step bias from misidentification is bounded by $(1 - p_{\text{id}}(s)) \log(1/q_0)$ under two-point support $Q_t \geq q_0$ (Section 5.4). Three regimes^[29] partition usable strength: A (intervention-rich, $\iota \approx 1$), B (partial, $\iota \in (0, 1)$), C (observation-only, $\iota \approx 0$).

4.2 The composition theorem

Theorem 4.1 (Composition: joint guarantees with conditional cumulative reduction). *Let $(\mathcal{S}, \mathcal{A}, P_t, r_t, N_h)$ be a non-stationary MDP with bounded reward, finite action space, and for each round t a deterministic optimum policy $\pi_t^* = \delta_{a_t^*}$ with action gap $\Delta_{\min} > 0$. Suppose the agent operates on a singular causal trajectory in the sense of Section 3 and updates its policy with per-element exponential forgetting at rates $\{\lambda_{ij}\}$. Let $K_t(s) := D_{\text{KL}}(\delta_{a_t^*(s)} \parallel Q_t(\cdot \mid s))$ be the per-state reverse-KL coordinate (in the bandit case $N_h = 1$ the state argument is suppressed). If*

(A1) **Metric on policy space.** $Q_t(a_t^* \mid s) > 0$ at every visited state and round.

(A2) **Multi-factor forgetting prerequisite.** Within Model (S) of Section 5.2, the bottleneck strategic tempo $\tau_\Sigma^{\text{bn,ss}} := \min_{(i,j) \in E} \nu_{ij} \cdot \iota_{ij} \cdot (1 - \lambda_{ij}) > \rho_\Sigma / R_\Sigma$.

(A3) **Identifiability regime.** The action mechanism is known and (C1)–(C3) of Section 5.3 hold; the agent operates in Regime A or Regime B with per-state argmax-correctness probability $p_{\text{id}}(s)$, with floor $p_{\text{id}} := \min_s p_{\text{id}}(s)$.

(A4) **Two-gap diagnostic.** The agent applies the satisfaction-gap / control-regret 2×2 to route corrective action.

(A5) **Piecewise-stationary block structure with restart-on-change identity-tight base learner.** The time axis partitions into $B_T + 1$ stationary intervals on which the MDP is stationary and $\pi_t^*(\cdot \mid s) = \delta_{a_t^*(s)}$ is fixed per state. The base learner is restarted at each block boundary; within each restarted interval of length L it satisfies $\sum_{t=1}^L \mathbb{E}[\overline{\text{TV}}_t] \leq 2c\sqrt{L}$, where $\overline{\text{TV}}_t := \frac{1}{N_h} \sum_{h=0}^{N_h-1} \mathbb{E}_{s_h \sim d_h^{Q_t}} [\text{TV}(\pi_t^*(\cdot \mid s_h), Q_t(\cdot \mid s_h))]$ is the occupancy-weighted per-round coordinate. (A non-restarting carryover variant is plausible but requires bounding cross-block occupancy shift; see Remark in Section 5.)

Then:

(i) **Per-round regret is two-sided identity-bounded.** At every visited state, $\Delta_{\min}(1 - e^{-K_t(s)}) \leq R(Q_t \mid s) \leq V_{\max}(1 - e^{-K_t(s)})$.

334 **(ii) Aggregate strategic mismatch is ultimately bounded.** Under (A2), the modeled mismatch $\|\delta_\Sigma\|$ is
335 ultimately bounded within $R_\Sigma^* = \rho_\Sigma / \mathcal{T}_\Sigma^{\text{bn,ss}}$.

336 **(iii) The KL coordinate is computable with controlled bias.** $-\log Q_t(a_{\text{ag},t}^*(s) \mid s)$ is computable
337 directly from Q_t . Per-state bias: $\mathbb{E}[\|\hat{D}_t(s) - D_t^{\text{true}}(s)\|] \leq (1 - p_{\text{id}}(s)) \log(1/q_0)$ under $Q_t(\cdot \mid s) \geq q_0$
338 at the two argmax candidates. Vanishes in Regime A.

339 **(iv) The 2×2 cell identifies the corrective-action class.** $(\delta_{\text{sat}}, \delta_{\text{regret}})$ routes to revise-policy / revise-
340 model-class-or-horizon / revise-objective.

341 **(v) Trajectory-level blockwise dynamic regret.** Under (A5), $\mathbb{E}[\text{DynReg}(T)] \quad :=$
342 $\mathbb{E}\left[\sum_{t=1}^T (V_t^{\pi_t^*} - V_t^{Q_t})\right] \leq 2c V_{\max} N_h \sqrt{(B_T + 1)T} + N_h(1 - p_{\text{id}}) \log(1/q_0) \cdot T.$

343 The proof composes the per-round identity (i) with the simulation lemma^[40–43], a block decomposition
344 at optimum-change events, Cauchy–Schwarz across the $B_T + 1$ blocks, and the bias bound from (iii).
345 Full proof in Appendix A; the four key lemmas are surfaced in Section 5.

346 4.3 Unpacking the conclusions

347 Theorem 4.1 asserts that the four components compose into a non-stationarity-aware convergence
348 theory with three properties that no existing strand has individually: it *handles non-stationarity*
349 (via (ii)/(A2)), has *explicit metric structure on policy space* (via (i)/Component 2), and is *learnable*
350 *on-policy* (via (iii)/(A3)). Conclusion (iv) is the connective tissue routing corrective action; conclusion
351 (v) is the cumulative regret rate that follows from composing the per-round metric coordinate with the
352 base-learner stationary rate.

353 We emphasize four things about this theorem.

354 *The per-round coordinate is exact, not an inequality.* Under the canonical scope (deterministic π^*), the
355 per-round identity-bound (i) is $\Delta_{\min}/V_{\max}$ -Lipschitz-equivalent to the regret. Both endpoints of the
356 Lipschitz envelope are *achieved* on extremal value landscapes ($\Delta(a) = V_{\max}$ for all $a \neq a^*$ saturates
357 the upper bound; $\Delta_{\min} = \max_{a \neq a^*} \Delta(a)$ saturates the lower). The per-round coordinate is therefore
358 *coordinate-optimal among bounds depending only on TV*: tightening would require value-landscape
359 information beyond TV. This is sharper than Pinsker’s $\sqrt{D_{\text{KL}}/2}$ — exact instead of approximate —
360 and sharper than the Bretagnolle–Huber envelope $\sqrt{1 - e^{-D_{\text{KL}}}}$ at this corner since $x < \sqrt{x}$ on $(0, 1)$.

361 *The cumulative rate is in the piecewise-stationary B_T family, not the continuous-variation V_T fam-*
362 *ily.* The lifted rate $\tilde{O}(N_h^2 \sqrt{SA(B_T + 1)T})$ that follows under UCRL2/UCBVI base learners (Ap-
363 pendix D) sits in the same block-shape family as the restart-on-change variants of^[1] — the route
364 here is through the per-round identity rather than direct optimism. It is *not directly comparable* to
365 the continuous-variation rate $\tilde{O}(SA V_T^{1/3} H T^{2/3})$ of^[3]: the parameter (B_T vs. V_T), the T -scaling
366 ($T^{1/2}$ vs. $T^{2/3}$), the S, A -scaling ($\sqrt{SA}$ vs. $(SA)^{1/3}$), and the horizon dependence (N_h^2 vs. N_h) all
367 differ. Under abrupt-magnitude- δ piecewise-stationarity the two regimes coincide up to constants
368 ($V_T \asymp \delta B_T$), and B-CS1 is sharper for fixed B_T ; outside that, Mao-style continuous-variation analysis
369 is sharper for fixed V_T . Generalizing the per-round-identity route to the continuous-variation regime
370 is open (Section 6).

371 *The bias term in (v) is well-controlled in Regimes A/B and saturates in Regime C.* The second term
372 $N_h(1 - p_{\text{id}}) \log(1/q_0) \cdot T$ is linear in T — vacuous as a rate. It vanishes in Regime A ($p_{\text{id}} \rightarrow 1$ under
373 (C1)–(C3)). In Regime B ($p_{\text{id}} \in (0, 1)$), the bias is controlled but linear; the first term ($\sqrt{(B_T + 1)T}$)
374 is sublinear, so the dominant rate term is the metric one provided $1 - p_{\text{id}}$ is small enough. In Regime
375 C ($p_{\text{id}} \rightarrow 0$), the bound is provable in principle but vacuous as an actionable rate; algorithm flags this
376 regime as out-of-scope for on-policy learnability (Appendix D).

377 *The horizon factor is the standard simulation-lemma penalty.* The N_h in conclusion (v) is the linear-
378 in- N_h per-round simulation-lemma coefficient — sharper than the N_h^2 of TRPO-style worst-state
379 bounds^[44]. The lifted N_h^2 in the UCRL2/UCBVI cumulative rate reflects the additional N_h in the
380 within-block base-learner constant c .

381 4.4 Necessity of the four components

382 A natural reviewer question: why exactly four? The composition is a *bundle of compatible guarantees*
 383 (each conclusion supported by its own assumptions), not a single integrated theorem in which every
 384 component is necessary for the rate. The rate (v) goes through (A5) and (i) alone; the joint properties
 385 — rate plus persistence plus learnability plus corrective-action routing — require all four.

386 **Without Component 2 (point-mass identity):** the per-round bound is Pinsker $V_{\max} \sqrt{D_{\text{KL}}/2}$ or BH
 387 $V_{\max} \sqrt{1 - e^{-D_{\text{KL}}}}$ — both strictly looser at this corner; Pinsker is vacuous for $D_{\text{KL}} > 2$. The metric
 388 layer loses exactness, and the “behavior-cloning loss against optimal trajectory” interpretation of the
 389 cumulative coordinate vanishes.

390 **Without Component 3 (strategic tempo):** the agent has no persistence guarantee. By the structural-
 391 class theorem on $\mathcal{A}_{\text{decay}}$, every count-accumulating update without forgetting eventually violates
 392 the threshold for any positive disturbance; the modeled mismatch $\|\delta_{\Sigma}\|$ grows unboundedly. The
 393 per-round bound (i) still holds, but $K_t(s)$ may itself drift — there is no guarantee the per-round
 394 coordinate stays bounded.

395 **Without Component 4 (loop-Level-2):** the bound’s RHS is computable from Q_t , but its interpreta-
 396 tion as regret-against-true-optimum requires identifying $a_t^*(s)$ — an interventional query under the
 397 environment’s transition kernel. In Regime C ($\iota \approx 0$), the bias saturates at $\log(1/q_0)$ per state per
 398 step, contributing $N_h \log(1/q_0) \cdot T$ to cumulative regret — the rate (v) becomes vacuous.

399 **Without Component 1 (two-gap diagnostic):** the regret signal δ_{regret} alone does not distinguish
 400 “policy revision is the right move” from “the goal is unattainable from M_t .” An agent might persist in
 401 policy revision when the model, policy class, or horizon is the actual problem (Capability-limit cell),
 402 or revise the goal when policy revision would have sufficed.

403 So while the rate (v) goes through (A5) and (i) alone, the *bundle* — rate plus persistence plus
 404 learnability plus corrective-action routing — requires all four. Each drop is a concrete loss of an
 405 epistemic property that the field treats as separate.

406 5 Mechanism

407 This section overviews the proof of Theorem 4.1 via four key lemmas chained through a block
 408 decomposition. Each lemma is stated formally; intuitions are given in plain English with the algebra
 409 deferred to the appendices. The rate (v) is the composition’s load-bearing conclusion; we sketch its
 410 derivation in four steps after surfacing the lemmas.

411 5.1 Key Lemma 1: Point-mass reverse-KL/TV identity

412 **Lemma 5.1** (Point-mass reverse-KL/TV identity). *For deterministic $\pi^* = \delta_{a^*}$ and any policy*
 413 *Q with $Q(a^*) > 0$, $D_{\text{KL}}(\pi^* \| Q) = -\log Q(a^*) = -\log(1 - \text{TV}(\pi^*, Q))$, equivalently*
 414 *$\text{TV}(\pi^*, Q) = 1 - e^{-D_{\text{KL}}(\pi^* \| Q)}$.*

415 *Intuition.* The reverse-KL collapses under the point-mass: $D_{\text{KL}}(\delta_{a^*} \| Q) = \sum_a \delta_{a^*}(a) \log \frac{\delta_{a^*}(a)}{Q(a)}$,
 416 where the convention $0 \log 0 = 0$ kills every $a \neq a^*$ term, leaving $\log \frac{1}{Q(a^*)} = -\log Q(a^*)$. Total
 417 variation collapses just as cleanly: $\text{TV}(\delta_{a^*}, Q) = \frac{1}{2} \sum_a |\delta_{a^*}(a) - Q(a)| = 1 - Q(a^*)$. Combining
 418 gives the identity. Substituting into the Bretagnolle–Huber inequality $\text{TV}(P, Q) \leq \sqrt{1 - e^{-D_{\text{KL}}(P \| Q)}}$
 419 yields $1 - e^{-D_{\text{KL}}} \leq \sqrt{1 - e^{-D_{\text{KL}}}}$, strictly on $D_{\text{KL}} > 0$. The identity is therefore not BH-at-equality;
 420 it is the *exact value* at the deterministic- π^* corner, sitting strictly below the BH envelope. Full proof
 421 in Appendix B.

422 *Per-round regret bound.* Composed with the textbook TV-regret bound $R(Q) \leq V_{\max} \text{TV}(\pi^*, Q)$
 423 and the matching action-gap lower bound $R(Q) \geq \Delta_{\min} \text{TV}(\pi^*, Q)$, the identity yields the two-sided
 424 characterization that gives Theorem 4.1’s conclusion (i). The bound is *coordinate-optimal among*
 425 *bounds depending only on TV*: any tighter bound on $R(Q)$ requires value-landscape information
 426 beyond TV.

427 *Perturbative extension.* The deterministic- π^* scope is not a hard wall. For ϵ -greedy π_ϵ^* with Q
428 satisfying full-support $Q(a) \geq q_0$, $\text{TV}(\pi_\epsilon^*, Q) = 1 - e^{-D_{\text{KL}}(\pi_\epsilon^* \| Q)} + O(\epsilon \log(1/\epsilon))$ uniformly in
429 Q ; the proof is a triangle-inequality + expansion argument given in Appendix B.2. For softmax-
430 regularized $\pi_\tau^* \propto \exp(Q_O/\tau)$, the off-optimum mass concentrates as $O(\exp(-\Delta_{\min}/\tau))$ and the
431 correction is $O(\tau^{-1} \exp(-\Delta_{\min}/\tau))$ — exponentially small with polynomial prefactor. The two-sided
432 regret bound transfers with the same correction order. Outside the perturbative regime — genuinely
433 high-entropy optima, tied-optimum, hard-exploration — the BH inequality becomes the relevant
434 general bound; Pinsker is the textbook fallback.

435 5.2 Key Lemma 2: Forgetting prerequisite

436 The strategic mismatch state $\delta_\Sigma = (\delta_{ij})_{(i,j) \in E}$ in the Euclidean norm $\|\delta_\Sigma\|^2 = \sum_{(i,j)} \delta_{ij}^2$ evolves
437 in discrete time as $\mathbb{E}[\Delta \delta_{ij} \mid \delta_{ij}] = -\alpha_{ij} \delta_{ij} + w_{ij}$, $|w_{ij}| \leq \rho_\Sigma / |E|^{1/2}$, where $\alpha_{ij} \geq 0$ is
438 a per-element correction rate (the strategic sector parameter), w_{ij} is a bounded disturbance, and
439 the disturbance budget satisfies $\sum_{(i,j)} w_{ij}^2 \leq \rho_\Sigma^2$. The strategic reserve R_Σ is the operative norm-
440 radius beyond which the agent loses tracking. Beta-Bernoulli edge updates instantiate this *Model (S)*
441 approximately; gradient-based decaying-step updates instantiate it under standard Robbins–Monro
442 conditions; we will refer to it as *Model (S)* throughout.

443 **Lemma 5.2** (Multi-factor forgetting prerequisite). *Within Model (S) with diagonal per-element*
444 *correction rates $\alpha_{ij} = \nu_{ij} \cdot \iota_{ij} \cdot \eta_{\text{edge},ij}$ (per-element observation rate $\nu_{ij} \in [0, 1]$, regime-adjusted*
445 *identifiability $\iota_{ij} \in [0, 1]$, edge-update gain $\eta_{\text{edge},ij} > 0$), under per-element exponential forgetting*
446 *at rate $\lambda_{ij} \in (0, 1)$ giving steady-state $\eta_{\text{edge},ij}^{\text{ss}} \approx 1 - \lambda_{ij}$, ultimate boundedness of $\|\delta_\Sigma\|$ within*
447 *$R_\Sigma^* = \rho_\Sigma / \mathcal{T}_\Sigma^{\text{bn,ss}}$ is sufficient under $\mathcal{T}_\Sigma^{\text{bn,ss}} := \min_{(i,j) \in E} (\nu_{ij} \cdot \iota_{ij} \cdot (1 - \lambda_{ij})) > \rho_\Sigma / R_\Sigma$. The*
448 *threshold is sharp inside the diagonal sector model: when the inequality reverses, an adversarial*
449 *disturbance concentrating on the bottleneck element drives $\|\delta_\Sigma\|$ beyond R_Σ within the modeled*
450 *dynamics.*

451 *Intuition (why bottleneck, not sum).* The Lyapunov sector product is $\delta_\Sigma^\top \mathbf{F}(\delta_\Sigma) = \sum_{(i,j)} \alpha_{ij}^{\text{ss}} \delta_{ij}^2$.
452 Bounding this from below uniformly in δ_Σ requires the largest constant c such that $\sum_{(i,j)} \alpha_{ij}^{\text{ss}} \delta_{ij}^2 \geq$
453 $c \cdot \|\delta_\Sigma\|^2$ holds for all δ_Σ . The largest such c is $\min_{(i,j)} \alpha_{ij}$, not $\sum_{(i,j)} \alpha_{ij}$ — adversarial concentration
454 of the disturbance on the *weakest* element saturates the bound at the bottleneck. The same effect
455 makes per-dimension persistence conditions sharper than scalar conditions in adaptive control^[45]:
456 scalar tempo can overstate the available margin because the worst-case disturbance concentrates on
457 the slowest-correcting dimension. Ultimate boundedness then follows from Khalil’s classical sector-
458 Lyapunov template applied to the strategic substate. Full proof in Appendix B; the structural-class
459 theorem on $\mathcal{A}_{\text{decay}}$ (gain-decay updates universally violate the threshold) is in Appendix C.6.

460 *ProST as a sector-level reduction.* The block-restart structure of ProST^[8] idealizes as an *impulsive*
461 *system* — between scheduled update times the policy is held and the mismatch evolves under dis-
462 turbance alone (continuous-destabilizing in the Lyapunov $V = \|\delta_\Sigma\|^2$); at each update time the
463 policy is revised, contracting the modeled mismatch by per-update impulse gain $\gamma \in (0, 1]$. The
464 reverse-average-dwell-time framework of^[46] (Theorem 1, reverse-ADT branch) gives the threshold
465 $\Delta_{\max} \cdot \rho_\Sigma / R_\Sigma < -\ln(1 - \gamma)^2$ where $\Delta_{\max}$ is the longest inter-update gap. Under uniform blocks
466 $\Delta_i = T/K$ this reduces to $K/T > \rho_\Sigma / (2\gamma R_\Sigma)$ to leading order in γ — recovering ProST’s K/T
467 form with the impulse gain made explicit. Under nonuniform schedules the impulsive form is *strictly*
468 *stronger*: the longest block sets the threshold, not the average. This makes a tradeoff between update
469 frequency $1/\Delta_{\max}$ and per-update impulse strength γ visible at the sector-Lyapunov level — the same
470 disturbance can be tolerated either by frequent weak updates or rare strong ones.

471 5.3 Key Lemma 3: Closed-loop interventional access

472 **Lemma 5.3** (Loop generates interventional samples). *Let H_t denote the agent’s information state*
473 *at round t (model state M_t , history of observations and actions). Under (C1) positivity — $\pi_t(a \mid$
474 $H_t) \geq p_{\min} > 0$ on the support of H_t for the actions of interest; (C2) sequential ignorability —
475 H_t blocks unobserved confounding of the action mechanism, so $a_t \perp\!\!\!\perp (\text{environment latents}) \mid H_t$;
476 (C3) known action mechanism — the behavior policy $\pi_t(a \mid H_t)$ is known: executed actions are
477 interventions relative to the environment’s transition kernel, and the loop generates samples from*

478 $P(o_{t+1} \mid \text{do}(a_t = a), H_t)$ — i.e., *Pearl Level-2 interventional kernels conditional on the decision*
 479 *context*.

480 *Intuition.* By temporal ordering and the singular-trajectory commitment of Section 3, a_t causally
 481 precedes o_{t+1} . Under (C1)–(C3), conditioning on the information state H_t blocks the action-selection
 482 confounding pathway, so the conditional distribution $P(o_{t+1} \mid a_t, H_t)$ coincides with the interven-
 483 tional distribution $P(o_{t+1} \mid \text{do}(a_t), H_t)$ on the support where $\pi_t(a_t \mid H_t) > 0$. The architecture
 484 of the agent (Q-learning, transformer, etc.) does not enter: the loop generates conditional-Level-2
 485 samples whenever (C1)–(C3) hold, irrespective of how the agent updates internally. Whether the
 486 agent *exploits* this data for Level-2 reasoning is a separate algorithmic question.

487 *Three regimes of identifiability.* Three regimes^[29] partition usable strength of identification: - **Regime**
 488 **A (intervention-rich, $\iota \approx 1$).** Software tests, controlled labs. do-effects identified cleanly; bound
 489 realizable on-policy. - **Regime B (partial, $\iota \in (0, 1)$).** Mixed observation-intervention. Bound
 490 holds for the model the agent identifies; bias $\propto 1 - \iota$. - **Regime C (observation-only, $\iota \approx 0$).**
 491 Passive monitoring. Bound provable analytically but not realizable on-policy; observer-on-sub-agent
 492 extensions can recover identifiability in subcases (Appendix C.8).

493 Action-generated data is Level 2 in *character*, but a clean estimate of $P(o \mid \text{do}(a), H_t)$ requires
 494 overcoming four typical obstacles: coverage (diverse actions tried), within-step confounding, delay
 495 (consequences past $t + 1$), and partial observability. We honor this distinction. Components 2 and
 496 4 are jointly load-bearing: the identity of Section 5.1 gives the analytic regret-KL relationship; the
 497 loop-Level-2 access here gives the data substrate from which π^* — and the bound’s RHS — can be
 498 empirically estimated under sufficient identifiability. Neither alone makes the bound learnable.

499 5.4 Key Lemma 4: Bias bound under partial identifiability

500 **Lemma 5.4** (Bias bound). Assume $Q_t(a \mid s) \geq q_0 > 0$ for both $a = a_{\text{ag},t}^*(s)$ (the agent’s identified
 501 optimum) and $a = \tilde{a}_t^*(s)$ (the true optimum). Let $\hat{D}_t(s) := -\log Q_t(a_{\text{ag},t}^*(s) \mid s)$ and $D_t^{\text{true}}(s) :=$
 502 $-\log Q_t(\tilde{a}_t^*(s) \mid s)$. Then $|\hat{D}_t(s) - D_t^{\text{true}}(s)| \leq \mathbf{1}[a_{\text{ag},t}^*(s) \neq \tilde{a}_t^*(s)] \cdot \log(1/q_0)$, and taking expecta-
 503 tions: $\mathbb{E}[|\hat{D}_t(s) - D_t^{\text{true}}(s)|] \leq (1 - p_{\text{id}}(s)) \log(1/q_0)$, where $p_{\text{id}}(s) := \Pr[a_{\text{ag},t}^*(s) = \tilde{a}_t^*(s)]$.

504 *Intuition.* The argument is an indicator decomposition. On the matching event, the difference is
 505 identically zero. On the misidentification event, $\log Q_t(a_{\text{ag},t}^*(s) \mid s)$ and $\log Q_t(\tilde{a}_t^*(s) \mid s)$ both lie
 506 in $[\log q_0, 0]$, so their absolute difference is at most $\log(1/q_0)$. The expectation bound follows from
 507 $\Pr[\text{misidentification event}] = 1 - p_{\text{id}}(s)$. Bias scales *linearly* in misidentification mass — vanishes
 508 in Regime A ($p_{\text{id}} \rightarrow 1$), controlled in Regime B, saturates at $\log(1/q_0)$ in Regime C. The argmax-
 509 correctness probability p_{id} relates to the per-edge identifiability ι via the policy DAG structure; in
 510 the simple-bandit case $p_{\text{id}} = \iota$. Note: the support condition is *two-point* (only at the two argmax
 511 candidates), strictly weaker than the *full-support* condition $Q(a) \geq q_0$ for all a that the perturbative
 512 identity (Section 5.1) requires — F.1 applies whenever the two candidates are positively-supported
 513 regardless of Q_t ’s mass on other actions. Full proof in Appendix B.

514 5.5 Combining the lemmas: cumulative regret in four steps

515 We sketch the proof of conclusion (v) of Theorem 4.1 — the cumulative trajectory-level dynamic
 516 regret rate — by chaining the four key lemmas through a block decomposition.

517 *Step 1 — Per-round bound via the simulation lemma.* For finite-horizon N_h non-stationary MDPs with
 518 bounded reward, the standard performance-difference / simulation lemma^[40–43] gives, for any two
 519 policies π_t^* and Q_t and initial state s_0 , $V_t^{\pi_t^*}(s_0) - V_t^{Q_t}(s_0) \leq V_{\max} \sum_{h=0}^{N_h-1} \mathbb{E}_{s_h \sim d_h^{Q_t}} [\text{TV}(\pi_t^*(\cdot \mid$
 520 $s_h), Q_t(\cdot \mid s_h))]$ = $V_{\max} \cdot N_h \cdot \overline{\text{TV}}_t$, where $d_h^{Q_t}$ is the round- t learner-induced state distribution and
 521 $\overline{\text{TV}}_t$ is the occupancy-weighted TV defined in (A5). By Key Lemma 1 applied per state, $\text{TV}(\pi_t^*(\cdot \mid$
 522 $s), Q_t(\cdot \mid s)) = 1 - e^{-K_t(s)}$ — strictly sharper than Pinsker — preserved underneath the occupancy
 523 aggregation. The horizon factor N_h here is the linear-in- N_h simulation-lemma penalty, sharper than
 524 the N_h^2 of TRPO-style worst-state bounds^[44].

525 *Step 2 — Block decomposition at optimum-change events.* Partition $[1, T]$ at optimum-change events
 526 into $0 = \tau_0 < \tau_1 < \dots < \tau_{B_T} \leq \tau_{B_T+1} = T$. Within each block $[\tau_i, \tau_{i+1})$ of length Δ_i

the MDP is stationary. (A5) — restart-on-change base learner — gives within-block guarantee $\sum_{t \in \text{block}_i} \mathbb{E}[\overline{\text{TV}}_t] \leq 2c\sqrt{\Delta_i}$.

Step 3 — Cauchy–Schwarz across blocks. Across the $B_T + 1$ blocks, $\sum_i \sqrt{\Delta_i} \leq \sqrt{(B_T + 1)T}$ by Cauchy–Schwarz. Multiplying by $V_{\max}N_h$ closes the rate term: $\sum_t \mathbb{E}[V_{\max}N_h \overline{\text{TV}}_t] \leq 2cV_{\max}N_h \sqrt{(B_T + 1)T}$.

Step 4 — Bias term via Key Lemma 4. The bias from misidentification at each visited state per horizon step contributes at most $(1 - p_{\text{id}}(s)) \log(1/q_0)$ per state-step (Key Lemma 4, per-state form). Aggregated linearly over the N_h horizon steps and the T rounds with the per-state floor $p_{\text{id}} := \min_s p_{\text{id}}(s)$, this contributes $N_h(1 - p_{\text{id}}) \log(1/q_0) \cdot T$ — vanishes in Regime A.

Combining Steps 1–4 gives conclusion (v): $\mathbb{E}[\text{DynReg}(T)] \leq 2cV_{\max}N_h \sqrt{(B_T + 1)T} + N_h(1 - p_{\text{id}}) \log(1/q_0) \cdot T$. The Cesàro tracking corollary $\frac{1}{T} \sum_t (V_t^* - V(\pi_t)) = O(N_h \sqrt{(B_T + 1)/T}) \rightarrow 0$ (when $B_T = o(T)$ and $p_{\text{id}} \rightarrow 1$) follows by dividing through by T .

Bandit case sharpening In the bandit special case ($N_h = 1$), Thompson sampling and UCB satisfy (A5) with the *stronger logarithmic rate* $\mathbb{E}[1 - e^{-K_t}] = O(\log t / (t \Delta_{\min}))$ per stationary block (Appendix D), giving cumulative regret $O(V_{\max}(B_T + 1) \log^2(T/(B_T + 1)) / \Delta_{\min})$ — sharper than $\sqrt{(B_T + 1)T}$ when $\Delta_{\min}$ is bounded away from zero. The square-root rate is the worst-case Cauchy–Schwarz bound; the logarithmic rate is the typical-case bound for stochastic-bandit base learners.

Carryover variant of (A5) If the base learner does *not* restart at block boundaries — e.g., a single UCRL2 instance run across the whole horizon — the within-block $\sqrt{L}$ guarantee can fail at switch events because $d_h^{Q_t}$ at the start of block $i + 1$ reflects the previous block’s policy rather than a fresh exploration distribution. Bounding the cross-block occupancy shift requires either (a) detecting switches and locally re-exploring (which recovers a restart-style analysis), or (b) showing the base learner’s regret is robust to mis-calibrated state distributions. Sliding-window or adaptive-window variants^[1,2] can be analyzed this way under additional structural assumptions. Theorem 4.1(v) is stated for the restart case; the carryover case is plausible but requires bounding the occupancy shift, deferred as future work.

Detailed full proof of (v) in Appendix A.

6 Conclusion

We assemble four components — the two-gap diagnostic, the point-mass reverse-KL/TV identity, the multi-factor strategic-tempo forgetting prerequisite, and closed-loop interventional access — into a unified non-stationary RL convergence theory. The composition delivers three properties no published strand has individually: it *handles non-stationarity* (via the strategic-tempo persistence inequality), has *explicit metric structure on policy space* (via the point-mass identity), and is *learnable on-policy* (via closed-loop interventional access). The cumulative dynamic regret rate $\tilde{O}(N_h \sqrt{(B_T + 1)T})$ matches the piecewise-stationary B_T family of^[1] as a corollary; the per-round identity coordinate is *exact* under deterministic π^* and strictly sharper than Pinsker and Bretagnolle–Huber at this corner. The point-mass identity is the technical anchor — absent, to our knowledge, from the prior RL/non-stationary corpus as an upper-bound regret coordinate.

We close with concluding observations and practitioner takeaways.

6.1 Concluding observations

On the optimal dependencies on N_h , S , A The lifted UCRL2/UCBVI rate $\tilde{O}(N_h^2 \sqrt{SA(B_T + 1)T})$ has linear-in- N_h per-round penalty (the simulation lemma) plus an additional N_h from the within-block constant c . Bernstein-type concentration in the base learner could potentially shave one $\sqrt{N_h}$ factor, by analogy with^[43] for tabular non-stationary MDPs. We believe the $\sqrt{N_h}$ gap between this and the corresponding tabular non-stationary lower bound is expected because the base-learner concentration used here is intrinsically Hoeffding-type.

574 **On the continuous-variation extension** The current rate is in the piecewise-stationary B_T fam-
 575 ily. Lifting to the continuous-variation V_T family of^[3] would require a different block structure —
 576 variation-aware windows with adaptive width^[2] — and a different analysis at the per-round level (con-
 577 tinuous variation cannot be partitioned at discrete optimum-change events). It remains an interesting
 578 open question whether the per-round-identity route can match the $T^{2/3}$ continuous-variation rate, or
 579 whether the regimes are structurally distinct.

580 **On joint uniqueness of the four-component composition** The metric layer is uniquely reverse-KL
 581 up to positive scaling under chain-rule additivity (Theorem C.3 in Appendix C.10) — any theory
 582 satisfying explicit metric structure plus chain-rule must use reverse-KL. We do *not* claim joint
 583 uniqueness across all three properties (non-stationarity / metric / learnable). Alternative correction
 584 architectures (non-sector-bounded revisions; non-diagonal correction) and alternative interventional-
 585 access taxonomies exist in principle, and a joint-uniqueness theorem would require further axioms
 586 (singular-trajectory + sector-boundedness + identifiability conditions) — open.

587 **On coupled-goal architectures** (C2)’s sequential-ignorability assumption presupposes architec-
 588 tural separation (conditional independence) between the model state M_t and the goal state. Goal-
 589 conditioned LLM policies and similar coupled-goal architectures need additional machinery to satisfy
 590 (C2); the conditional-independence relaxation results of^[47] are a starting point. We treat this as an
 591 interesting boundary case that the current framework does not directly cover.

592 **On strict tightness of the strategic-tempo threshold** The bottleneck threshold $\mathcal{T}_{\Sigma}^{\text{bn,ss}} > \rho_{\Sigma}/R_{\Sigma}$ is
 593 *sufficient and sharp inside the diagonal sector model* — when reversed, an adversarial disturbance
 594 concentrating on the bottleneck element drives the modeled mismatch beyond R_{Σ} . The theorem
 595 is silent about correction architectures *outside* the sector-Lyapunov reduction (e.g., non-diagonal
 596 correction, alternative non-sector-bounded revisions, or stabilization produced by mechanisms other
 597 than the modeled α_{Σ} -correction). Whether the threshold extends to those architectures is a useful
 598 open question in the adaptive-control-RL bridge.

599 6.2 Practitioner takeaways

600 For a practitioner seeking to apply the four-component theory:

- 601 1. **The point-mass identity at the deterministic- π^* corner is a corner-case sharpening**
 602 **of Bretagnolle–Huber that hasn’t been deployed in RL regret upper bounds.** Where
 603 deterministic π^* applies — most tabular finite-MDP regret analyses with isolated optima —
 604 the identity yields *exact* regret-vs-KL Lipschitz equivalence rather than the loose Pinsker /
 605 BH inequalities. The per-round coordinate $1 - e^{-D_{\text{KL}}}$ is computable directly from Q and is
 606 strictly sharper than $\sqrt{D_{\text{KL}}/2}$ at every $D_{\text{KL}} > 0$.
- 607 2. **The multi-factor strategic tempo lifts ProST’s tempo-as-hyperparameter framing into**
 608 **a structural survival inequality.** Three factors (observation rate ν , identifiability ι , discount
 609 rate $1 - \lambda$) are independently load-bearing per element. A practitioner can use the aggregate
 610 necessary-condition $\mathcal{T}_{\Sigma}^{\text{agg,ss}} > |E| \cdot \rho_{\Sigma}/R_{\Sigma}$ as a fail-fast pre-check and the variation-budget
 611 estimation literature^[1,2,39] to estimate the environment-side ρ_{Σ}/R_{Σ} from observed reward
 612 and transition variation.
- 613 3. **Closed-loop interventional access makes regret bounds learnable on-policy under suffi-**
 614 **cient identifiability.** Regime A (intervention-rich, $\iota \approx 1$) makes the bound directly realizable;
 615 Regime B (partial) carries quantified bias linear in $1 - p_{\text{id}}$; Regime C (observation-only)
 616 requires observer-on-sub-agent extensions (Appendix C.8) or flags the regime as out-of-
 617 scope. The 2×2 diagnostic routes corrective action when regret is non-zero — distinguishing
 618 “policy revision” from “model / policy class / horizon revision” from “objective revision.”

619 We hope this work contributes a unifying lens for the four-track non-stationary RL landscape, and
 620 that the point-mass identity is a useful tool in subsequent regret analyses where deterministic π^* is
 621 the canonical scope.

References

- [1] W. C Cheung, D Simchi-Levi, and R Zhu. Reinforcement learning for non-stationary markov decision processes: the blessing of (more) optimism. *arXiv:2006.14389*, 2020.
- [2] Chen-Yu Wei and Haipeng Luo. Non-stationary reinforcement learning without prior knowledge: An optimal black-box approach, 2021.
- [3] W Mao, K Zhang, R Zhu, D Simchi-Levi, and T Başar. Near-optimal model-free rl in non-stationary episodic mdps. In *ICML*, 2021.
- [4] P Gajane, R Ortner, and P Auer. Variational regret bounds for reinforcement learning. In *UAI*, 2019.
- [5] Long-Fei Li, Peng Zhao, and Zhi-Hua Zhou. Dynamic regret of adversarial MDPs with unknown transition and linear function approximation. In *Proceedings of the AAAI Conference on Artificial Intelligence*, volume 38, 2024. doi: 10.1609/aaai.v38i12.29261.
- [6] Y Fei, Z Yang, Z Wang, and Q Xie. Dynamic regret of policy optimization in non-stationary environments. *arXiv:2007.00148*, 2020.
- [7] F. E Stradi, A Lunghi, M Castiglioni, A Marchesi, and N Gatti. Learning cmdps with non-stationary rewards and constraints. *arXiv:2405.14372*, 2024.
- [8] Hyunin Lee, Yuhao Ding, Jongmin Lee, Ming Jin, Javad Lavaei, and Somayeh Sojoudi. Tempo adaptation in non-stationary reinforcement learning. In *Advances in Neural Information Processing Systems 36 (NeurIPS 2023)*, 2023. arXiv 2309.14989; OpenReview 7R8noSP4vL.
- [9] Hyunin Lee, Ming Jin, Javad Lavaei, and Somayeh Sojoudi. Pausing policy learning in non-stationary reinforcement learning. In *Proceedings of the 41st International Conference on Machine Learning (ICML 2024)*, Proceedings of Machine Learning Research, 2024. arXiv 2405.16053.
- [10] A Touati and P Vincent. Efficient learning in non-stationary linear mdps (opt-wlsvi). *arXiv:2010.12870*, 2020.
- [11] Y Russac, C Vernade, and O Cappé. Weighted linear bandits for non-stationary environments. In *NeurIPS*, 2019.
- [12] A Garivier and E Moulines. On upper-confidence bound policies for non-stationary bandit problems. *arXiv:0805.3415*, 2008.
- [13] Junzhe Zhang and E Bareinboim. Mdps with unobserved confounders: a causal approach. (*Tech. report; Columbia.*), 2016.
- [14] Junzhe Zhang and Elias Bareinboim. Online reinforcement learning for mixed policy scopes. In *Advances in Neural Information Processing Systems*, 2022.
- [15] Y Lu, A Meisami, and A Tewari. Causal mdps: learning good interventions efficiently. *arXiv:2102.07663*, 2021.
- [16] Y Lu, A Meisami, and A Tewari. Efficient rl with prior causal knowledge. *CLEaR*, 2022.
- [17] L Wang, Z Yang, and Z Wang. Provably efficient causal rl with confounded observational data (dovi). In *NeurIPS*. *arXiv:2006.12311 (2020)*, 2021.
- [18] Junzhe Zhang. Designing optimal dynamic treatment regimes: a causal rl approach. In *ICML*, 2020.
- [19] D Russo and B Van Roy. An information-theoretic analysis of thompson sampling. *J. Mach. Learn. Res.*, 2014.
- [20] Daniel Russo and Benjamin Van Roy. Learning to optimize via information-directed sampling. *Operations Research*, 2014.

- [21] X Lu and B Van Roy. Information-theoretic confidence bounds for reinforcement learning. In *NeurIPS*, 2019.
- [22] Seungki Min and Daniel Russo. An information-theoretic analysis of nonstationary bandit learning. In *International Conference on Machine Learning (ICML 2023)*, 2023.
- [23] T Lattimore and A György. Mirror descent and the information ratio. In *COLT. (PMLR vol 134)*; *arXiv:2009.12228 (2020)*, 2021.
- [24] Sham Kakade, Akshay Krishnamurthy, Kendall Lowrey, Motoya Ohnishi, and Wen Sun. Information theoretic regret bounds for online nonlinear control. In *Advances in Neural Information Processing Systems*, 2020. arXiv:2006.12466.
- [25] Jean Bretagnolle and Catherine Huber. Estimation des densités: Risque minimax. In C. Del-lacherie, P. A. Meyer, and M. Weil, editors, *Séminaire de Probabilités XII*, volume 649 of *Lecture Notes in Mathematics*, pages 342–363. Springer, 1978. doi: 10.1007/BFb0064609.
- [26] Alexandre B. Tsybakov. *Introduction to Nonparametric Estimation*. Springer Series in Statistics. Springer, New York, 2009. doi: 10.1007/b13794.
- [27] T Lattimore and C Szepesvári. *Bandit Algorithms*. Cambridge., 2020.
- [28] Hassan K. Khalil. *Nonlinear Systems*. Prentice-Hall, Upper Saddle River, NJ, 3rd edition, 2002.
- [29] Elias Bareinboim, Juan D. Correa, Duligur Ibeling, and Thomas Icard. On Pearl’s hierarchy and the foundations of causal inference. In Hector Geffner, Rina Dechter, and Joseph Y. Halpern, editors, *Probabilistic and Causal Inference: The Works of Judea Pearl*, pages 507–556. ACM Books, 2022. doi: 10.1145/3501714.3501743.
- [30] Karl Friston, Thomas FitzGerald, Francesco Rigoli, Philipp Schwartenbeck, and Giovanni Pezzulo. Active inference: a process theory. *Neural Computation*, 29(1):1–49, 2017.
- [31] T Parr, G Pezzulo, and K. J Friston. *Active Inference: The Free Energy Principle in Mind, Brain, and Behavior*. MIT Press., 2022.
- [32] Wang Chi Cheung, David Simchi-Levi, and Ruihao Zhu. Reinforcement learning for non-stationary Markov decision processes: The blessing of (more) optimism. *Management Science*, 68(3):1697–1716, 2022. doi: 10.1287/mnsc.2021.4123. ArXiv 2020.
- [33] Y Yang, Z Zheng, M Tomizuka, C Liu, and S Li. The feasibility theory of constrained reinforcement learning: A tutorial study. *arXiv:2404.10064*, 2024.
- [34] Oliver Schulte and Pascal Poupart. When should reinforcement learning use causal reasoning? *Transactions on Machine Learning Research*, 2025. Published 2025-05-09; OpenReview D1PPuk8ZBI; Survey Certification.
- [35] J Pearl. *Causality: Models, Reasoning, and Inference*. (2nd ed.). Cambridge., 2009.
- [36] C Canonne. A short note on an inequality between kl and tv. *arXiv:2202.07198*, 2022.
- [37] H Hajiabolhassan and R Ortner. Online regret bounds for satisficing in markov decision processes. *Mathematics of Operations Research*, 2025.
- [38] Y. Zhang, R Zhu, and Q Xie. On the peril of (even a little) non-stationarity in satisficing regret minimization. (*contemporaneous*; March 2026), 2026.
- [39] A Gerogiannis, Huang, Y.-H, and V Veeravalli. Darling: Detection augmented reinforcement learning with non-stationary guarantees. (*contemporaneous*; April 2026), 2026.
- [40] S Kakade and J Langford. Approximately optimal approximate reinforcement learning (cpi / performance-difference lemma). In *ICML*, 2002.
- [41] R Munos. Error bounds for approximate policy iteration. In *ICML*, 2003.
- [42] S Ross and J. A Bagnell. Efficient reductions for imitation learning (dagger). In *AISTATS*, 2010.

- [43] M. G Azar, I Osband, and R Munos. Minimax regret bounds for reinforcement learning (ucbvi). In *ICML. arXiv:1703.05449*, 2017.
- [44] J Schulman, S Levine, P Moritz, M. I Jordan, and P Abbeel. Trust region policy optimization (trpo). In *ICML. arXiv:1502.05477*, 2015.
- [45] Brian D. O. Anderson. Adaptive systems, lack of persistency of excitation and bursting phenomena. *Automatica*, 21(3):247–258, 1985. doi: 10.1016/0005-1098(85)90058-5.
- [46] J. P Hespanha, D Liberzon, and A. R Teel. Lyapunov conditions for input-to-state stability of impulsive systems. *Automatica*. 44(11):2735–2744, 2008.
- [47] Jelle Bruineberg, Krzysztof Dolega, Joe Dewhurst, and Manuel Baltieri. The Emperor’s new Markov blankets. *Behavioral and Brain Sciences*, 45:e183, 2022. doi: 10.1017/S0140525X21002351.
- [48] S Levine. Reinforcement learning and control as probabilistic inference. *arXiv:1805.00909*, 2018.
- [49] N Wiener. *Cybernetics*. MIT Press., 1948.
- [50] R Conant and W. R Ashby. Every good regulator of a system must be a model of that system. *Int. J. Systems Sci.*, 1970.
- [51] J Aczél and Z Daróczy. *On Measures of Information and Their Characterizations*. Academic Press., 1975.
- [52] P Auer, T Jaksch, and R Ortner. Near-optimal regret bounds for reinforcement learning (ucrl2). *J. Mach. Learn. Res.*, 11:1563–1600, 2010.

A Proof of the Composition Theorem

This appendix gives the full proof of Theorem 4.1. The proof composes the four key lemmas of Section 5 (full proofs of those in Appendix B) with a block decomposition at optimum-change events plus Cauchy–Schwarz across blocks.

Notation Throughout this section: $K_t(s) := D_{\text{KL}}(\delta_{a_t^*(s)} \parallel Q_t(\cdot \mid s))$ is the per-state reverse-KL coordinate; $d_h^{Q_t}(\cdot \mid s_0)$ is the round- t learner-induced state distribution at horizon step h starting from s_0 ; $\overline{\text{TV}}_t := \frac{1}{N_h} \sum_{h=0}^{N_h-1} \mathbb{E}_{s_h \sim d_h^{Q_t}} [\text{TV}(\pi_t^*(\cdot \mid s_h), Q_t(\cdot \mid s_h))]$ is the occupancy-weighted per-round coordinate of (A5); $B_T = |\{t : a_t^* \neq a_{t-1}^*\}|$ counts optimum-change events; the partition is $0 = \tau_0 < \tau_1 < \dots < \tau_{B_T} \leq \tau_{B_T+1} = T$ with block i length $\Delta_i := \tau_{i+1} - \tau_i$.

A.1 Proof of (i) — Per-round identity-bounded regret

For deterministic $\pi_t^* = \delta_{a_t^*(s)}$ and any $Q_t(\cdot \mid s)$ with $Q_t(a_t^*(s) \mid s) > 0$ at every visited state (which is (A1)), Key Lemma 1 (Lemma 5.1, full proof in Appendix B.1) gives the exact identity $\text{TV}(\delta_{a_t^*(s)}, Q_t(\cdot \mid s)) = 1 - e^{-K_t(s)}$. Composed with the per-state TV-regret bounds (full proofs in Appendix B.1 and Appendix C.4): $\Delta_{\min} \cdot \text{TV}(\delta_{a_t^*(s)}, Q_t(\cdot \mid s)) \leq R(Q_t \mid s) \leq V_{\max} \cdot \text{TV}(\delta_{a_t^*(s)}, Q_t(\cdot \mid s))$, substituting the identity yields conclusion (i): $\Delta_{\min}(1 - e^{-K_t(s)}) \leq R(Q_t \mid s) \leq V_{\max}(1 - e^{-K_t(s)})$. $\square$

A.2 Proof of (ii) — Aggregate mismatch ultimate boundedness

Under (A2), Key Lemma 2 (Lemma 5.2, full proof in Appendix B.3) applied to the diagonal sector model with $\alpha_{ij}^{\text{ss}} = \nu_{ij} \cdot \iota_{ij} \cdot (1 - \lambda_{ij})$ gives ultimate boundedness of $\|\delta_\Sigma\|$ within $R_\Sigma^* = \rho_\Sigma / \mathcal{T}_\Sigma^{\text{bn,ss}}$. $\square$

749 A.3 Proof of (iii) — KL coordinate bias bound

750 Key Lemma 4 (Lemma 5.4, full proof in Appendix B.6) gives the per-state bias bound directly, under
 751 (A3)’s identifiability assumption (which gives $p_{\text{id}}(s)$ for each s) and the two-point support condition
 752 $Q_t(\cdot | s) \geq q_0$ at $a_{\text{ag},t}^*(s)$ and $\tilde{a}_t^*(s)$. $\square$

753 A.4 Proof of (iv) — 2×2 corrective-action routing

754 (A4) states the agent applies the satisfaction-gap / control-regret 2×2 (Section 4.1 Component 1).
 755 Conclusion (iv) is a direct restatement of the 2×2 cell partition. The decomposition is convention-
 756 dependent (the values of δ_{sat} and δ_{regret} depend on the continuation convention; see Appendix C.1 for
 757 monotonicity across C1/C2/C3); the 2×2 structure is preserved across all three conventions. $\square$

758 A.5 Proof of (v) — Trajectory-level cumulative regret

759 We chain Key Lemmas 1 and 4 with the simulation lemma and a Cauchy–Schwarz block argument.
 760 The proof is in four steps; the high-level sketch is in Section 5.5.

761 *Step 1: Simulation / performance-difference lemma.* For any two policies π
 762 and π' , any state s , and finite-horizon non-discounted MDP with horizon N_h ,
 763 the standard performance-difference lemma^[40–43] gives $V_t^{\pi'}(s) - V_t^\pi(s) =$
 764 $\mathbb{E}\left[\sum_{h=0}^{N_h-1} \left(\mathbb{E}_{a' \sim \pi'(\cdot | s_h)}[Q_h^{\pi'}(s_h, a')] - Q_h^\pi(s_h, a_h)\right) \mid s_0 = s, a_h \sim \pi\right]$. Each per-step bracket
 765 is bounded by value range times TV: $|\mathbb{E}_{a' \sim \pi'}[Q_h^{\pi'}] - \mathbb{E}_{a \sim \pi}[Q_h^\pi]| \leq V_{\max} \cdot \text{TV}(\pi'(\cdot | s_h), \pi(\cdot | s_h))$. Applied at round t with $\pi' = \pi_t^*$, $\pi = Q_t$: $V_t^{\pi_t^*}(s_0) - V_t^{Q_t}(s_0) \leq$
 766 $V_{\max} \sum_{h=0}^{N_h-1} \mathbb{E}_{s_h \sim d_h^{Q_t}(\cdot | s_0)}[\text{TV}(\pi_t^*(\cdot | s_h), Q_t(\cdot | s_h))] = V_{\max} \cdot N_h \cdot \overline{\text{TV}}_t$. The horizon factor N_h
 767 here is the *linear-in- N_h* simulation-lemma penalty — sharper than the N_h^2 of TRPO-style worst-state
 768 bounds^[44], which bound max-state TV rather than occupancy-weighted TV.

770 *Step 2: Per-state TV identity along the trajectory.* By the deterministic-per-state optimum scope of
 771 Theorem 4.1 ($\pi_t^*(\cdot | s) = \delta_{a_t^*(s)}$), at every visited state Key Lemma 1 gives $\text{TV}(\pi_t^*(\cdot | s), Q_t(\cdot | s)) = 1 - Q_t(a_t^*(s) | s) = 1 - e^{-K_t(s)}$. Summing over the trajectory and taking expectations:
 772 $\overline{\text{TV}}_t = \frac{1}{N_h} \sum_{h=0}^{N_h-1} \mathbb{E}_{s_h \sim d_h^{Q_t}}[1 - e^{-K_t(s_h)}]$. The per-state KL/TV identity is preserved underneath
 773 the occupancy aggregation; the per-round coordinate is sharper than Pinsker/BH at every visited state.

775 *Step 3: Block-sum + Cauchy–Schwarz across blocks.* Within each block $[\tau_i, \tau_{i+1})$ of length
 776 Δ_i , (A5)’s restart-on-change identity-tight base learner gives $\sum_{t \in \text{block}_i} \mathbb{E}[\overline{\text{TV}}_t] \leq 2c\sqrt{\Delta_i}$. Sum-
 777 ming over the $B_T + 1$ blocks: $\sum_{t=1}^T \mathbb{E}[\overline{\text{TV}}_t] = \sum_{i=0}^{B_T} \sum_{t \in \text{block}_i} \mathbb{E}[\overline{\text{TV}}_t] \leq 2c \sum_{i=0}^{B_T} \sqrt{\Delta_i}$.
 778 Cauchy–Schwarz gives $\sum_{i=0}^{B_T} \sqrt{\Delta_i} \leq \sqrt{(B_T + 1) \cdot \sum_{i=0}^{B_T} \Delta_i} = \sqrt{(B_T + 1)T}$. Multiplying by
 779 $V_{\max} N_h$: $\sum_{t=1}^T \mathbb{E}[V_{\max} N_h \overline{\text{TV}}_t] \leq 2c V_{\max} N_h \sqrt{(B_T + 1)T}$. This is the rate term.

780 *Step 4: Bias term via per-state, per-horizon-step lift of Key Lemma 4.* Key Lemma 4 (Lemma 5.4)
 781 applies pointwise at any visited state s_h along a horizon- N_h trajectory: $\mathbb{E}[|\hat{D}_t(s_h) - D_t^{\text{true}}(s_h)|] \leq$
 782 $(1 - p_{\text{id}}(s_h)) \log(1/q_0)$. Aggregating linearly over the N_h horizon steps and assuming the per-state
 783 argmax-correctness floor $p_{\text{id}} := \min_s p_{\text{id}}(s)$, the per-round bias contribution to the trajectory-level
 784 value gap is at most $N_h(1 - p_{\text{id}}) \log(1/q_0)$. Summing over T rounds: $N_h(1 - p_{\text{id}}) \log(1/q_0) \cdot T$.
 785 This is the bias term.

786 *Combining.* Adding Steps 1–4: $\mathbb{E}[\text{DynReg}(T)] \leq 2c V_{\max} N_h \sqrt{(B_T + 1)T} + N_h(1 -$
 787 $p_{\text{id}}) \log(1/q_0) \cdot T$, which is conclusion (v). $\square$

788 **Cesàro tracking corollary** Dividing through by T : $\frac{1}{T} \sum_t (V_t^{\pi_t^*} - V_t^{Q_t}) \leq$
 789 $2c V_{\max} N_h \sqrt{(B_T + 1)/T} + N_h(1 - p_{\text{id}}) \log(1/q_0)$. As $T \rightarrow \infty$ with $B_T = o(T)$ and $p_{\text{id}} \rightarrow 1$, the
 790 right side tends to zero — the agent’s average value over the trajectory tracks the moving optimum’s
 791 average value. $\square$

792 **Bandit case sharpening (Remark)** In the bandit special case ($N_h = 1$, so $\overline{\text{TV}}_t = \mathbb{E}[1 - e^{-K_t}]$),
 793 Thompson sampling and UCB (Appendix D) satisfy (A5) with the *stronger logarithmic* per-block rate

794 $\mathbb{E}[1 - e^{-K_t}] = O(\log t / (t \Delta_{\min}))$. Substituting and summing gives cumulative regret $O(V_{\max}(B_T +$
795 $1) \log^2(T/(B_T + 1))/\Delta_{\min})$ — sharper than $\sqrt{(B_T + 1)T}$ when $\Delta_{\min}$ is bounded away from zero.
796 The square-root rate is the worst-case Cauchy–Schwarz bound; the logarithmic rate is the typical-case
797 bound for stochastic-bandit base learners. For $N_h > 1$ MDPs, UCRL2 and UCBVI provide the
798 natural occupancy-weighted instantiation (Appendix D) with c of order $\tilde{O}(\sqrt{N_h SA})$, lifting cleanly
799 to $\tilde{O}(N_h^2 \sqrt{SA(B_T + 1)T})$ — same piecewise-stationary B_T family as restart-on-change variants
800 of^[1].

801 B Proofs of the Key Lemmas

802 This appendix gives full proofs of the four key lemmas surfaced in Section 5, plus the perturbative
803 extension and the ProST sector-level reduction.

804 B.1 Proof of Key Lemma 1 (Point-mass reverse-KL/TV identity)

805 For deterministic $\pi^* = \delta_{a^*}$ and any policy Q with $Q(a^*) > 0$:

806 The reverse-KL collapses under the point-mass: $D_{\text{KL}}(\delta_{a^*} \| Q) = \sum_a \delta_{a^*}(a) \log \frac{\delta_{a^*}(a)}{Q(a)} =$
807 $\log \frac{1}{Q(a^*)} = -\log Q(a^*)$, where the convention $0 \log 0 = 0$ disposes of the $a \neq a^*$ terms.
808 Combined with the textbook total-variation calculation $\text{TV}(\delta_{a^*}, Q) = \frac{1}{2} \sum_a |\delta_{a^*}(a) - Q(a)| =$
809 $\frac{1}{2} (|1 - Q(a^*)| + \sum_{a \neq a^*} Q(a)) = \frac{1}{2} ((1 - Q(a^*)) + (1 - Q(a^*))) = 1 - Q(a^*)$, we have
810 $-\log Q(a^*) = -\log(1 - \text{TV})$. Solving for TV gives $\text{TV} = 1 - e^{-D_{\text{KL}}}$. $\square$

811 **Strict-improvement substitution** Substituting the identity into the Bretagnolle–Huber inequality
812 $\text{TV}(P, Q) \leq \sqrt{1 - e^{-D_{\text{KL}}(P \| Q)}}$ at $P = \delta_{a^*}$ yields $1 - e^{-D_{\text{KL}}} \leq \sqrt{1 - e^{-D_{\text{KL}}}}$, which holds *strictly*
813 on $D_{\text{KL}} > 0$ since $x < \sqrt{x}$ on $(0, 1)$. The identity therefore supplies the *exact* TV value at the
814 deterministic- π^* corner and lies strictly below the BH envelope there.

815 **Two-sided regret bound** (stated as Component 2’s two-sided characterization in Section 4.1 and as Key
816 Lemma 1’s *Per-round regret bound* paragraph in Section 5.1). From the regret expression $R(Q) =$
817 $\sum_{a \neq a^*} Q(a) \Delta(a)$ and $\Delta(a) \leq V_{\max}$: $R(Q) \leq V_{\max} \sum_{a \neq a^*} Q(a) = V_{\max}(1 - Q(a^*)) =$
818 $V_{\max} \text{TV}(\delta_{a^*}, Q) = V_{\max}(1 - e^{-D_{\text{KL}}})$. The matching action-gap lower bound (full statement in
819 Appendix C.4) gives $R(Q) = \sum_{a \neq a^*} Q(a) \Delta(a) \geq \Delta_{\min} \sum_{a \neq a^*} Q(a) = \Delta_{\min} \text{TV}(\delta_{a^*}, Q) =$
820 $\Delta_{\min}(1 - e^{-D_{\text{KL}}})$. Combining gives the two-sided characterization. The bound is *Lipschitz-equivalent*
821 with constants $\Delta_{\min}/V_{\max}$ and 1, both achieved on extremal value landscapes — therefore *coordinate-*
822 *optimal among bounds depending only on TV*. $\square$

823 B.2 Proof of the perturbative extension — ϵ -stochastic and softmax-regularized

824 We establish the perturbative identity via a uniform perturbation argument rather than an alignment-
825 specific calculation.

826 **ϵ -greedy stochastic optimum** Let $\pi_{\epsilon}^*(a^*) = 1 - \epsilon$ and $\pi_{\epsilon}^*(a) = \epsilon/(|\mathcal{A}| - 1)$ for $a \neq a^*$, and
827 assume the full-support lower bound $Q(a) \geq q_0 > 0$ for all $a \in \mathcal{A}$. (Full support is required since π_{ϵ}^*
828 has positive mass on every action; without $Q(a) > 0$ on the full support, $D_{\text{KL}}(\pi_{\epsilon}^* \| Q) = +\infty$.)

829 *Step 1 — TV is Lipschitz under the point-mass perturbation.* $\text{TV}(\pi_{\epsilon}^*, \delta_{a^*}) = \epsilon$. By the triangle
830 inequality for total variation, $|\text{TV}(\pi_{\epsilon}^*, Q) - \text{TV}(\delta_{a^*}, Q)| \leq \text{TV}(\pi_{\epsilon}^*, \delta_{a^*}) = \epsilon$, uniformly in Q —
831 no alignment hypothesis needed.

832 *Step 2 — Reverse-KL admits a uniform expansion.* Direct computation:

$$833 \begin{aligned} D_{\text{KL}}(\pi_{\epsilon}^* \| Q) - D_{\text{KL}}(\delta_{a^*} \| Q) &= (1 - \epsilon) \log \frac{1 - \epsilon}{Q(a^*)} + \frac{\epsilon}{|\mathcal{A}| - 1} \sum_{a \neq a^*} \log \frac{\epsilon/(|\mathcal{A}| - 1)}{Q(a)} - (-\log Q(a^*)) \\ &= (1 - \epsilon) \log(1 - \epsilon) + \epsilon \log Q(a^*) + \epsilon \log \frac{\epsilon}{|\mathcal{A}| - 1} - \frac{\epsilon}{|\mathcal{A}| - 1} \sum_{a \neq a^*} \log Q(a), \end{aligned}$$

834 where the second equality collects the $\log Q(a^*)$ contributions via $-(1 - \epsilon) \log Q(a^*) + \log Q(a^*) =$

835 $\epsilon \log Q(a^*)$ and splits the constant $\log$ term off the conditional sum. Expanding
836 $(1 - \epsilon) \log(1 - \epsilon) = -\epsilon + O(\epsilon^2)$ and $\epsilon \log(\epsilon/(|\mathcal{A}| - 1)) = \epsilon \log \epsilon - \epsilon \log(|\mathcal{A}| - 1)$,
837 the leading-order behavior is $D_{\text{KL}}(\pi_\epsilon^* \| Q) = D_{\text{KL}}(\delta_{a^*} \| Q) + \epsilon \log \epsilon - \epsilon(\log(|\mathcal{A}| - 1) + 1) + \epsilon \log Q(a^*) - \frac{\epsilon}{|\mathcal{A}| - 1} \sum_{a \neq a^*} \log Q(a) + O(\epsilon^2)$. Under $Q(a) \geq q_0$, each
838 $|\log Q(a)| \leq \log(1/q_0)$ is uniformly bounded; the dominant correction is the $\epsilon \log \epsilon$ leading
839 term, giving $|D_{\text{KL}}(\pi_\epsilon^* \| Q) - D_{\text{KL}}(\delta_{a^*} \| Q)| = O(\epsilon \log(1/\epsilon) + \epsilon \log(1/q_0)) = O(\epsilon \log(1/\epsilon))$
840 absorbing $\epsilon \log(1/q_0)$ into the leading $\epsilon \log(1/\epsilon)$ term as $\epsilon \rightarrow 0$ (constants depend on q_0 and $|\mathcal{A}|$).

842 **Step 3 — Map through $1 - e^{-x}$ via Lipschitzness.** The map $x \mapsto 1 - e^{-x}$ is 1-Lipschitz on $[0, \infty)$,
843 so $|(1 - e^{-D_{\text{KL}}(\pi_\epsilon^* \| Q)}) - (1 - e^{-D_{\text{KL}}(\delta_{a^*} \| Q)})| = O(\epsilon \log(1/\epsilon))$. Combining Steps 1 and 3 with
844 the unperturbed identity $\text{TV}(\delta_{a^*}, Q) = 1 - e^{-D_{\text{KL}}(\delta_{a^*} \| Q)}$: $\text{TV}(\pi_\epsilon^*, Q) = 1 - e^{-D_{\text{KL}}(\pi_\epsilon^* \| Q)} +$
845 $O(\epsilon \log(1/\epsilon))$ uniformly in Q over the class $\{Q : Q(a) \geq q_0\}$, with constants depending on q_0 and
846 $|\mathcal{A}|$. This establishes the ϵ -greedy form of the perturbative identity. $\square$

847 **Softmax-regularized optimum** Let $\pi_\tau^*(a) \propto \exp(Q_O(a)/\tau)$. Under isolated optimum with
848 action gap $\Delta_{\min}$, the off-optimum mass concentrates exponentially in $1/\tau$: $1 - \pi_\tau^*(a^*) =$
849 $\frac{\sum_{a \neq a^*} \exp(-(Q_O(a^*) - Q_O(a))/\tau)}{1 + \sum_{a \neq a^*} \exp(-(Q_O(a^*) - Q_O(a))/\tau)} \leq (|\mathcal{A}| - 1) \exp(-\Delta_{\min}/\tau)$. Treating π_τ^* as π_ϵ^* with ef-
850 fective $\epsilon = O(\exp(-\Delta_{\min}/\tau))$, the ϵ -greedy expansion above applies with $\epsilon \log(1/\epsilon) =$
851 $O((\Delta_{\min}/\tau) \exp(-\Delta_{\min}/\tau)) = O(\tau^{-1} \exp(-\Delta_{\min}/\tau))$, exponentially small with a polynomial
852 prefactor (equivalently, $O(\exp(-c\Delta_{\min}/\tau))$ for any fixed $c \in (0, 1)$ since the exponential dominates
853 the polynomial). $\square$

854 **Two-sided regret bound under perturbation** Composing with the TV-regret bounds: $\Delta_{\min}(1 -$
855 $e^{-D_{\text{KL}}(\pi_\epsilon^* \| Q)}) - O(\epsilon \log(1/\epsilon)) \leq R(Q) \leq V_{\max}(1 - e^{-D_{\text{KL}}(\pi_\epsilon^* \| Q)}) + O(\epsilon \log(1/\epsilon))$.

856 B.3 Proof of Key Lemma 2 (Multi-factor forgetting prerequisite)

857 Within Model (S) of Section 5.2 with diagonal correction architecture, the Lyapunov function
858 $V(\delta_\Sigma) := \|\delta_\Sigma\|^2$ satisfies $\mathbb{E}[\Delta V \mid \delta_\Sigma] = -2\delta_\Sigma^\top \mathbf{F}(\delta_\Sigma) + \mathbb{E}[\|\mathbf{w}\|^2 \mid \delta_\Sigma]$, where $\mathbf{F}(\delta_\Sigma) :=$
859 $(\alpha_{ij}^{\text{ss}} \delta_{ij})_{(i,j) \in E}$ is the sector correction. The first term is the *sector product*: $\delta_\Sigma^\top \mathbf{F}(\delta_\Sigma) =$
860 $\sum_{(i,j)} \alpha_{ij}^{\text{ss}} \delta_{ij}^2 \geq \min_{(i,j)} \alpha_{ij}^{\text{ss}} \cdot \|\delta_\Sigma\|^2 = \mathcal{T}_\Sigma^{\text{bn,ss}} \cdot \|\delta_\Sigma\|^2$. The bound is *tight*: setting δ_Σ pro-
861 portional to the indicator $\mathbf{e}_{(i^*, j^*)}$ where $(i^*, j^*) \in \arg \min_{(i,j)} \alpha_{ij}^{\text{ss}}$ saturates the inequality. The
862 disturbance second-moment $\mathbb{E}[\|\mathbf{w}\|^2] \leq \rho_\Sigma^2$.

863 Plugging in: $\mathbb{E}[\Delta V \mid \delta_\Sigma] \leq -2\mathcal{T}_\Sigma^{\text{bn,ss}} V + \rho_\Sigma^2$. This is the standard Khalil sector-Lyapunov ultimate-
864 boundedness setup^[28] (Lemma 9.2 / Theorem 9.1 of Chapter 9). Whenever $V > \rho_\Sigma^2 / (2\mathcal{T}_\Sigma^{\text{bn,ss}})$,
865 the expected drift is negative; iterating gives ultimate boundedness with $V \leq R_\Sigma^{*2}$ where $R_\Sigma^{*2} =$
866 $\rho_\Sigma^2 / \mathcal{T}_\Sigma^{\text{bn,ss} \cdot 2}$ to leading order, i.e., $R_\Sigma^* = \rho_\Sigma / \mathcal{T}_\Sigma^{\text{bn,ss}}$. So whenever $\mathcal{T}_\Sigma^{\text{bn,ss}} > \rho_\Sigma / R_\Sigma$, we have $R_\Sigma^* \leq R_\Sigma$
867 and the modeled mismatch is ultimately bounded within the strategic reserve. $\square$

868 **Sharpness** When the inequality reverses ($\mathcal{T}_\Sigma^{\text{bn,ss}} \leq \rho_\Sigma / R_\Sigma$), an adversarial disturbance concentrat-
869 ing on the bottleneck element (i^*, j^*) with $|w_{i^* j^*}| = \rho_\Sigma$ pushes $\|\delta_\Sigma\| > R_\Sigma$ within the modeled
870 dynamics — the threshold is sharp inside the diagonal sector model. The theorem is silent about
871 non-diagonal correction architectures or stabilization mechanisms outside Model (S). $\square$

872 B.4 Proof of the impulsive ProST sector-level reduction

873 Within Model (S), idealize ProST^[8] as an impulsive system: between scheduled update times
874 $\{t_1, \dots, t_K\} \subset [0, T]$ the agent's policy is held fixed and the strategic mismatch evolves under
875 disturbance budget ρ_Σ alone (continuous-destabilizing in the Lyapunov $V = \|\delta_\Sigma\|^2$); at each update
876 time t_i the policy is revised, contracting the modeled mismatch by per-update impulse gain $\gamma \in (0, 1]$
877 so $\|\delta_\Sigma(t_i^+) \| \leq (1 - \gamma) \|\delta_\Sigma(t_i^-) \|$. Then $V(t_i^+) \leq (1 - \gamma)^2 V(t_i^-)$.

878 This fits the Hespanha–Liberzon–Teel^[46] (Theorem 1) framework for impulsive systems with desta-
879 bilizing continuous evolution and stabilizing impulses (the *reverse-ADT* branch). The continuous-
880 evolution rate is $\lambda_c = \rho_\Sigma / R_\Sigma$ (from linearizing the disturbance contribution to $\dot{V} \approx +\lambda_c V + \gamma_c(\rho_\Sigma)$).

881 The impulse contraction in V is $\mu = (1-\gamma)^2$. The reverse-ADT condition for ultimate boundedness —
 882 impulses cannot be too sparse — gives $\Delta_{\max} \cdot \lambda_c < -\ln \mu \iff \Delta_{\max} \cdot \rho_{\Sigma}/R_{\Sigma} < -\ln(1-\gamma)^2$,
 883 where $\Delta_{\max} := \sup_i (t_{i+1} - t_i)$ is the longest inter-update gap. Under the condition the modeled
 884 mismatch is ultimately bounded within $R_{\Sigma}^* \approx \rho_{\Sigma} \Delta_{\max}/\gamma$ (leading order in γ).

885 Under uniform blocks $\Delta_i = T/K$ the condition reduces to $K/T > \rho_{\Sigma}/(2\gamma R_{\Sigma})$ to leading order in γ
 886 — recovering ProST’s K/T form with the impulse gain made explicit. Under nonuniform schedules
 887 the impulsive form is *strictly stronger*: the longest block sets the threshold, not the average. The
 888 lemma is rigorous for the modeled impulsive sector dynamics; the per-update impulse gain γ is a
 889 domain quantity (depends on within-block sample size and update-time discount) and is left abstract.
 890 $\square$

891 B.5 Proof of Key Lemma 3 (Loop generates interventional samples)

892 By temporal ordering and the singular-trajectory commitment of Section 3, a_t causally precedes o_{t+1} .
 893 Under (C1) positivity, (C2) sequential ignorability, and (C3) known action mechanism, conditioning
 894 on the information state H_t blocks the action-selection confounding pathway. Specifically, (C2) gives
 895 $a_t \perp\!\!\!\perp (\text{environment latents}) \mid H_t$, so for any subset of environment-latent variables $\mathbf{U}$, $P(o_{t+1} \mid$
 896 $a_t, H_t, \mathbf{U}) = P(o_{t+1} \mid a_t, H_t)$, on the support where $\pi_t(a_t \mid H_t) > 0$ (which (C1) guarantees). The
 897 interventional distribution on the same support is, by Pearl’s do-calculus^[35], $P(o_{t+1} \mid \text{do}(a_t), H_t) =$
 898 $\sum_{\mathbf{U}} P(o_{t+1} \mid a_t, H_t, \mathbf{U}) \cdot P(\mathbf{U} \mid H_t) = \sum_{\mathbf{U}} P(o_{t+1} \mid a_t, H_t) \cdot P(\mathbf{U} \mid H_t) = P(o_{t+1} \mid$
 899 $a_t, H_t)$, where the second equality uses the conditional-independence factorization above. So the
 900 conditional and interventional distributions coincide on the support. (C3) makes the agent’s action-
 901 mechanism identification well-defined; the architecture of the agent does not enter — the loop
 902 generates conditional-Level-2 samples whenever (C1)–(C3) hold. $\square$

903 B.6 Proof of Key Lemma 4 (Bias bound)

904 On the matching event $\{a_{\text{ag},t}^*(s) = \tilde{a}_t^*(s)\}$, the difference $\hat{D}_t(s) - D_t^{\text{true}}(s)$ vanishes identically (both
 905 quantities equal $-\log Q_t$ at the same action). On the complementary misidentification event, the
 906 support condition $Q_t(\cdot \mid s) \geq q_0$ at *both* $a_{\text{ag},t}^*(s)$ and $\tilde{a}_t^*(s)$ implies $\log Q_t(a_{\text{ag},t}^*(s) \mid s) \in [\log q_0, 0]$
 907 and $\log Q_t(\tilde{a}_t^*(s) \mid s) \in [\log q_0, 0]$, so their absolute difference is at most $\log(1/q_0)$. Therefore
 908 $|\hat{D}_t(s) - D_t^{\text{true}}(s)| \leq \mathbf{1}[a_{\text{ag},t}^*(s) \neq \tilde{a}_t^*(s)] \cdot \log(1/q_0)$. Taking expectations: $\mathbb{E}[|\hat{D}_t(s) - D_t^{\text{true}}(s)|] \leq$
 909 $(1 - p_{\text{id}}(s)) \log(1/q_0)$ by definition of $p_{\text{id}}(s)$. $\square$

910 **High-probability form** With probability $1 - \delta$ over a single round’s identification event, $|\hat{D}_t(s) -$
 911 $D_t^{\text{true}}(s)| \leq \log(1/q_0) \cdot \mathbf{1}[\text{misdevent}]$ where $\Pr[\text{misdevent}] = 1 - p_{\text{id}}(s)$. The bias decomposes
 912 cleanly: in Regime A ($p_{\text{id}}(s) \rightarrow 1$) it vanishes; in Regime B ($p_{\text{id}}(s) \in (0, 1)$) it is controlled by
 913 misidentification mass; in Regime C ($p_{\text{id}}(s) \rightarrow 0$) it saturates at $\log(1/q_0)$. This proves Key Lemma
 914 4. $\square$

915 **Per-state, per-horizon-step lift (used in Theorem 4.1(v)).** The lemma applies pointwise at any
 916 visited state s_h along a horizon- N_h trajectory: $\mathbb{E}[|\hat{D}_t(s_h) - D_t^{\text{true}}(s_h)|] \leq (1 - p_{\text{id}}(s_h)) \log(1/q_0)$,
 917 with $p_{\text{id}}(s_h)$ the per-state argmax-correctness probability. Aggregating linearly over the horizon and
 918 assuming the per-state floor $p_{\text{id}} := \min_s p_{\text{id}}(s)$, the per-round bias contribution to the trajectory-level
 919 value gap is at most $N_h(1 - p_{\text{id}}) \log(1/q_0)$, contributing the second term of Theorem 4.1(v).

920 **Remarks** - Both $a_{\text{ag},t}^*(s)$ and $\tilde{a}_t^*(s)$ must satisfy $Q_t \geq q_0$ for the $\log(1/q_0)$ ceiling to apply; the
 921 support condition is on *both* points, not just the agent’s estimate. This is *strictly weaker* than the
 922 perturbative identity’s full-support condition $Q(a) \geq q_0$ for *all* a (Appendix B.2) — the perturbative
 923 identity needs full support because π_{ϵ}^* has positive mass on every action; the bias bound here needs
 924 only two-point support. - This is *not* an estimation-from-samples lemma. If Q_t is unavailable internally
 925 and must be estimated from on-policy rollouts, an additional concentration argument (Hoeffding
 926 under iid fixed-policy samples, or martingale concentration for adaptive policies) is needed; in the
 927 architectures this paper considers, Q_t is the agent’s policy and is internally available.

928 C Auxiliary Mathematical Material

929 This appendix collects supporting material used by the main results: convention hierarchy details, the
 930 admissible-divergence family, direction-forcing argument, action-gap matching lower bound, tied-
 931 optimum extensions, the structural-class theorem on gain-decay updates, strategic-tempo topologies,
 932 deployment modes, and the chain-rule uniqueness theorem.

933 C.1 Convention hierarchy: C1, C2, C3

934 Three named continuation conventions form a hierarchy of increasing diagnostic power and computa-
 935 tional cost.

936 **C1 — One-step improvement (canonical default)** $\pi_{\text{cont}} = \pi_{\text{current}}$. Each action is evaluated
 937 assuming current behavior continues afterward. No fixed-point computation, no global-optimality
 938 assumption. Cheapest; weakest diagnostic.

939 **C2 — Receding-horizon (N_r -step replanning)** At each future step, re-optimize
 940 over a horizon of N_r steps using the model available at that step: $\pi_{\text{RH}}(M_\tau) =$
 941 $\arg \max_\pi V_O(M_\tau, \pi; N_r)$ applied at each τ . Captures multi-step recovery: a goal that ap-
 942 pears unattainable under frozen continuation may be reachable with replanning. Moderate cost;
 943 moderate diagnostic.

944 **C3 — Bellman** $\pi_{\text{cont}} = \pi^*$ where $\pi^* = \arg \max_\pi V_O(M_t, \pi; N_h)$. The continuation is the optimal
 945 policy — a fixed-point equation. Strongest diagnostic; most expensive.

946 **Proposition C.1** (Monotonicity). *For any model M_t , horizon N_h , and policy class Π :*
 947 $A_O^{(1)}(M_t; \Pi, N_h) \leq A_O^{\text{RH}}(M_t; \Pi, N_r, N_h) \leq A_O^{\text{B}}(M_t; \Pi, N_h)$.

948 *Proof.* C1 freezes continuation at π_{current} (possibly suboptimal); C2 re-optimizes periodically, so
 949 $\pi_{\text{RH}} \succeq \pi_{\text{current}}$ at each future step; C3 uses the globally optimal π^* . A weakly better continuation
 950 yields a weakly higher expected trajectory value; taking the supremum over the first action preserves
 951 the ordering. $\square$

952 **Corollary** $\delta_{\text{sat}}^{\text{B}} \leq \delta_{\text{sat}}^{\text{RH}} \leq \delta_{\text{sat}}^{(1)}$ (since $\delta_{\text{sat}} = V_O^{\min} - A_O$, higher A_O means lower δ_{sat}). And
 953 $\delta_{\text{regret}}^{(1)} \leq \delta_{\text{regret}}^{\text{RH}} \leq \delta_{\text{regret}}^{\text{B}}$ (since $\delta_{\text{regret}} = A_O - V_O(M_t, \pi_{\text{current}}; N_h)$, higher A_O means higher δ_{regret}).

954 C1 is the most conservative diagnostic (most likely to diagnose “locally unattainable”); C3 is the most
 955 accurate (least false “unattainable” diagnoses). The 2×2 diagnostic structure is preserved under all
 956 three; only the inferential force varies.

957 C.2 Admissible-divergence family for the regret bound

958 The reverse-KL direction is forced by the deterministic- π^* regime (Appendix C.3). Within the
 959 direction-forced family, multiple smooth f -divergences yield valid regret bounds:

	Divergence	Bound on $R(Q)$	Tightness	Finite under det. π^* ?
	$\text{TV}(\pi^*, Q)$	$V_{\max} \cdot \text{TV}$	Tight (extremal V)	Yes
	$D_{\text{KL}}(\pi^* \ Q)$ via Pinsker	$V_{\max} \sqrt{D_{\text{KL}}/2}$	Loose by $\sqrt{\cdot}$	Yes
960	$D_{\text{KL}}(\pi^* \ Q)$ via point-mass identity	$V_{\max}(1 - e^{-D_{\text{KL}}})$	Tight (this paper)	Yes
	$\chi^2(\pi^* \ Q)$ (Le Cam)	$V_{\max} \cdot \frac{1}{2} \sqrt{\chi^2}$	Looser than Pinsker	$\chi^2 = 1/Q(a^*) - 1$
	$D_\alpha(\pi^* \ Q)$ (Rényi, $\alpha \geq 1$)	Various	Interpolates	Yes for $\alpha \geq 1$
	$D_{\text{KL}}(Q \ \pi^*)$ (forward-KL)	$+\infty$	Vacuous	No

961 Reverse-KL is selected uniquely within the direction-forced family by the chain-rule axiom (Ap-
 962 pendix C.10). The point-mass identity supplies the *exact* bound on reverse-KL under determinis-
 963 tic π^* ; the table reflects different bound shapes on the same divergence (with the BH inequality
 964 $V_{\max} \sqrt{1 - e^{-D_{\text{KL}}}}$ as the looser general form to which the identity reduces outside the deterministic- π^*
 965 corner).

966 C.3 Direction-forcing argument

967 For deterministic $\pi^* = \delta_{a^*}$ and any Q with $Q(a) > 0$ for some $a \neq a^*$: $D_{\text{KL}}(Q \| \pi^*) =$
 968 $\sum_a Q(a) \log \frac{Q(a)}{\pi^*(a)} = \sum_{a \neq a^*} Q(a) \log \frac{Q(a)}{0} = +\infty$. A bound “ $R \leq +\infty$ ” is vacuous. The
 969 reverse direction $D_{\text{KL}}(\pi^* \| Q)$ is finite (and equal to $-\log Q(a^*)$) whenever $Q(a^*) > 0$. The asym-
 970 metry is forced by the regime: regret is a one-sided quantity (contributes only from Q ’s off-optimum
 971 mass; π^* has no support off a^*); divergences whose role is to bound this one-sided quantity must
 972 themselves be one-sided. Symmetric divergences (squared Hellinger, Jensen-Shannon, symmetrized
 973 KL) introduce a vacuous symmetric term.

974 C.4 Action-gap matching lower bound

975 For any Q with $\Delta_{\min} = \min_{a \neq a^*} \Delta(a) > 0$: $R(Q) = \sum_{a \neq a^*} Q(a) \Delta(a) \geq \Delta_{\min} \sum_{a \neq a^*} Q(a) =$
 976 $\Delta_{\min} \cdot (1 - Q(a^*)) = \Delta_{\min} \cdot \text{TV}(\pi^*, Q)$. By the point-mass identity (Lemma 5.1), $\text{TV}(\pi^*, Q) =$
 977 $1 - e^{-D_{\text{KL}}(\pi^* \| Q)}$, giving the matching lower bound used in the two-sided characterization.

978 The lower bound is tight when sub-optimal actions are uniformly bad ($\Delta_{\min} = \max_{a \neq a^*} \Delta(a)$). For
 979 typical landscapes the gap between upper and lower bound is $V_{\max} - \Delta_{\min}$, controlled by the *spread*
 980 of action gaps.

981 C.5 Tied-optimum and softmax-smoothed extensions

982 **Tied-optimum** If π^* has support on a tied-optimum set $\mathcal{A}^* = \{a : Q_O(a) = Q_O(a^*)\}$
 983 with uniform mass, reverse-KL is finite whenever Q covers $\mathcal{A}^*$. The regret bound extends:
 984 $R(Q) = \sum_{a \notin \mathcal{A}^*} Q(a) \Delta(a) \leq V_{\max} \cdot \mathbb{P}_Q(a \notin \mathcal{A}^*)$. A multi-action analog of the point-mass
 985 identity holds with $\pi^*(a^*) = 1/|\mathcal{A}^*|$ replacing the point-mass form, yielding $D_{\text{KL}}(\pi^* \| Q) =$
 986 $-\sum_{a \in \mathcal{A}^*} (1/|\mathcal{A}^*|) \log(|\mathcal{A}^*| Q(a))$, a multi-action analog of $-\log Q(a^*)$. The two-sided bound be-
 987 comes one-sided in general; the upper bound holds with the modified KL form.

988 **Softmax-smoothed π^* with temperature $\tau \rightarrow 0$** Replace $\pi^* = \delta_{a^*}$ with $\pi_\tau^* \propto \exp(Q_O/\tau)$. As
 989 $\tau \rightarrow 0$, $\pi_\tau^* \rightarrow \delta_{a^*}$. For finite τ , the *perturbative* identity (Appendix B.2) applies: $\text{TV}(\pi_\tau^*, Q) =$
 990 $1 - e^{-D_{\text{KL}}(\pi_\tau^* \| Q)} + O(\tau^{-1} \exp(-\Delta_{\min}/\tau))$ — exponentially small with polynomial prefactor. The
 991 two-sided regret bound transfers with the same correction order. Outside the perturbative regime —
 992 for genuinely high-entropy optima — the BH inequality $\text{TV} \leq \sqrt{1 - e^{-D_{\text{KL}}}}$ becomes the relevant
 993 general bound; Pinsker also applies, with a one-sided regret form.

994 C.6 The gain-decay structural class $\mathcal{A}_{\text{decay}}$

995 The persistence inequality $\mathcal{T}_{\Sigma}^{\text{bn,ss}} > \rho_{\Sigma}/R_{\Sigma}$ (Section 5.2) is an *instantaneous* check at the current
 996 operating point. Define the structural class by gain-decay rather than sample-counting: $\mathcal{A}_{\text{decay}} :=$
 997 $\{\text{updates whose effective sector gain } \alpha_{ij}^{(t)} \rightarrow 0 \text{ as } t \rightarrow \infty \text{ on every revisable element}\}$.

998 **Theorem C.2** (Universal failure of $\mathcal{A}_{\text{decay}}$). *For any fixed $(\rho_{\Sigma}, R_{\Sigma})$ with $\rho_{\Sigma} > 0$, every agent in*
 999 *$\mathcal{A}_{\text{decay}}$ eventually violates the persistence threshold $\mathcal{T}_{\Sigma}^{\text{bn,ss}} > \rho_{\Sigma}/R_{\Sigma}$.*

1000 *Proof.* At every element (i, j) , $\alpha_{ij}^{(t)} \rightarrow 0$ by definition of $\mathcal{A}_{\text{decay}}$, so $\min_{(i,j)} \alpha_{ij}^{(t)} \rightarrow 0$, so the
 1001 bottleneck $\mathcal{T}_{\Sigma}^{\text{bn}} \rightarrow 0$ with experience. Once $\mathcal{T}_{\Sigma}^{\text{bn,ss}} < \rho_{\Sigma}/R_{\Sigma}$, the threshold is violated. This holds
 1002 *regardless of prior calibration*: at any finite calibration level, the gain decays below threshold once
 1003 enough experience accumulates. $\square$

1004 This is a *structural* failure of the class, not a tuning problem. The class includes: - Count-accumulating
 1005 Bayesian updates without forgetting (e.g., Beta-Bernoulli with $\eta_{\text{edge},ij} = 1/(n_{ij} + 1)$ where $n_{ij} \rightarrow \infty$).
 1006 - Bounded-memory schemes with growing memory. - Observation-aggregating schemes without restart.
 1007 - Gradient-based methods with vanishing step sizes ($\eta_t = c/t^{\beta}$ for $\beta \in (0, 1]$, the Robbins–Monro
 1008 regime).

1009 Mechanisms *outside* $\mathcal{A}_{\text{decay}}$ — constant-step-size stochastic approximation, sliding-window updates,
 1010 bounded-memory learners, block-restart schemes, Kalman filters with bounded process noise —
 1011 maintain a finite gain ceiling and escape asymptotic decay, but then face a *bidirectional* threshold:

Table 1: Bidirectional thresholds for non-accumulating mechanisms outside $\mathcal{A}_{\text{decay}}$.

Mechanism	$\alpha_{\Sigma}^{\text{ss}}$	Prerequisite holds iff
Exponential forgetting with λ	$1 - \lambda$	$1 - \lambda > \rho_{\Sigma}/R_{\Sigma}$
Constant-step SA with rate η	η	$\eta > \rho_{\Sigma}/R_{\Sigma}$
Fixed-window mechanisms (sliding-window size W , bounded-memory size K , block-restart period T_R)	$1/W, 1/K, 1/T_R$	window/memory/period $< R_{\Sigma}/\rho_{\Sigma}$

1012 For each row, the threshold is *bidirectional* and sharp within the sector-Lyapunov reduction.

1013 C.7 Strategic-tempo across canonical topologies

1014 The bottleneck strategic tempo $\mathcal{T}_{\Sigma}^{\text{bn,ss}}$ of Lemma 5.2 is verified — and the resulting per-topology
 1015 threshold derived — across four canonical topologies under Beta-Bernoulli edge updates with per-
 1016 element forgetting:

- 1017 • **(T1) Single edge.** $\mathcal{T}_{\Sigma}^{\text{bn,ss}} = \nu \cdot \iota \cdot (1 - \lambda)$. All three factors load-bearing; setting any one to
 1018 zero collapses the bottleneck.
- 1019 • **(T2) Two-edge AND chain, observable intermediate.** $\mathcal{T}_{\Sigma}^{\text{bn,ss}} = \min(\nu_1(1 - \lambda_1), \nu_1\theta_1(1 -$
 1020 $\lambda_2))$. Edge 2’s effective observation rate is gated by edge 1’s success probability θ_1 —
 1021 *depth-gated attenuation*. For depth- d chains, the bottleneck is $\min_k \prod_{j < k} \theta_j(1 - \lambda_k)$.
- 1022 • **(T3) Two-edge AND chain, unobservable intermediate.** Per-edge tempo ill-defined; plan-
 1023 level bottleneck $\mathcal{T}_{\Sigma, \text{plan}}^{\text{bn,ss}} = \nu(1 - \lambda_{\Phi})$ over a single tracked plan-level quantity $\hat{\Phi} = p_1 p_2$.
- 1024 • **(T4) Two-arm OR node, ε -greedy.** $\mathcal{T}_{\Sigma}^{\text{bn,ss}} = \min((1 - \varepsilon)(1 - \lambda_1), \varepsilon(1 - \lambda_2))$. Action
 1025 selection controls rate allocation; pure greedy ($\varepsilon = 0$) collapses the unexplored arm’s
 1026 bottleneck — *exploration-gated*.

1027 The structural decomposition: AND-chains exhibit *depth-gated* geometric attenuation $\nu_k =$
 1028 $\nu \prod_{j < k} \theta_j$; OR-nodes exhibit *exploration-gated* allocation. Mixed AND/OR DAGs interleave both.
 1029 Each topology surfaces a different factor as the binding bottleneck.

1030 **Regime-C contamination is structurally fatal** $\iota_{ij} = 0$ at any single element $\Rightarrow \mathcal{T}_{\Sigma}^{\text{bn,ss}} = 0$. An
 1031 agent with even one fully-confounded element cannot meet the threshold by forgetting alone.

1032 C.8 Three deployment modes of loop-Level-2

1033 The Pearl-do structure of closed-loop interventional access manifests at distinct layers with semanti-
 1034 cally distinct interventional mechanisms:

- 1035 • **Mode 1 — agent-self-intervention.** The agent performs do-actions on its own action space
 1036 as part of its ordinary adaptive loop. Intervention is on the agent’s own action; target is the
 1037 environment’s response. This is Lemma 5.3’s primary content.
- 1038 • **Mode 2 — observer-on-sub-agent.** An observer external to a composite performs do-
 1039 interventions on one sub-agent’s action space; target is another sub-agent’s response. Reveals
 1040 cross-coupling structure that component-marginal observation cannot identify.
- 1041 • **Mode 3 — observer-on-agent-input** (sketched, future work). An observer intervenes on
 1042 the agent’s observation channel; target is the agent’s subsequent policy. Useful for offline
 1043 architectural-class diagnosis.

1044 Modes share the Pearl-do structure but differ in who intervenes on what. The unification is at the
 1045 pattern level; the mechanism is semantically distinct per layer.

C.9 Distinction from active inference and causal-RL precursors

Action-perception-loop frameworks — active inference^[30,31], control-as-inference^[48], cybernetics^[49,50] — implicitly use the action-causes-observation observation. Our distinctive moves:

- *Bareinboim-hierarchy connection*. Active inference / cybernetics rest on Bayesian-network (Level 1) generative models; we invoke the causal-hierarchy theorem^[29] to position the policy DAG as causal with do-conditioning in Q_O .
- *Regime-indexed identifiability (A/B/C)*. Active inference literature treats causal identifiability uniformly within modeling assumptions; we surface the regime split at framework level.
- *Scope honesty*. We distinguish “data generated under intervention” from “cleanly identified do-estimates”;^[47] documents that the active-inference literature sometimes elides this.

The causal-RL line^[13–18] is the direct ancestor for regime-indexed identifiability and on-policy interventional access; all are stationary-MDP. Composition with non-stationarity is, to our knowledge, novel.

C.10 Chain-rule uniqueness of reverse-KL

Theorem C.3 (Chain-rule uniqueness (Hobson 1969; Csiszár 1991)). *Let $D_f(P \parallel Q) = \sum_x Q(x) f(P(x)/Q(x))$ be a smooth f -divergence with f convex and $f(1) = 0$. The chain rule $D_f(P_{XY} \parallel Q_{XY}) = D_f(P_X \parallel Q_X) + \mathbb{E}_{P_X}[D_f(P_{Y|X} \parallel Q_{Y|X})]$ holds for all joint distributions if and only if $f(t) = c \cdot t \log t$ for some $c > 0$ — i.e., D_f is reverse-KL up to positive scaling.*

Proof. Writing $r_x = P(x)/Q(x)$ and $s_{y|x} = P(y|x)/Q(y|x)$, the chain rule reduces to the functional equation $f(rs) = f(r) + rf(s)$ for all $r, s > 0$. With $f(1) = 0$ and convexity, the unique solution is $f(t) = c \cdot t \log t$ for $c > 0$ ^[51] (§4). $\square$

References Hobson 1969 (“A new theorem of information theory,” *J. Stat. Phys.*); Csiszár 1991 (“Why least squares and maximum entropy?” *Annals of Statistics*; Theorem 3 corollary, Theorem 5); Shore-Johnson 1980 (“Axiomatic derivation of the principle of maximum entropy,” *IEEE Trans. Info. Theory*, system-independence axiom); Sanov 1957 (large-deviation rate function); Aczél-Daróczy 1975 (functional-equation machinery).

These references give *structurally equivalent reformulations* of the same axiom. The Cauchy functional equation each reduces to is the common content. No known uniqueness route outside the independence-on-sub-problems family exists.

Why other family members fail the chain rule Concrete counterexample for χ^2 : take Q_X uniform on $\{x_1, x_2\}$, $P_X = (3/4, 1/4)$, $Q(y|x)$ uniform, $P(y|x) = (3/4, 1/4)$. Direct calculation gives $\chi^2(P_{XY} \parallel Q_{XY}) = 9/16$ while $\chi^2(P_X \parallel Q_X) + \mathbb{E}_{P_X}[\chi^2(P_{Y|X} \parallel Q_{Y|X})] = 1/4 + 1/4 = 8/16$. Non-additive. Rényi- α for $\alpha \neq 1$ fails analogously; squared Hellinger likewise fails.

D Algorithm Sketch and Base-Learner Instantiation

A practical algorithm achieving the joint guarantees of Theorem 4.1 requires four components, each translating one of the assumptions (A1)–(A5) into operational form.

(a) Identity-tight base learner satisfying (A5) In the stationary deterministic- π^* bandit case ($N_h = 1$), Thompson sampling^[19] and UCB^[27] both achieve the per-round coordinate $\mathbb{E}[1 - e^{-K_t}] = \mathbb{E}[1 - Q_t(a^*)] = O\left(\frac{\log t}{t \Delta_{\min}}\right)$, matching the identity form $V_{\max}(1 - e^{-D_{\text{KL}}})$ of Lemma 5.1 up to constants and a $\log t$ factor. This satisfies (A5) of Theorem 4.1 with the *stronger logarithmic* rate, yielding cumulative dynamic regret $O(V_{\max}(B_T + 1) \log^2(T/(B_T + 1))/\Delta_{\min})$. For $N_h > 1$ MDPs, UCRL2^[52] and UCBVI^[43] give per-block cumulative trajectory-level regret $\tilde{O}(N_h^{3/2} \sqrt{SAL})$ with S, A the state/action sizes; rewriting in the occupancy-weighted coordinate satisfies (A5) with c of order $\tilde{O}(\sqrt{N_h SA})$, lifting cleanly to the $\tilde{O}(N_h^2 \sqrt{SA(B_T + 1)T})$ cumulative rate — same

1090 piecewise-stationary B_T family as restart-on-change variants of^[1] (the continuous-variation regime
1091 of^[3] is structurally distinct; see Section 4 *Comparator regime*). The occupancy-weighted form is the
1092 natural object for trajectory-level base learners; the per-state KL/TV identity is preserved underneath.
1093 Information-directed sampling^[20] requires a different analysis here because $H(\pi^*) = 0$ collapses its
1094 information ratio at the deterministic- π^* corner; we leave the IDS analysis as future work.

1095 **(b) Multi-factor forgetting schedule satisfying (A2)** Choose per-element discount rates $\{\lambda_{ij}\}$
1096 such that the bottleneck $\mathcal{T}_{\Sigma}^{\text{bn,ss}} = \min_{(i,j)} \nu_{ij} \iota_{ij} (1 - \lambda_{ij})$ exceeds ρ_{Σ}/R_{Σ} . When ρ_{Σ} is unknown,
1097 the^[2] black-box reduction gives an adaptive choice (applied per element if heterogeneous); when ρ_{Σ}
1098 is known via a variation-budget^[1] or change-point detection^[39], the choice is direct. As a fail-fast
1099 pre-check, $\mathcal{T}_{\Sigma}^{\text{agg,ss}} > |E| \cdot \rho_{\Sigma}/R_{\Sigma}$ is necessary.

1100 **(c) Identifiability check satisfying (A3)** Test whether the loop data is in Regime A (full identifiability),
1101 Regime B (partial), or Regime C (none). In Regime A, no adjustment is needed. In Regime
1102 B, apply^[17] (DOVI)-style confounding-bias adjustment. In Regime C, the bound is provable but not
1103 realizable; algorithm flags the regime as out-of-scope.

1104 **(d) Two-gap diagnostic readout satisfying (A4)** Compute δ_{sat} (requires A_O , intractable in general;
1105 estimable via the policy-class supremum over recent rounds) and δ_{regret} (requires V_O at current policy;
1106 tractable in simulation). Route to corrective action class.

1107 The algorithm is sketch-grade; full instantiation and empirical evaluation are deferred to a follow-up
1108 paper.

1109 E Pinsker vs. Point-Mass Identity Numerical Comparison

1110 We compare $V_{\max}(1 - e^{-D_{\text{KL}}})$ (point-mass identity bound, exact under deterministic π^* per Lemma 5.1
1111 and the two-sided regret theorem) with $V_{\max}\sqrt{D_{\text{KL}}/2}$ (Pinsker bound, also valid under deterministic
1112 π^* but loose), $V_{\max}\sqrt{1 - e^{-D_{\text{KL}}}}$ (Bretagnolle–Huber inequality, the general envelope below which
1113 the identity sits at this corner), and the trivial envelope $V_{\max}$. Set $V_{\max} = 1$ for normalization.

Table 2: Identity vs. Pinsker vs. BH at $V_{\max} = 1$.

D_{KL}	$1 - e^{-D_{\text{KL}}}$ (identity)	$\sqrt{1 - e^{-D_{\text{KL}}}}$ (BH)	$\sqrt{D_{\text{KL}}/2}$ (Pinsker)	ratio
0.01	0.00995	0.0998	0.0707	7.10×
0.1	0.0952	0.308	0.224	2.35×
0.5	0.393	0.627	0.500	1.27×
1.0	0.632	0.795	0.707	1.12×
2.0	0.865	0.930	1.000	1.16×
4.0	0.982	0.991	1.414	1.02×
10.0	0.9995	0.9998	2.236	1.00×

1114 The point-mass identity bound is uniformly tighter than Pinsker, by a factor of $7\times$ at $D_{\text{KL}} = 0.01$ and
1115 converging to 1 as D_{KL} grows large (where both saturate at $V_{\max}$). It is also uniformly tighter than the
1116 BH inequality $\sqrt{1 - e^{-D_{\text{KL}}}}$ — the relation $x < \sqrt{x}$ on $(0, 1)$ — although both saturate together at
1117 $V_{\max}$ for large D_{KL} . Pinsker, by contrast, hits the trivial envelope $V_{\max} = 1$ at $D_{\text{KL}} = 2$ and exceeds it
1118 beyond — vacuous from $D_{\text{KL}} = 4$ onward (where the unclipped form returns 1.414), fully vacuous
1119 by $D_{\text{KL}} = 10$.

1120 The matching lower bound $\Delta_{\min}(1 - e^{-D_{\text{KL}}})$ has the same shape, scaled by $\Delta_{\min}/V_{\max}$. In the extremal
1121 value landscape ($\Delta_{\min} = V_{\max}$), the upper and lower bounds coincide and the regret is *identified*
1122 *exactly* by the KL coordinate. For typical landscapes, regret and KL are Lipschitz-equivalent with
1123 constants $\Delta_{\min}/V_{\max}$ and 1.

1124 F Prior-Art Retrieval Summary

1125 We conducted a structured prior-art retrieval (via Undermind) with the goal of identifying any frame-
1126 work composing all four target elements: (1) goal-feasibility-vs-policy-quality two-gap diagnostic;
1127 (2) an exact reverse-KL/TV regret identity (point-mass) under deterministic- π^* , or — more loosely
1128 — Bretagnolle–Huber-style use in regret analysis; (3) strategic-tempo / forgetting prerequisite tying
1129 convergence to rate of useful policy revision; (4) closed-loop interventional access making regret
1130 bounds learnable from on-policy data.

1131 **Result: 63 papers retrieved, abstract-level coverage estimated 75%, no direct anticipation in**
1132 **the retrieved set** The landscape splits into four largely separate strands corresponding to our four
1133 components; we did not find a published framework composing all four in this retrieval. Retrieval
1134 queries used the structured prior-art tool Undermind; the four-strand cite-and-distinguish in Section 2
1135 surfaces the closest neighbor in each strand.

1136 **Strongest negative signal (in the retrieved corpus): the Bretagnolle–Huber inequality is absent**
1137 **as an upper-bound coordinate for regret analysis in the retrieved RL / non-stationary literature**
1138 The information-theoretic regret line —^[19–24] — uses entropy / mutual information / information
1139 ratio / Pinsker / Hellinger; we did not find the BH inequality deployed *as an upper-bound regret*
1140 *coordinate*. (BH is well-known in bandit *lower-bound* proofs via Le Cam’s method and Tsybakov’s
1141 framework^[26], including in standard RL textbook treatments^[27]; that use is not contested here.) The
1142 *exact point-mass identity at the deterministic- π corner** (Lemma 5.1) — which strictly improves both
1143 Pinsker and the BH envelope at this corner — was likewise not found in the retrieved corpus, in either
1144 upper-bound or lower-bound use. We frame this as a literature-search observation against an explicitly
1145 bounded retrieval (63 papers, ~75% abstract-level coverage), not a formal absence claim against all
1146 published RL work; reviewers should read it as “to our knowledge / in the retrieved set” rather than
1147 “absolutely novel.”

1148 **Closest neighbors per strand:**^[1] (Strand 1 — variation-budget dynamic regret);^[5] (Strand 2 —
1149 closest two-term decomposition, different axis);^[8] (Strand 3 — closest tempo result, different form);^[14]
1150 (Strand 4 — closest causal-RL, stationary only).

1151 NeurIPS Paper Checklist

1152 1. Claims

1153 Question: Do the main claims made in the abstract and introduction accurately reflect the
1154 paper’s contributions and scope?

1155 Answer: [Yes]

1156 Justification: The abstract and the contribution subsection of the introduction state the
1157 four-component composition (two-gap diagnostic, point-mass reverse-KL/TV identity, multi-
1158 factor strategic tempo with forgetting prerequisite, closed-loop interventional access) and
1159 the cumulative dynamic regret rate. Each component is developed in the main-result and
1160 mechanism sections and proved in full in the appendices.

1161 Guidelines:

- 1162 • The answer [N/A] means that the abstract and introduction do not include the claims
1163 made in the paper.
- 1164 • The abstract and/or introduction should clearly state the claims made, including the
1165 contributions made in the paper and important assumptions and limitations. A [No] or
1166 [N/A] answer to this question will not be perceived well by the reviewers.
- 1167 • The claims made should match theoretical and experimental results, and reflect how
1168 much the results can be expected to generalize to other settings.
- 1169 • It is fine to include aspirational goals as motivation as long as it is clear that these goals
1170 are not attained by the paper.

1171 2. Limitations

1172 Question: Does the paper discuss the limitations of the work performed by the authors?

1173 Answer: [Yes]

1174 Justification: The conclusion section discusses theory-only scope, canonical-scope assump-
1175 tions (deterministic π^* , bounded value range, isolated optimum, singular trajectory), per-state
1176 vs. cumulative regret scope, partial uniqueness at the metric layer, and the dependence on
1177 architectural separation between M_t and goal state for coupled-goal architectures.

1178 Guidelines:

- 1179 • The answer [N/A] means that the paper has no limitation while the answer [No] means
1180 that the paper has limitations, but those are not discussed in the paper.
- 1181 • The authors are encouraged to create a separate “Limitations” section in their paper.
- 1182 • The paper should point out any strong assumptions and how robust the results are to
1183 violations of these assumptions (e.g., independence assumptions, noiseless settings,
1184 model well-specification, asymptotic approximations only holding locally). The authors
1185 should reflect on how these assumptions might be violated in practice and what the
1186 implications would be.
- 1187 • The authors should reflect on the scope of the claims made, e.g., if the approach was
1188 only tested on a few datasets or with a few runs. In general, empirical results often
1189 depend on implicit assumptions, which should be articulated.
- 1190 • The authors should reflect on the factors that influence the performance of the approach.
1191 For example, a facial recognition algorithm may perform poorly when image resolution
1192 is low or images are taken in low lighting. Or a speech-to-text system might not be
1193 used reliably to provide closed captions for online lectures because it fails to handle
1194 technical jargon.
- 1195 • The authors should discuss the computational efficiency of the proposed algorithms
1196 and how they scale with dataset size.
- 1197 • If applicable, the authors should discuss possible limitations of their approach to address
1198 problems of privacy and fairness.
- 1199 • While the authors might fear that complete honesty about limitations might be used by
1200 reviewers as grounds for rejection, a worse outcome might be that reviewers discover
1201 limitations that aren’t acknowledged in the paper. The authors should use their best
1202 judgment and recognize that individual actions in favor of transparency play an important
1203 role in developing norms that preserve the integrity of the community. Reviewers will
1204 be specifically instructed to not penalize honesty concerning limitations.

3. Theory assumptions and proofs

Question: For each theoretical result, does the paper provide the full set of assumptions and a complete (and correct) proof?

Answer: [\[Yes\]](#)

Justification: All theorems and lemmas are numbered, stated with their assumptions in-place, and cross-referenced. Main proofs (Theorems 4.1, 4.2, 4.3, 5.1, 6.1, 7.1, F.1, C.1, Lemma 5.2) appear in the body or appendices A, C, F, G; proof sketches in the body link to the full proofs in the appendices.

Guidelines:

- The answer [\[N/A\]](#) means that the paper does not include theoretical results.
- All the theorems, formulas, and proofs in the paper should be numbered and cross-referenced.
- All assumptions should be clearly stated or referenced in the statement of any theorems.
- The proofs can either appear in the main paper or the supplemental material, but if they appear in the supplemental material, the authors are encouraged to provide a short proof sketch to provide intuition.
- Inversely, any informal proof provided in the core of the paper should be complemented by formal proofs provided in appendix or supplemental material.
- Theorems and Lemmas that the proof relies upon should be properly referenced.

4. Experimental result reproducibility

Question: Does the paper fully disclose all the information needed to reproduce the main experimental results of the paper to the extent that it affects the main claims and/or conclusions of the paper (regardless of whether the code and data are provided or not)?

Answer: [\[N/A\]](#)

Justification: Theory-only paper; no original experiments. The worked-example reduction to ProST (Lemma 5.2) draws on Lee et al. 2023^[8]'s published experiments; reproducibility there is theirs.

Guidelines:

- The answer [\[N/A\]](#) means that the paper does not include experiments.
- If the paper includes experiments, a [\[No\]](#) answer to this question will not be perceived well by the reviewers: Making the paper reproducible is important, regardless of whether the code and data are provided or not.
- If the contribution is a dataset and/or model, the authors should describe the steps taken to make their results reproducible or verifiable.
- Depending on the contribution, reproducibility can be accomplished in various ways. For example, if the contribution is a novel architecture, describing the architecture fully might suffice, or if the contribution is a specific model and empirical evaluation, it may be necessary to either make it possible for others to replicate the model with the same dataset, or provide access to the model. In general, releasing code and data is often one good way to accomplish this, but reproducibility can also be provided via detailed instructions for how to replicate the results, access to a hosted model (e.g., in the case of a large language model), releasing of a model checkpoint, or other means that are appropriate to the research performed.
- While NeurIPS does not require releasing code, the conference does require all submissions to provide some reasonable avenue for reproducibility, which may depend on the nature of the contribution. For example
 - (a) If the contribution is primarily a new algorithm, the paper should make it clear how to reproduce that algorithm.
 - (b) If the contribution is primarily a new model architecture, the paper should describe the architecture clearly and fully.
 - (c) If the contribution is a new model (e.g., a large language model), then there should either be a way to access this model for reproducing the results or a way to reproduce the model (e.g., with an open-source dataset or instructions for how to construct the dataset).

1259 (d) We recognize that reproducibility may be tricky in some cases, in which case authors
1260 are welcome to describe the particular way they provide for reproducibility. In the
1261 case of closed-source models, it may be that access to the model is limited in some
1262 way (e.g., to registered users), but it should be possible for other researchers to have
1263 some path to reproducing or verifying the results.

1264 5. Open access to data and code

1265 Question: Does the paper provide open access to the data and code, with sufficient instructions
1266 to faithfully reproduce the main experimental results, as described in supplemental material?

1267 Answer: [N/A]

1268 Justification: Theory-only paper; no code or data accompanies the contribution.

1269 Guidelines:

- 1270 • The answer [N/A] means that paper does not include experiments requiring code.
- 1271 • Please see the NeurIPS code and data submission guidelines (<https://neurips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- 1272 • While we encourage the release of code and data, we understand that this might not
1273 be possible, so [No] is an acceptable answer. Papers cannot be rejected simply for not
1274 including code, unless this is central to the contribution (e.g., for a new open-source
1275 benchmark).
- 1276 • The instructions should contain the exact command and environment needed to run to
1277 reproduce the results. See the NeurIPS code and data submission guidelines (<https://neurips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- 1278 • The authors should provide instructions on data access and preparation, including how
1279 to access the raw data, preprocessed data, intermediate data, and generated data, etc.
- 1280 • The authors should provide scripts to reproduce all experimental results for the new
1281 proposed method and baselines. If only a subset of experiments are reproducible, they
1282 should state which ones are omitted from the script and why.
- 1283 • At submission time, to preserve anonymity, the authors should release anonymized
1284 versions (if applicable).
- 1285 • Providing as much information as possible in supplemental material (appended to the
1286 paper) is recommended, but including URLs to data and code is permitted.

1289 6. Experimental setting/details

1290 Question: Does the paper specify all the training and test details (e.g., data splits, hyperpa-
1291 rameters, how they were chosen, type of optimizer) necessary to understand the results?

1292 Answer: [N/A]

1293 Justification: Theory-only paper; no experiments.

1294 Guidelines:

- 1295 • The answer [N/A] means that the paper does not include experiments.
- 1296 • The experimental setting should be presented in the core of the paper to a level of detail
1297 that is necessary to appreciate the results and make sense of them.
- 1298 • The full details can be provided either with the code, in appendix, or as supplemental
1299 material.

1300 7. Experiment statistical significance

1301 Question: Does the paper report error bars suitably and correctly defined or other appropriate
1302 information about the statistical significance of the experiments?

1303 Answer: [N/A]

1304 Justification: Theory-only paper; no experiments.

1305 Guidelines:

- 1306 • The answer [N/A] means that the paper does not include experiments.
- 1307 • The authors should answer [Yes] if the results are accompanied by error bars, confidence
1308 intervals, or statistical significance tests, at least for the experiments that support the
1309 main claims of the paper.

- The factors of variability that the error bars are capturing should be clearly stated (for example, train/test split, initialization, random drawing of some parameter, or overall run with given experimental conditions).
- The method for calculating the error bars should be explained (closed form formula, call to a library function, bootstrap, etc.)
- The assumptions made should be given (e.g., Normally distributed errors).
- It should be clear whether the error bar is the standard deviation or the standard error of the mean.
- It is OK to report 1-sigma error bars, but one should state it. The authors should preferably report a 2-sigma error bar than state that they have a 96% CI, if the hypothesis of Normality of errors is not verified.
- For asymmetric distributions, the authors should be careful not to show in tables or figures symmetric error bars that would yield results that are out of range (e.g., negative error rates).
- If error bars are reported in tables or plots, the authors should explain in the text how they were calculated and reference the corresponding figures or tables in the text.

8. Experiments compute resources

Question: For each experiment, does the paper provide sufficient information on the computer resources (type of compute workers, memory, time of execution) needed to reproduce the experiments?

Answer: [N/A]

Justification: Theory-only paper; no experiments.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- The paper should indicate the type of compute workers CPU or GPU, internal cluster, or cloud provider, including relevant memory and storage.
- The paper should provide the amount of compute required for each of the individual experimental runs as well as estimate the total compute.
- The paper should disclose whether the full research project required more compute than the experiments reported in the paper (e.g., preliminary or failed experiments that didn't make it into the paper).

9. Code of ethics

Question: Does the research conducted in the paper conform, in every respect, with the NeurIPS Code of Ethics <https://neurips.cc/public/EthicsGuidelines>?

Answer: [Yes]

Justification: Theory paper without human subjects, datasets, or model release; conforms with the NeurIPS Code of Ethics throughout.

Guidelines:

- The answer [N/A] means that the authors have not reviewed the NeurIPS Code of Ethics.
- If the authors answer [No], they should explain the special circumstances that require a deviation from the Code of Ethics.
- The authors should make sure to preserve anonymity (e.g., if there is a special consideration due to laws or regulations in their jurisdiction).

10. Broader impacts

Question: Does the paper discuss both potential positive societal impacts and negative societal impacts of the work performed?

Answer: [N/A]

Justification: This is foundational theory work on convergence guarantees for non-stationary reinforcement learning; the contribution is a regret-bound analysis tool, not a deployable algorithm or model. We see no direct path from sharper non-stationary-RL convergence theory to either positive or negative societal applications that would not equally apply to the broader RL literature, and we do not release any models, datasets, or code that would carry deployment risks.

Guidelines:

- The answer [N/A] means that there is no societal impact of the work performed.
- If the authors answer [N/A] or [No], they should explain why their work has no societal impact or why the paper does not address societal impact.
- Examples of negative societal impacts include potential malicious or unintended uses (e.g., disinformation, generating fake profiles, surveillance), fairness considerations (e.g., deployment of technologies that could make decisions that unfairly impact specific groups), privacy considerations, and security considerations.
- The conference expects that many papers will be foundational research and not tied to particular applications, let alone deployments. However, if there is a direct path to any negative applications, the authors should point it out. For example, it is legitimate to point out that an improvement in the quality of generative models could be used to generate Deepfakes for disinformation. On the other hand, it is not needed to point out that a generic algorithm for optimizing neural networks could enable people to train models that generate Deepfakes faster.
- The authors should consider possible harms that could arise when the technology is being used as intended and functioning correctly, harms that could arise when the technology is being used as intended but gives incorrect results, and harms following from (intentional or unintentional) misuse of the technology.
- If there are negative societal impacts, the authors could also discuss possible mitigation strategies (e.g., gated release of models, providing defenses in addition to attacks, mechanisms for monitoring misuse, mechanisms to monitor how a system learns from feedback over time, improving the efficiency and accessibility of ML).

11. Safeguards

Question: Does the paper describe safeguards that have been put in place for responsible release of data or models that have a high risk for misuse (e.g., pre-trained language models, image generators, or scraped datasets)?

Answer: [N/A]

Justification: No data or models are released with this paper.

Guidelines:

- The answer [N/A] means that the paper poses no such risks.
- Released models that have a high risk for misuse or dual-use should be released with necessary safeguards to allow for controlled use of the model, for example by requiring that users adhere to usage guidelines or restrictions to access the model or implementing safety filters.
- Datasets that have been scraped from the Internet could pose safety risks. The authors should describe how they avoided releasing unsafe images.
- We recognize that providing effective safeguards is challenging, and many papers do not require this, but we encourage authors to take this into account and make a best faith effort.

12. Licenses for existing assets

Question: Are the creators or original owners of assets (e.g., code, data, models), used in the paper, properly credited and are the license and terms of use explicitly mentioned and properly respected?

Answer: [N/A]

Justification: No external code, datasets, or models are used; all cited prior work is credited via the references list.

Guidelines:

- The answer [N/A] means that the paper does not use existing assets.
- The authors should cite the original paper that produced the code package or dataset.
- The authors should state which version of the asset is used and, if possible, include a URL.
- The name of the license (e.g., CC-BY 4.0) should be included for each asset.

- For scraped data from a particular source (e.g., website), the copyright and terms of service of that source should be provided.
- If assets are released, the license, copyright information, and terms of use in the package should be provided. For popular datasets, paperswithcode.com/datasets has curated licenses for some datasets. Their licensing guide can help determine the license of a dataset.
- For existing datasets that are re-packaged, both the original license and the license of the derived asset (if it has changed) should be provided.
- If this information is not available online, the authors are encouraged to reach out to the asset's creators.

13. New assets

Question: Are new assets introduced in the paper well documented and is the documentation provided alongside the assets?

Answer: [N/A]

Justification: No new assets (datasets, code, models) are released with this paper.

Guidelines:

- The answer [N/A] means that the paper does not release new assets.
- Researchers should communicate the details of the dataset/code/model as part of their submissions via structured templates. This includes details about training, license, limitations, etc.
- The paper should discuss whether and how consent was obtained from people whose asset is used.
- At submission time, remember to anonymize your assets (if applicable). You can either create an anonymized URL or include an anonymized zip file.

14. Crowdsourcing and research with human subjects

Question: For crowdsourcing experiments and research with human subjects, does the paper include the full text of instructions given to participants and screenshots, if applicable, as well as details about compensation (if any)?

Answer: [N/A]

Justification: No crowdsourcing or research with human subjects.

Guidelines:

- The answer [N/A] means that the paper does not involve crowdsourcing nor research with human subjects.
- Including this information in the supplemental material is fine, but if the main contribution of the paper involves human subjects, then as much detail as possible should be included in the main paper.
- According to the NeurIPS Code of Ethics, workers involved in data collection, curation, or other labor should be paid at least the minimum wage in the country of the data collector.

15. Institutional review board (IRB) approvals or equivalent for research with human subjects

Question: Does the paper describe potential risks incurred by study participants, whether such risks were disclosed to the subjects, and whether Institutional Review Board (IRB) approvals (or an equivalent approval/review based on the requirements of your country or institution) were obtained?

Answer: [N/A]

Justification: No human subjects.

Guidelines:

- The answer [N/A] means that the paper does not involve crowdsourcing nor research with human subjects.

- Depending on the country in which research is conducted, IRB approval (or equivalent) may be required for any human subjects research. If you obtained IRB approval, you should clearly state this in the paper.
- We recognize that the procedures for this may vary significantly between institutions and locations, and we expect authors to adhere to the NeurIPS Code of Ethics and the guidelines for their institution.
- For initial submissions, do not include any information that would break anonymity (if applicable), such as the institution conducting the review.

16. Declaration of LLM usage

Question: Does the paper describe the usage of LLMs if it is an important, original, or non-standard component of the core methods in this research? Note that if the LLM is used only for writing, editing, or formatting purposes and does *not* impact the core methodology, scientific rigor, or originality of the research, declaration is not required.

Answer: [N/A]

Justification: LLMs were used only for writing assistance, formatting, and citation-verification cross-checks against published sources. The core mathematical content (theorem statements, proofs, the four-component composition, the point-mass identity, the sector-Lyapunov reduction) is author-developed and author-verified; the LLM did not originate methodological or theoretical contributions.

Guidelines:

- The answer [N/A] means that the core method development in this research does not involve LLMs as any important, original, or non-standard components.
- Please refer to our LLM policy in the NeurIPS handbook for what should or should not be described.