
A Unified Convergence Theory for Non-Stationary Reinforcement Learning

Anonymous Author(s)

Affiliation

Address

email

Abstract

1 Non-stationary reinforcement learning lacks a unified convergence theory: variation-
2 budget regret bounds, sliding-window methods, tempo-and-forgetting analyses,
3 two-term decompositions, and causal-RL frameworks each address fragments with-
4 out composing. We present a *compatible bundle of guarantees* that delivers all
5 three properties — non-stationarity persistence, explicit metric structure on policy
6 space, and on-policy learnability — *jointly*. Each of four components supplies
7 a distinct epistemic property: (i) a two-gap diagnostic separating goal-feasibility
8 from policy-quality routes corrective action; (ii) a point-mass reverse-KL/TV iden-
9 tity under deterministic optimum, $\text{TV}(\delta_{a^*}, Q) = 1 - e^{-D_{\text{KL}}(\delta_{a^*} \parallel Q)}$ — an *exact*
10 algebraic identity that lies strictly below the Pinsker and Bretagnolle–Huber en-
11 velopes at this corner — supplies a coordinate-optimal regret coordinate; (iii) a
12 multi-factor strategic tempo $\bar{\mathcal{T}}_\Sigma = \min_{(i,j)} \nu_{ij} \cdot \iota_{ij} \cdot (1 - \lambda_{ij})$ with structural sur-
13 vival inequality $\mathcal{T}_\Sigma > \rho_\Sigma / R_\Sigma$ supplies persistence; (iv) closed-loop interventional
14 access under sequential ignorability supplies on-policy learnability via the Pearl
15 causal hierarchy. Composition yields a Best-of-Both-Worlds dynamic regret theo-
16 rem $\tilde{O}(\min\{V_{\max} N_h \sqrt{(B_T + 1)T}, V_{\max} N_h (V_T + 1)^{1/3} T^{2/3}\})$ — automatically
17 adapting between piecewise-stationary segment count B_T and continuous variation
18 budget V_T without prior knowledge of either; the V_T exponent is information-
19 theoretically near-optimal at the deterministic- π^* corner. A structural failure-
20 class theorem identifies the *gain-decay* updates (count-accumulating Bayesian
21 schemes, Robbins–Monro gradient methods) as universally violating the forgetting
22 prerequisite for any positive disturbance rate; finite-gain mechanisms (constant-
23 step stochastic approximation, sliding windows, bounded memory, block restart)
24 face *bidirectional* thresholds instead. The deterministic-optimum scope extends
25 perturbatively to ϵ -stochastic and softmax-regularized policies with $O(\epsilon \log(1/\epsilon))$
26 and $O(\tau^{-1} \exp(-\Delta_{\min}/\tau))$ corrections (exponentially small with polynomial pref-
27 actor). Theory only; we make the ProST reduction rigorous at the steady-state
28 sector-parameter level, lifting it from a single-factor hyperparameter claim to a
29 multi-factor structural-threshold claim.

30 1 Introduction

31 Non-stationary reinforcement learning has matured along four largely separate research tracks. Each
32 has produced rigorous results; none has been composed with the others into a single convergence
33 theory.

34 The four tracks are visible in the recent literature as parallel lineages. **Variation-budget dynamic**
35 **regret under drift**^[1–4] gives sublinear dynamic regret under bounded total variation in reward
36 and transition dynamics. **Two-term regret decompositions**^[5–7] split the dynamic regret along an

37 *exploration-vs-adaptation* axis, isolating the cost of confidence-set construction from the cost of
 38 tracking a moving optimum. **Tempo and forgetting analyses**^[8–12] make policy-update timing or
 39 evidence-discount rate an explicit convergence variable. **Causal and interventional access**^[13–17]
 40 uses causal-graph structure to sharpen sample complexity, but operates in stationary settings only.

41 The information-theoretic regret literature^[18–23] traverses these tracks at a different layer (entropy / mu-
 42 tual information / information ratio / Pinsker / Hellinger) — but the *exact* reverse-KL / total-variation
 43 identity at the deterministic-optimum corner, strictly improving both Pinsker and the Bretagnolle–Hu-
 44 ber inequality^[24] at that corner, is absent from the retrieved corpus (Appendix F).

45 These four tracks share an evident common ancestor — the dynamic-regret analysis of online MDPs
 46 — but no published framework composes them along the *goal-feasibility-vs-policy-quality* axis with
 47 an exact point-mass identity, structural survival inequality, and closed-loop causal access. The
 48 fundamental question remains open:

49 *Can we design a unified convergence theory for non-stationary reinforcement*
 50 *learning that handles non-stationarity, has explicit metric structure on policy space,*
 51 *and is learnable from on-policy data?*

52 We present here a composition that delivers all three properties jointly. Composing four structural
 53 elements — a two-gap diagnostic, a point-mass reverse-KL/TV identity, a multi-factor strategic tempo
 54 with forgetting prerequisite, and closed-loop interventional access — yields a cumulative dynamic
 55 regret theorem with rate $\text{DynReg}(T) \leq \tilde{O}(V_{\max} N_h \sqrt{(B_T + 1)T})$ in the piecewise-stationary
 56 regime, where B_T counts piecewise-stationary segment boundaries (kernel-and-reward changes)
 57 along the realized trajectory and N_h is the planning horizon. The per-round coordinate of this rate is
 58 sharper than both Pinsker and Bretagnolle–Huber at the deterministic- π^* corner. The technical anchor
 59 is an exact algebraic identity: $\text{TV}(\delta_{a^*}, Q) = 1 - e^{-D_{\text{KL}}(\delta_{a^*} \parallel Q)}$, which holds for any policy Q
 60 with $Q(a^*) > 0$ and lies *strictly below* both the Pinsker envelope $\sqrt{D_{\text{KL}}/2}$ and the Bretagnolle–Huber
 61 envelope $\sqrt{1 - e^{-D_{\text{KL}}}}$ at this corner.

62 1.1 Contribution

63 Each of four components supplies a *distinct epistemic property* the others cannot deliver; together
 64 they form a *compatible bundle of guarantees*.

65 **Component 1 — corrective-action routing** The *two-gap diagnostic* splits the regret signal along the
 66 *goal-feasibility-vs-policy-quality* axis — orthogonal to the exploration-vs-adaptation decompositions
 67 of^[5–7] — routing four regimes (success, strategy problem, capability limit, both) to revise-policy /
 68 revise-model-class-or-horizon / revise-objective (Section 4).

69 **Component 2 — metric layer on policy space** Under deterministic optimum $\pi^* = \delta_{a^*}$, the per-
 70 round regret coordinate is *exact*: $\text{TV}(\pi^*, Q) = 1 - e^{-D_{\text{KL}}(\pi^* \parallel Q)}$, *coordinate-optimal among bounds*
 71 *depending only on TV* and strictly below both Pinsker and Bretagnolle–Huber envelopes at this corner
 72 (Appendix E). Deterministic scope extends perturbatively to ϵ -greedy and softmax-regularized optima
 73 with quantified corrections (Section 5.1).

74 **Component 3 — non-stationarity persistence** A *multi-factor strategic tempo* composed per
 75 element of observation rate ν , identifiability ι , and update gain $1 - \lambda$ yields a survival inequality $\mathcal{T}_{\Sigma} >$
 76 ρ_{Σ}/R_{Σ} — below threshold, no schedule preserves the agent’s strategic substate against disturbance
 77 (Section 5.2). A structural-class theorem identifies *gain-decay* updates (count-accumulating Bayesian,
 78 vanishing-step Robbins–Monro) as universally violating the threshold; finite-gain mechanisms face
 79 bidirectional ceilings (Appendix C.6). *Lifts* Lee et al. 2023^[81]’s tempo result from hyperparameter to
 80 structural-threshold claim.

81 **Component 4 — on-policy learnability** Under (C1) positivity, (C2) sequential ignorability (H_t
 82 d-separates a_t from o_{t+1} in $G_{\overline{a_t}}$), and (C3) known action mechanism, the closed-loop generates Pearl
 83 Level-2 data by construction^[25,26], making Component 2’s regret bound *learnable on-policy* under
 84 sufficient identifiability (Section 5.3). Three regimes A/B/C partition usable strength; coupled-goal
 85 architectures violate (C2) (Section 6).

86 **Composition** Via block-Cauchy–Schwarz in piecewise-stationary case and Wei and Luo 2021^[2]
87 MASTER wrapping in continuous variation, the four components compose into a Best-of-Both-Worlds
88 theorem $\mathbb{E}[\text{DynReg}(T)] \leq \tilde{O}\left(V_{\max}N_h \cdot \min\left\{\sqrt{(B_T+1)T}, (V_T+1)^{1/3}T^{2/3}\right\}\right) + V_{\max}N_h(1 -$
89 $p_{\text{id}})T$, *automatically adapting* to B_T or V_T without prior knowledge; matching Mao et al. 2021^[3]’s
90 near-optimal V_T rate, with Besbes et al. 2014^[27]’s lower bound applying at the deterministic- π^*
91 corner. The bias term vanishes in Regime A.

92 1.2 Scope and limitations

93 **What “convergence” means here** We do not claim pointwise policy convergence — under genuine
94 non-stationarity the optimum is itself moving. By *convergence* we mean: per-round regret-coordinate
95 convergence within stationary blocks; cumulative dynamic regret with sublinear / vanishing-average
96 rate when $B_T = o(T)$ (Cesàro tracking $\frac{1}{T} \sum_t (V_t^* - V(\pi_t)) \rightarrow 0$); and ultimate boundedness of
97 the strategic mismatch $\|\delta_\Sigma\|$ under the tempo threshold. Pointwise $V(\pi_t) \rightarrow V^*$ is structurally
98 unavailable under genuine non-stationarity; the Cesàro tracking statement is the right replacement.

99 **Theory only** No experiments. The point-mass identity is mathematically airtight (a two-line direct
100 calculation that specializes the BH family at its deterministic- π^* corner). The empirical grounding
101 for the strategic-tempo claim is provided indirectly through reduction to^[8] ProST as a special case
102 (Section 5); at the steady-state sector-parameter level the reduction is rigorous.

103 **Canonical scope** The composition theorem holds under deterministic optimum policy, bounded
104 value range, isolated optimum (so $\Delta_{\min} > 0$), and a singular causal trajectory in the sense made
105 precise in Section 3. The deterministic- π^* scope extends *perturbatively* to ϵ -stochastic and softmax-
106 regularized optima (Section 5.1); tied-optimum is a degraded one-sided form (Appendix C.5). Each
107 axis of the canonical scope degrades cleanly outside it (we discuss the failure modes in Section 6).

108 **Comparator regime** The cumulative rate covers both the piecewise-stationary B_T family and the
109 continuous-variation V_T family of^[3] via Best-of-Both-Worlds wrapping (Section 6).

110 2 Related Work

111 The four tracks named in Section 1 organize the literature.

112 **Variation-budget dynamic regret** ^[1–4] establish dynamic regret bounds under a *variation budget*
113 V_T on cumulative reward and transition variation.^[3] achieves the near-optimal continuous-variation
114 rate $\tilde{O}((SA)^{1/3}V_T^{1/3}HT^{2/3})$ via RestartQ-UCB;^[1] treats the piecewise-stationary case with restart-
115 on-change. Our cumulative rate $\tilde{O}(N_h\sqrt{(B_T+1)T})$ sits in the piecewise-stationary family — same
116 block-shape as Cheung-style restart-on-change — recovered as a corollary of composition with the
117 per-round identity. The novelty is the *threshold form* of the strategic-tempo inequality (environment
118 regimes where no $\lambda \in (0, 1)$ stabilizes the sector-Lyapunov mismatch), which dynamic-regret-
119 optimization analyses do not surface.

120 **Two-term regret decompositions along the exploration-vs-adaptation axis** ^[5–7] decompose the
121 dynamic regret into a *confidence-set construction* term and an *adaptation-under-non-stationarity*
122 term.^[28] treats a structurally adjacent question — constraint-region feasibility — using a two-term
123 decomposition along that axis. Same shape, different axis: ours is goal-feasibility-vs-policy-quality
124 (the satisfaction gap distinguishes “the goal is unattainable from M_t ” from “the current policy is
125 suboptimal”); theirs is uncertainty-source. Neither subsumes the other: a Long-Fei Li-style exploration
126 term can be present in any of our 2×2 cells, and an adaptation term likewise. The two axes are
127 orthogonal.

128 **Tempo and forgetting** ^[8] (ProST) defines an agent tempo as the schedule of policy update times
129 $\{t_1, \dots, t_K\}$ and computes the schedule that minimizes a dynamic regret upper bound.^[9] (Pausing
130 Policy Learning) refines this with non-zero hold durations.^[10–12] establish forgetting / sliding-window
131 analyses for non-stationary linear MDPs and bandits. We *lift* the ProST move along two axes
132 simultaneously: single-factor $\rightarrow$ multi-factor (per-element $\nu \cdot \iota \cdot (1 - \lambda)$ rather than scalar update
133 frequency); and hyperparameter-optimization $\rightarrow$ structural-survival inequality (the threshold below

which no schedule persists, irrespective of which schedule is chosen). ProST recovers as the Beta-Bernoulli special case $|E| = 1$, $\nu = \iota = 1$.

Causal and interventional access ^[13–17,29] use causal-graph structure to sharpen sample complexity in RL. Their setting is *stationary*. We compose the closed-loop interventional-access argument with non-stationarity. The connection between Pearl’s causal hierarchy^[25] and the data-generating process of an agent in a feedback loop^[26] is the technical bridge.

Information-theoretic regret (cross-cutting) ^[18–23,30] traverse the above tracks at a different layer, using entropy / mutual information / information ratio / Pinsker / Hellinger as the divergence backbone. The Russo–Van Roy information ratio uses Pinsker as its Cauchy–Schwarz follow-up; the IDS line and its descendants use mutual-information bounds. The Bretagnolle–Huber inequality^[24] — well-known in lower-bound proofs via Le Cam’s method^[31,32] — does not appear as an *upper-bound regret coordinate* in this corpus, and its exact point-mass corner (our Lemma 5.1) is *a fortiori* absent. We document the retrieval scope and methodology in Appendix F.

Adjacent (satisficing / feasibility) ^[33,34] use *satisficing* to mean “any policy above acceptance level β is acceptable.” Our δ_{sat} is structurally distinct: $V_O^{\min}$ is a property of the objective (set by the application domain), and $\delta_{\text{sat}} > 0$ signals goal-relative *infeasibility* (no policy in Π meets the threshold), not policy-relative satisficing.

Contemporaneous (post-March 2026) ^[34,35]; cited and distinguished without empirical-comparison demand per the contemporaneous-work cutoff convention.

3 Preliminaries

We adopt strict minimalism: only the notation and assumptions used directly in the main result and the mechanism statements appear here. Convention-hierarchy details (one-step vs. receding-horizon vs. Bellman continuation), the action-gap matching argument, tied-optimum extensions, and the deterministic-trajectory deployment modes are deferred to Appendix C.

Episodic non-stationary MDP A finite-horizon non-stationary MDP $(\mathcal{S}, \mathcal{A}, P_t, r_t, N_h)$ with state space $\mathcal{S}$, finite action space $\mathcal{A}$, time-indexed transition kernels $P_t(\cdot \mid s, a)$ and reward functions $r_t(s, a)$ allowed to vary across rounds t , and finite planning horizon N_h . Standard non-stationary-RL boundedness assumptions^[11]: per-step reward in $[0, 1]$.

Variation regimes A variation budget V_T measures continuous variation in transitions and rewards. We work in the *piecewise-stationary* specialization: time partitions into $B_T + 1$ maximal stationary intervals on which (P_t, r_t) is constant, with $B_T := |\{t : (P_t, r_t) \neq (P_{t-1}, r_{t-1})\}|$ counting *kernel-stationary-segment boundaries* (the right counter for the base learner’s per-block stationary-regret guarantee in (A5); see Appendix A.6). The two variation regimes (B_T vs. continuous V_T) coincide up to constants under abrupt-magnitude- δ piecewise-stationarity ($V_T \asymp \delta B_T$); cumulative results carry through to continuous V_T via Wei-Luo MASTER black-box wrapping (Section 6).

Policy distribution The agent’s policy at round t is a distribution $Q_t(\cdot \mid s)$ over actions; we write Q when state and round are understood. We adopt the canonical scope: the optimum policy at (M_t, s) is *deterministic*, $\pi^*(\cdot \mid M_t, s) = \delta_{a^*(s)}$ where $a^*(s) := \arg \max_a Q^\pi(s, a)$ is the optimal action under the agent’s current model M_t . Deterministic π^* is canonical for finite-MDP RL with isolated optima^[32]; the perturbative extension to ϵ -stochastic and softmax-regularized optima (Section 5.1, full derivation Appendix B.2) covers smooth deviations.

Value object Given an objective O , model M_t , policy π , and horizon N_h : $Q_O(M_t, a; \pi_{\text{cont}}, N_h) = \mathbb{E}[V_O(\tau) \mid M_t, \text{do}(a_t = a), a_{t+1} \sim \pi_{\text{cont}}]$. The $\text{do}(\cdot)$ ^[25] flags intervention on a_t , not conditioning. We adopt the **one-step improvement** convention $\pi_{\text{cont}} = \pi_{\text{current}}$ as default: most conservative, comparable across rounds, fixed-point-free. Receding-horizon and Bellman alternatives are in Appendix C.1.

Bounded value range $V_{\max}(M_t) := \max_a Q_O(M_t, a) - \min_a Q_O(M_t, a)$, finite under bounded reward and finite horizon.

181 **Action gap (isolated optimum)** $\Delta(a) := Q_O(a^*) - Q_O(a) \in [0, V_{\max}]$, $\Delta_{\min} := \min_{a \neq a^*} \Delta(a) >$
182 0 whenever the optimum is isolated.

183 **Strategic regret** For policy distribution Q : $R(Q) := Q_O(a^*) - \mathbb{E}_{a \sim Q}[Q_O(a)] = \sum_{a \neq a^*} Q(a) \cdot$
184 $\Delta(a)$. Three classical forms — $\mathbb{E}_{\pi^*}[V] - \mathbb{E}_Q[V]$, $V(a^*) - \mathbb{E}_Q[V]$, $\mathbb{E}_{\pi^*}[V - V_Q]$ — coincide under
185 deterministic π^* .

186 **Total variation and reverse-KL** For finite action measures P, Q on $\mathcal{A}$, $\text{TV}(P, Q) :=$
187 $\frac{1}{2} \sum_a |P(a) - Q(a)|$, $D_{\text{KL}}(P \| Q) := \sum_a P(a) \log \frac{P(a)}{Q(a)}$ with the convention $0 \log 0 = 0$.
188 Our bounds use the *reverse* KL direction (with $P = \pi^*$); the forward direction $D_{\text{KL}}(Q \| \pi^*)$ is $+\infty$
189 whenever Q has off-optimum mass and is therefore vacuous as a regret coordinate (full direction-
190 forcing argument in Appendix C.3).

191 **Singular trajectory** The agent operates on a single, non-forkable causal trajectory: a_t causally
192 precedes o_{t+1} , and the model state M_t 's sufficiency is defined relative to *this* trajectory rather than to
193 a model-state equivalence class. This is the operational ground for the do-operator above and for the
194 closed-loop interventional-access argument (Section 5.3).

195 4 Main Result

196 We present the four components first as motivation, then state the formal composition theorem, then
197 unpack what its conclusions mean.

198 4.1 The four components

199 Each component supplies a *distinct epistemic property* that the others cannot deliver. The composition
200 is necessary for the joint properties (rate + persistence + learnability + corrective-action routing), not
201 for the rate alone. The component-by-component failure-mode analysis is in Appendix A.6 below.

202 **Component 1: Two-gap diagnostic — corrective-action routing** Let Π be the policy class and
203 $V_O^{\min}$ the trajectory-value threshold above which the objective is met. Define the *satisfaction gap*
204 $\delta_{\text{sat}} := V_O^{\min} - \sup_{\pi \in \Pi} V_O(M_t, \pi; N_h)$ (goal-feasibility: positive means unmet under the best available
205 policy) and the *control regret* $\delta_{\text{regret}} := \sup_{\pi \in \Pi} V_O(M_t, \pi; N_h) - V_O(M_t, \pi_{\text{current}}; N_h) \geq 0$ (policy-
206 quality). The 2×2 cross routes four regimes — *Success* / *Strategy problem* (revise the policy) /
207 *Capability limit* (revise model, policy class, or horizon) / *Both* — along an axis *orthogonal* to the
208 exploration-vs-adaptation decompositions of^[5–7].

209 **Component 2: Point-mass reverse-KL/TV identity — metric layer on policy space** Un-
210 der deterministic $\pi^* = \delta_{a^*}$, the reverse-KL and total variation collapse to an *exact identity*
211 $\text{TV}(\pi^*, Q) = 1 - e^{-D_{\text{KL}}(\pi^* \| Q)}$, yielding the two-sided regret characterization $\Delta_{\min}(1 - e^{-K}) \leq$
212 $R(Q) \leq V_{\max}(1 - e^{-K})$ — *Lipschitz-equivalent* with constants $\Delta_{\min}/V_{\max}$ and 1, and *coordinate-*
213 *optimal among bounds depending only on TV*. The identity sits *strictly below* both the Pinsker $\sqrt{D_{\text{KL}}/2}$
214 and Bretagnolle–Huber $\sqrt{1 - e^{-D_{\text{KL}}}}$ envelopes at this corner (Appendix E); the deterministic- π^*
215 scope extends perturbatively to ϵ -greedy and softmax-regularized optima with quantified corrections
216 (Section 5.1 + Appendix B.2).

217 **Component 3: Multi-factor strategic tempo with forgetting prerequisite — non-stationarity**
218 **persistence** A multi-factor strategic tempo composed per revisable policy element of an observation
219 rate ν , an identifiability coefficient $\iota \in [0, 1]$, and an update gain $1 - \lambda$ yields a structural survival
220 inequality $\mathcal{T}_{\Sigma}^{\text{bn,ss}} := \min_{(i,j) \in E} \nu_{ij} \cdot \iota_{ij} \cdot (1 - \lambda_{ij}) > \rho_{\Sigma}/R_{\Sigma}$, the *bottleneck* (per-element
221 minimum) being the right aggregator since adversarial disturbance concentrates on the weakest
222 element. A structural-class theorem on the *gain-decay* class $\mathcal{A}_{\text{decay}}$ (count-accumulating Bayesian
223 without forgetting, observation-aggregating without restart, gradient-based with vanishing step) shows
224 every member eventually violates the threshold for any positive disturbance rate (Appendix C.6);
225 finite-gain mechanisms (constant-step SA, sliding-window, bounded-memory, block-restart) face
226 *bidirectional* ceilings instead. This *lifts* the tempo result of Lee et al. 2023^[8] from hyperparameter
227 optimization to a structural-threshold claim. Full Model (Σ) and derivation in Section 5.2.

Component 4: Closed-loop interventional access — on-policy learnability Under (C1) *positivity* ($\pi_t(a \mid H_t) \geq p_{\min} > 0$), (C2) *sequential ignorability* (H_t d-separates a_t from o_{t+1} in the mutilated graph $G_{\bar{a}_t}$), and (C3) *known action mechanism*, the loop’s data-generating process is *interventional* — the conditional $P(o_{t+1} \mid a_t, H_t)$ coincides with the do-distribution $P(o_{t+1} \mid \text{do}(a_t), H_t)$ on the realized-action support (Section 5.3). This makes Component 2’s regret bound *learnable on-policy* in regimes with sufficient identifiability, with bias from misidentification quantified by $1 - p_{\text{id}}$ (Section 5.4). Three regimes^[26] partition usable strength: A (intervention-rich, $\iota \approx 1$), B (partial), C (observation-only).

4.2 The composition theorem

Theorem 4.1 (Composition: joint guarantees with conditional cumulative reduction). *Let $(\mathcal{S}, \mathcal{A}, P_t, r_t, N_h)$ be a non-stationary MDP with bounded reward, finite action space, and for each round t a deterministic optimum policy $\pi_t^* = \delta_{a_t^*}$ with action gap $\Delta_{\min} > 0$. Suppose the agent operates on a singular causal trajectory in the sense of Section 3 and updates its policy with per-element exponential forgetting at rates $\{\lambda_{ij}\}$. Let $K_t(s) := D_{\text{KL}}(\delta_{a_t^*(s)} \parallel Q_t(\cdot \mid s))$ be the per-state reverse-KL coordinate (in the bandit case $N_h = 1$ the state argument is suppressed). If*

(A1) Metric extension. *The identity $\text{TV}(\delta_{a_t^*(s)}, Q_t(\cdot \mid s)) = 1 - e^{-K_t(s)}$ of Lemma 5.1 is read in the extended real sense: when $Q_t(a_t^*(s) \mid s) = 0$, $K_t(s) = +\infty$ and $1 - e^{-K_t(s)} = 1$ recovers the trivial bound $R(Q_t \mid s) \leq V_{\max}$. No pointwise positivity is required; vanilla deterministic UCB / UCBVI / UCRL2 deployed at the optimism-argmax are in scope without modification.*

(A2) Multi-factor forgetting prerequisite. *Within Model (Σ) of Section 5.2, the bottleneck strategic tempo $\mathcal{T}_{\Sigma}^{\text{bn,ss}} := \min_{(i,j) \in E} \nu_{ij} \cdot \iota_{ij} \cdot (1 - \lambda_{ij}) > \rho_{\Sigma} / R_{\Sigma}$.*

(A3) Identifiability regime. *The action mechanism is known and (C1)–(C3) of Section 5.3 hold; the agent operates in Regime A or Regime B with per-state argmax-correctness probability $p_{\text{id}}(s)$, with floor $p_{\text{id}} := \min_s p_{\text{id}}(s)$.*

(A4) Two-gap diagnostic. *The agent applies the satisfaction-gap / control-regret 2×2 to route corrective action.*

(A5) Piecewise-stationary block structure with restart-on-change identity-tight base learner. *The time axis partitions into $B_T + 1$ stationary intervals on which the MDP is stationary and $\pi_t^*(\cdot \mid s) = \delta_{a_t^*(s)}$ is fixed per state. The base learner is restarted at each block boundary; within each restarted interval of length L it satisfies $\sum_{t=1}^L \mathbb{E}[\overline{\text{TV}}_t] \leq 2c\sqrt{L}$, where $\overline{\text{TV}}_t := \frac{1}{N_h} \sum_{h=0}^{N_h-1} \mathbb{E}_{s_h \sim d_h^{Q_t}} [\text{TV}(\pi_t^*(\cdot \mid s_h), Q_t(\cdot \mid s_h))]$ is the occupancy-weighted per-round coordinate. (A non-restarting carryover variant is plausible but requires bounding cross-block occupancy shift; see Remark in Section 5.)*

Then:

(i) Per-round regret is two-sided identity-bounded. *At every visited state, $\Delta_{\min}(1 - e^{-K_t(s)}) \leq R(Q_t \mid s) \leq V_{\max}(1 - e^{-K_t(s)})$.*

(ii) Aggregate strategic mismatch is ultimately bounded. *Under (A2), the modeled mismatch $\|\delta_{\Sigma}\|$ is ultimately bounded within $R_{\Sigma}^* = \rho_{\Sigma} / \mathcal{T}_{\Sigma}^{\text{bn,ss}}$.*

(iii) The KL coordinate is computable with controlled bias. *$-\log Q_t(a_{\text{ag},t}^*(s) \mid s)$ is computable directly from Q_t . Per-state bias: $\mathbb{E}[|\hat{D}_t(s) - D_t^{\text{true}}(s)|] \leq (1 - p_{\text{id}}(s)) \log(1/q_0)$ under $Q_t(\cdot \mid s) \geq q_0$ at the two argmax candidates. Vanishes in Regime A.*

(iv) The 2×2 cell identifies the corrective-action class. *$(\delta_{\text{sat}}, \delta_{\text{regret}})$ routes to revise-policy / revise-model-class-or-horizon / revise-objective.*

(v) Trajectory-level blockwise dynamic regret. *Under (A5), $\mathbb{E}[\text{DynReg}(T)] := \mathbb{E}\left[\sum_{t=1}^T (V_t^{\pi_t^*} - V_t^{Q_t})\right] \leq 2c V_{\max} N_h \sqrt{(B_T + 1)T} + V_{\max} N_h (1 - p_{\text{id}}) \cdot T$.*

The proof composes the per-round identity (i) with the simulation lemma^[36–39], a block decomposition at optimum-change events, Cauchy–Schwarz across the $B_T + 1$ blocks, and the bias bound from (iii). Full proof in Appendix A; the four key lemmas are surfaced in Section 5.

275 4.3 Unpacking the conclusions

276 Theorem 4.1 asserts a non-stationarity-aware convergence theory with three jointly-novel properties:
 277 *non-stationarity persistence* (via (ii)/(A2)), *explicit metric structure on policy space* (via (i)/Component
 278 2), *on-policy learnability* (via (iii)/(A3)). Conclusion (iv) is the connective tissue; conclusion (v) the
 279 cumulative rate from composing per-round metric with base-learner stationary rate.

280 *The per-round coordinate is exact, not an inequality.* Under deterministic π^* , the identity-bound (i) is
 281 $\Delta_{\min}/V_{\max}$ -Lipschitz-equivalent to regret with both endpoints achieved on extremal value landscapes
 282 — *coordinate-optimal among bounds depending only on TV*. Strictly sharper than Pinsker $\sqrt{D_{\text{KL}}/2}$
 283 (exact, not approximate) and Bretagnolle–Huber $\sqrt{1 - e^{-D_{\text{KL}}}}$ at this corner ($x < \sqrt{x}$ on $(0, 1)$).

284 *Cumulative rate: piecewise-stationary B_T family by direct aggregation, continuous-variation V_T
 285 family by MASTER wrapping.* The lifted UCRL2/UCBVI rate $\tilde{O}(N_h^2 \sqrt{SA(B_T + 1)T})$ (Appendix D)
 286 sits in the restart-on-change family of^[1]; via^[2]’s MASTER black-box, the same identity-routed base
 287 learner achieves Best-of-Both-Worlds $\tilde{O}(\min\{\sqrt{(B_T + 1)T}, (V_T + 1)^{1/3}T^{2/3}\})$ automatically —
 288 matching^[3]’s near-optimal V_T exponent (parameter $(B_T$ vs $V_T)$, T -scaling, S , A -scaling differ from
 289 Mao’s RestartQ-UCB; full development in Section 6 + (A5’)-BoBW Remark in Appendix A).

290 *The bias term in (v) is a value-coordinate misidentification penalty.* The second term $V_{\max}N_h(1 -$
 291 $p_{\text{id}}) \cdot T$ is the expected value gap between true and agent-identified optima, $V_{\max}$ -bounded on the
 292 misidentification event of probability $1 - p_{\text{id}}$ aggregated over N_h horizon steps. Vanishes in Regime A;
 293 in Regime B the dominant rate term is the metric one when $1 - p_{\text{id}}$ is small relative to $\sqrt{(B_T + 1)T}$;
 294 in Regime C saturates at the trivial $V_{\max}N_h \cdot T$ envelope (Appendix D flags Regime C out-of-scope).
 295 *Strict assumption-set weakening:* (v) does not require Lemma 4’s two-point support $Q_t \geq q_0$ — that
 296 condition is needed only for conclusion (iii)’s diagnostic computability.

297 *The bundle is compatible, not integrated.* (v)’s rate goes through (A5) and (i) alone — Component
 298 2 supplies the metric, (A5) the base-learner regret, the $\sqrt{(B_T + 1)T}$ aggregation falls out. The
 299 other components carry the *other* guarantees: Component 1 the routing (iv); Component 3 the
 300 persistence (ii) (without it, K_t may itself drift unboundedly under $\mathcal{A}_{\text{decay}}$); Component 4 the on-policy
 301 learnability via (iii) (without it, (v) becomes vacuous in Regime C). The horizon factor N_h in (v) is
 302 the linear-in- N_h simulation-lemma penalty, sharper than the N_h^2 of TRPO-style worst-state bounds^[40].
 303 Component-by-component failure-mode analysis in Appendix A.6.

304 5 Mechanism

305 We surface four key lemmas and chain them through a block decomposition. Algebra is deferred to
 306 the appendices.

307 5.1 Key Lemma 1: Point-mass reverse-KL/TV identity

308 **Lemma 5.1** (Point-mass reverse-KL/TV identity). *For deterministic $\pi^* = \delta_{a^*}$ (per visited state,
 309 with $a^* = a^*(s)$ and $Q = Q(\cdot \mid s)$) and any policy Q , $D_{\text{KL}}(\pi^* \parallel Q) = -\log Q(a^*) =$
 310 $-\log(1 - \text{TV}(\pi^*, Q))$, equivalently $\text{TV}(\pi^*, Q) = 1 - e^{-D_{\text{KL}}(\pi^* \parallel Q)}$. The identity is read in the
 311 extended real sense: when $Q(a^*) = 0$ both sides equal 1, with the natural convention $D_{\text{KL}} = +\infty$
 312 and $e^{-\infty} = 0$ — the identity holds for all policies Q , including Diracs at suboptimal actions.*

313 *Intuition.* Both D_{KL} and TV collapse cleanly under the point-mass — every $a \neq a^*$ term vanishes
 314 by $0 \log 0 = 0$. Substituting into BH gives $1 - e^{-D_{\text{KL}}} \leq \sqrt{1 - e^{-D_{\text{KL}}}}$, strict on $D_{\text{KL}} > 0$: the
 315 identity is the *exact value* at this corner rather than BH-at-equality. Extends perturbatively to ϵ -greedy
 316 and softmax-regularized optima with $O(\epsilon \log(1/\epsilon))$ and $O(\tau^{-1} \exp(-\Delta_{\min}/\tau))$ corrections under
 317 full-support $Q \geq q_0$ (Appendix B.2). Full proof in Appendix B.

318 5.2 Key Lemma 2: Forgetting prerequisite

319 The strategic mismatch state $\delta_{\Sigma} = (\delta_{ij})_{(i,j) \in E}$ in the Euclidean norm $\|\delta_{\Sigma}\|^2 = \sum_{(i,j)} \delta_{ij}^2$ evolves
 320 in discrete time as $\mathbb{E}[\Delta \delta_{ij} \mid \delta_{ij}] = -\alpha_{ij} \delta_{ij} + w_{ij}$, $|w_{ij}| \leq \rho_{\Sigma}/|E|^{1/2}$, where $\alpha_{ij} \geq 0$ is a per-
 321 element correction rate, w_{ij} a bounded disturbance with budget $\sum w_{ij}^2 \leq \rho_{\Sigma}^2$, and R_{Σ} the operative

norm-radius beyond which tracking is lost. Beta-Bernoulli edge updates and Robbins–Monro gradient updates both instantiate this *Model* (Σ).

Lemma 5.2 (Multi-factor forgetting prerequisite). *Within Model (Σ) with diagonal per-element correction rates $\alpha_{ij} = \nu_{ij} \cdot \iota_{ij} \cdot \eta_{\text{edge},ij}$ (per-element observation rate $\nu_{ij} \in [0, 1]$, regime-adjusted identifiability $\iota_{ij} \in [0, 1]$, edge-update gain $\eta_{\text{edge},ij} > 0$), under per-element exponential forgetting at rate $\lambda_{ij} \in (0, 1)$ giving steady-state $\eta_{\text{edge},ij}^{\text{ss}} \approx 1 - \lambda_{ij}$, ultimate boundedness of $\|\delta_\Sigma\|$ within $R_\Sigma^* = \rho_\Sigma / \mathcal{T}_\Sigma^{\text{bn,ss}}$ is sufficient under $\mathcal{T}_\Sigma^{\text{bn,ss}} := \min_{(i,j) \in E} (\nu_{ij} \cdot \iota_{ij} \cdot (1 - \lambda_{ij})) > \rho_\Sigma / R_\Sigma$. The threshold is sharp inside the diagonal sector model: when the inequality reverses, an adversarial disturbance concentrating on the bottleneck element drives $\|\delta_\Sigma\|$ beyond R_Σ within the modeled dynamics.*

Intuition (why bottleneck, not sum). Bounding $\sum_{(i,j)} \alpha_{ij} \delta_{ij}^2 \geq c \|\delta_\Sigma\|^2$ uniformly forces $c = \min \alpha_{ij}$ — adversarial disturbance concentrating on the weakest element saturates at the bottleneck (the same effect makes per-dimension persistence conditions sharper than scalar conditions in adaptive control^[41]). Ultimate boundedness then follows from Khalil’s sector-Lyapunov template; full proof in Appendix B, structural-class theorem on $\mathcal{A}_{\text{decay}}$ in Appendix C.6.

ProST as a sector-level reduction. The block-restart structure of ProST^[8] recovers as the *impulsive limit* of Model (Σ): between updates the mismatch evolves under disturbance alone, and each update contracts the mismatch by per-update gain γ . The Hespanha–Liberzon–Teel reverse-ADT framework^[42] gives the threshold (full development in Appendix B.4); under uniform schedules this recovers ProST’s K/T form with the impulse gain γ made explicit. The reframe lifts ProST’s *tempo-as-hyperparameter* framing into *tempo-as-stability-margin* — the longest inter-update gap, not the average, sets the threshold, exposing the tradeoff between update frequency and per-update impulse strength at the sector-Lyapunov level.

5.3 Key Lemma 3: Closed-loop interventional access

Lemma 5.3 (Loop generates interventional samples). *Let H_t denote the agent’s information state at round t (model state M_t , history of observations and actions). Under (C1) positivity — $\pi_t(a \mid H_t) \geq p_{\min} > 0$ on the support of H_t for the actions of interest; (C2) sequential ignorability — H_t d -separates a_t from o_{t+1} in the mutilated graph $G_{\overline{a_t}}$ (equivalently, $a_t \perp\!\!\!\perp Y^{(a_t)} \mid H_t$ in potential-outcome notation); (C3) known action mechanism — the behavior policy $\pi_t(a \mid H_t)$ is known: executed actions are interventions relative to the environment’s transition kernel, and the loop generates samples from $P(o_{t+1} \mid \text{do}(a_t = a), H_t)$ — i.e., Pearl Level-2 interventional kernels conditional on the decision context. The substantive content for the identification claim is (C2); under the singular-trajectory + agent-as-policy architecture of Section 3, (C1) reduces to realized-action positivity (automatic) and (C3) to agent-policy-queryability (architectural). (C1) and (C3) are stated explicitly for forward-compatibility with off-policy / counterfactual extensions.*

Intuition. By temporal ordering, a_t causally precedes o_{t+1} . Under (C1)–(C3) the conditional $P(o_{t+1} \mid a_t, H_t)$ coincides with the interventional $P(o_{t+1} \mid \text{do}(a_t), H_t)$ on the support where $\pi_t(a_t \mid H_t) > 0$ — the loop generates conditional-Level-2 samples regardless of how the agent updates internally. Three regimes^[26] partition usable strength: A ($\iota \approx 1$, do-effects identified cleanly), B ($\iota \in (0, 1)$, bias $\propto 1 - \iota$), C (observation-only, not on-policy realizable; observer-on-sub-agent extensions in Appendix C.8). Components 2 and 4 are jointly load-bearing — neither alone makes the bound learnable.

5.4 Key Lemma 4: Bias bound under partial identifiability

Lemma 5.4 (Bias bound). *Assume $Q_t(a \mid s) \geq q_0 > 0$ for both $a = a_{\text{ag},t}^*(s)$ (the agent’s identified optimum) and $a = \tilde{a}_t^*(s)$ (the true optimum). Let $\hat{D}_t(s) := -\log Q_t(a_{\text{ag},t}^*(s) \mid s)$ and $D_t^{\text{true}}(s) := -\log Q_t(\tilde{a}_t^*(s) \mid s)$. Then $|\hat{D}_t(s) - D_t^{\text{true}}(s)| \leq \mathbf{1}[a_{\text{ag},t}^*(s) \neq \tilde{a}_t^*(s)] \cdot \log(1/q_0)$, and taking expectations: $\mathbb{E}[|\hat{D}_t(s) - D_t^{\text{true}}(s)|] \leq (1 - p_{\text{id}}(s)) \log(1/q_0)$, where $p_{\text{id}}(s) := \Pr[a_{\text{ag},t}^*(s) = \tilde{a}_t^*(s)]$.*

Intuition. Indicator decomposition: zero on the matching event; on the misidentification event both $\log Q_t$ values lie in $[\log q_0, 0]$ so their absolute difference is at most $\log(1/q_0)$. Bias scales linearly in misidentification mass — vanishes in Regime A, controlled in Regime B, saturates at $\log(1/q_0)$ in

372 Regime C. The two-point support condition (at the two argmax candidates) is strictly weaker than the
 373 perturbative identity’s full-support condition (Section 5.1). Full proof in Appendix B.

374 5.5 Combining the lemmas: cumulative regret in four steps

375 *Step 1 — Per-round bound via the simulation lemma.* The performance-difference lemma^[36–39] gives,
 376 for any policies π_t^* and Q_t , $V_t^{\pi_t^*}(s_0) - V_t^{Q_t}(s_0) \leq V_{\max} \cdot N_h \cdot \overline{\text{TV}}_t$, where $\overline{\text{TV}}_t$ is the occupancy-
 377 weighted TV of (A5). By Key Lemma 1 applied per state, $\text{TV}(\pi_t^*, Q_t) = 1 - e^{-K_t(s)}$ — sharper
 378 than Pinsker, preserved under occupancy aggregation. The N_h is the linear-in- N_h simulation-lemma
 379 penalty, sharper than the N_h^2 of TRPO-style worst-state bounds^[40].

380 *Step 2 — Block decomposition.* Partition $[1, T]$ at (P, r) -change events into blocks of lengths Δ_i (the
 381 kernel-stationary-segment boundaries of Section 3). (A5) — restart-on-change base learner — gives
 382 within-block $\sum_{t \in \text{block}_i} \mathbb{E}[\overline{\text{TV}}_t] \leq 2c\sqrt{\Delta_i}$.

383 *Step 3 — Cauchy–Schwarz across $B_T + 1$ blocks.* $\sum_i \sqrt{\Delta_i} \leq \sqrt{(B_T + 1)T}$, giving
 384 $\sum_t \mathbb{E}[V_{\max} N_h \overline{\text{TV}}_t] \leq 2cV_{\max} N_h \sqrt{(B_T + 1)T}$.

385 *Step 4 — Misidentification penalty.* Steps 1–3 give the rate term against the agent’s *identified* optimum
 386 $\pi_{\text{ag},t}^*$ (what (A5) delivers); (v) compares to the *true* optimum $\tilde{\pi}_t^*$. Applying the simulation lemma to
 387 the two point-mass policies, each per-step bracket reduces to $Q_h^{\tilde{\pi}_t^*}(s_h, \tilde{a}^*) - Q_h^{\pi_{\text{ag},t}^*}(s_h, a_{\text{ag}}^*) \in [0, V_{\max}]$
 388 on the misidentification event and zero on the matching event. Aggregated over N_h horizon steps
 389 with floor $p_{\text{id}} := \min_s p_{\text{id}}(s)$, the per-round penalty is at most $V_{\max} N_h (1 - p_{\text{id}})$ — value-coordinate
 390 throughout, vanishing in Regime A. The support condition $Q_t \geq q_0$ is needed only for conclusion
 391 (iii)’s diagnostic readout (Lemma 4), not for the rate term.

392 **Combining:** $\mathbb{E}[\text{DynReg}(T)] \leq 2cV_{\max} N_h \sqrt{(B_T + 1)T} + V_{\max} N_h (1 - p_{\text{id}}) \cdot T$, with Cesàro tracking
 393 $\frac{1}{T} \sum_t (V_t^{\tilde{\pi}_t^*} - V_t^{Q_t}) \rightarrow 0$ when $B_T = o(T)$ and $p_{\text{id}} \rightarrow 1$.

394 In the bandit special case the rate sharpens logarithmically; the non-restarting carryover variant of
 395 (A5) is deferred. Full proof of (v) in Appendix A; bandit-case sharpening + algorithmic instantiation
 396 in Appendix D.

397 6 Conclusion

398 The four-component composition delivers three jointly-novel properties — non-stationarity persistence,
 399 explicit metric structure on policy space, on-policy learnability — and a Best-of-Both-Worlds rate
 400 $\tilde{O}(\min\{V_{\max} N_h \sqrt{(B_T + 1)T}, V_{\max} N_h (V_T + 1)^{1/3} T^{2/3}\})$ adapting between B_T and V_T via Wei-
 401 Luo MASTER wrapping (Appendix A). The V_T exponent is information-theoretically near-optimal at
 402 the deterministic- π^* corner per^[27]; the per-round identity sharpens constants and per-round form.

403 **On coupled-goal architectures** (C2) presupposes architectural separation between M_t and the goal
 404 state — goal-conditioned LLM policies violate this and need additional machinery^[43] as starting
 405 point).

406 **Open directions** An *agent-experienced* variation budget $V_T^{\text{eff}, \Sigma}$ or identity-coordinate variation
 407 $V_T^{(K)}$ would explain why agents in slow-drift environments outperform B_T -pessimistic predictions;
 408 Bernstein-type concentration^[39] could shave one $\sqrt{N_h}$ factor from the lifted rate; joint uniqueness
 409 across all three properties and strategic-tempo beyond Model (Σ) remain open (Theorem C.3).

410 **Practitioner takeaway** Use $\mathcal{T}_{\Sigma}^{\text{agg}, \text{ss}} > |E| \rho_{\Sigma} / R_{\Sigma}$ as a fail-fast pre-check; route corrective action
 411 via the 2×2 diagnostic; apply the per-round identity wherever deterministic- π^* scope holds.

412 References

413 [1] Wang Chi Cheung, David Simchi-Levi, and Ruihao Zhu. Reinforcement learning for non-
 414 stationary Markov decision processes: The blessing of (more) optimism. *Management Science*,
 415 68(3):1697–1716, 2022. doi: 10.1287/mnsc.2021.4123. ArXiv 2020.

- [2] Chen-Yu Wei and Haipeng Luo. Non-stationary reinforcement learning without prior knowledge: An optimal black-box approach, 2021.
- [3] W Mao, K Zhang, R Zhu, D Simchi-Levi, and T Başar. Near-optimal model-free RL in non-stationary episodic MDPs. In *ICML*, 2021.
- [4] P Gajane, R Ortner, and P Auer. Variational regret bounds for reinforcement learning. In *UAI*, 2019.
- [5] Long-Fei Li, Peng Zhao, and Zhi-Hua Zhou. Dynamic regret of adversarial MDPs with unknown transition and linear function approximation. In *Proceedings of the AAAI Conference on Artificial Intelligence*, volume 38, 2024. doi: 10.1609/aaai.v38i12.29261.
- [6] Y Fei, Z Yang, Z Wang, and Q Xie. Dynamic regret of policy optimization in non-stationary environments. *arXiv:2007.00148*, 2020.
- [7] F. E Stradi, A Lunghi, M Castiglioni, A Marchesi, and N Gatti. Learning CMDPs with non-stationary rewards and constraints. *arXiv:2405.14372*, 2024.
- [8] Hyunin Lee, Yuhao Ding, Jongmin Lee, Ming Jin, Javad Lavaei, and Somayeh Sojoudi. Tempo adaptation in non-stationary reinforcement learning. In *Advances in Neural Information Processing Systems 36 (NeurIPS 2023)*, 2023. arXiv 2309.14989; OpenReview 7R8noSP4vL.
- [9] Hyunin Lee, Ming Jin, Javad Lavaei, and Somayeh Sojoudi. Pausing policy learning in non-stationary reinforcement learning. In *Proceedings of the 41st International Conference on Machine Learning (ICML 2024)*, Proceedings of Machine Learning Research, 2024. arXiv 2405.16053.
- [10] A Touati and P Vincent. Efficient learning in non-stationary linear MDPs (OPT-WLSVI). *arXiv:2010.12870*, 2020.
- [11] Yoan Russac, Claire Vernade, and Olivier Cappé. Weighted Linear Bandits for Non-Stationary Environments. In *Advances in Neural Information Processing Systems 32 (NeurIPS 2019)*, 2019.
- [12] A Garivier and E Moulines. On upper-confidence bound policies for non-stationary bandit problems. *arXiv:0805.3415*, 2008.
- [13] Junzhe Zhang and E Bareinboim. MDPs with unobserved confounders: a causal approach. (*Tech. report; Columbia.*), 2016.
- [14] Junzhe Zhang and Elias Bareinboim. Online reinforcement learning for mixed policy scopes. In *Advances in Neural Information Processing Systems*, 2022.
- [15] Y Lu, A Meisami, and A Tewari. Efficient RL with prior causal knowledge. *CLEaR*, 2022.
- [16] L Wang, Z Yang, and Z Wang. Provably efficient causal RL with confounded observational data (DOVI). In *NeurIPS. arXiv:2006.12311 (2020)*, 2021.
- [17] Junzhe Zhang. Designing optimal dynamic treatment regimes: a causal RL approach. In *ICML*, 2020.
- [18] D Russo and B Van Roy. An information-theoretic analysis of Thompson sampling. *J. Mach. Learn. Res.*, 2014.
- [19] Daniel Russo and Benjamin Van Roy. Learning to optimize via information-directed sampling. *Operations Research*, 2014.
- [20] X Lu and B Van Roy. Information-theoretic confidence bounds for reinforcement learning. In *NeurIPS*, 2019.
- [21] Seungki Min and Daniel Russo. An information-theoretic analysis of nonstationary bandit learning. In *International Conference on Machine Learning (ICML 2023)*, 2023.
- [22] T Lattimore and A György. Mirror descent and the information ratio. In *COLT. (PMLR vol 134); arXiv:2009.12228 (2020)*, 2021.

- [23] Sham Kakade, Akshay Krishnamurthy, Kendall Lowrey, Motoya Ohnishi, and Wen Sun. Information theoretic regret bounds for online nonlinear control. In *Advances in Neural Information Processing Systems*, 2020. arXiv:2006.12466.
- [24] Jean Bretagnolle and Catherine Huber. Estimation des densités: Risque minimax. In C. Del-lacherie, P. A. Meyer, and M. Weil, editors, *Séminaire de Probabilités XII*, volume 649 of *Lecture Notes in Mathematics*, pages 342–363. Springer, 1978. doi: 10.1007/BFb0064609.
- [25] J Pearl. *Causality: Models, Reasoning, and Inference*. (2nd ed.). Cambridge., 2009.
- [26] Elias Bareinboim, Juan D. Correa, Duligur Ibeling, and Thomas Icard. On Pearl’s hierarchy and the foundations of causal inference. In Hector Geffner, Rina Dechter, and Joseph Y. Halpern, editors, *Probabilistic and Causal Inference: The Works of Judea Pearl*, pages 507–556. ACM Books, 2022. doi: 10.1145/3501714.3501743.
- [27] Omar Besbes, Yonatan Gur, and Assaf Zeevi. Stochastic Multi-Armed-Bandit Problem with Non-Stationary Rewards. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2014. URL <https://arxiv.org/abs/1405.3316>.
- [28] Y Yang, Z Zheng, M Tomizuka, C Liu, and S Li. The Feasibility Theory of Constrained Reinforcement Learning: A Tutorial Study. *arXiv:2404.10064*, 2024.
- [29] Oliver Schulte and Pascal Poupart. When should reinforcement learning use causal reasoning? *Transactions on Machine Learning Research*, 2025. Published 2025-05-09; OpenReview D1PPuk8ZBI; Survey Certification.
- [30] C Canonne. A short note on an inequality between KL and TV. *arXiv:2202.07198*, 2022.
- [31] Alexandre B. Tsybakov. *Introduction to Nonparametric Estimation*. Springer Series in Statistics. Springer, New York, 2009. doi: 10.1007/b13794.
- [32] T Lattimore and C Szepesvári. *Bandit Algorithms*. Cambridge., 2020.
- [33] H Hajiabolhassan and R Ortner. Online regret bounds for satisficing in Markov decision processes. *Mathematics of Operations Research*, 2025.
- [34] Y. Zhang, R Zhu, and Q Xie. On the peril of (even a little) non-stationarity in satisficing regret minimization. (*contemporaneous*; March 2026), 2026.
- [35] A Gerogiannis, Y.-H. Huang, and V Veeravalli. DARLING: Detection augmented reinforcement learning with non-stationary guarantees. (*contemporaneous*; April 2026), 2026.
- [36] S Kakade and J Langford. Approximately optimal approximate reinforcement learning (CPI / performance-difference lemma). In *ICML*, 2002.
- [37] R Munos. Error bounds for approximate policy iteration. In *ICML*, 2003.
- [38] S Ross and J. A Bagnell. Efficient reductions for imitation learning (Dagger). In *AISTATS*, 2010.
- [39] M. G Azar, I Osband, and R Munos. Minimax regret bounds for reinforcement learning (UCBVI). In *ICML*. *arXiv:1703.05449*, 2017.
- [40] J Schulman, S Levine, P Moritz, M. I Jordan, and P Abbeel. Trust region policy optimization (TRPO). In *ICML*. *arXiv:1502.05477*, 2015.
- [41] Brian D. O. Anderson. Adaptive systems, lack of persistency of excitation and bursting phenomena. *Automatica*, 21(3):247–258, 1985. doi: 10.1016/0005-1098(85)90058-5.
- [42] J. P Hespanha, D Liberzon, and A. R Teel. Lyapunov conditions for input-to-state stability of impulsive systems. *Automatica*. 44(11):2735–2744, 2008.
- [43] Jelle Bruineberg, Krzysztof Dolega, Joe Dewhurst, and Manuel Baltieri. The Emperor’s new Markov blankets. *Behavioral and Brain Sciences*, 45:e183, 2022. doi: 10.1017/S0140525X21002351.

- [44] Peter Auer, Nicolò Cesa-Bianchi, and Paul Fischer. Finite-time analysis of the multiarmed bandit problem. *Machine Learning*, 47(2–3):235–256, 2002. doi: 10.1023/A:1013689704352.
- [45] Hassan K. Khalil. *Nonlinear Systems*. Prentice-Hall, Upper Saddle River, NJ, 3rd edition, 2002.
- [46] Karl Friston, Thomas FitzGerald, Francesco Rigoli, Philipp Schwartenbeck, and Giovanni Pezzulo. Active inference: a process theory. *Neural Computation*, 29(1):1–49, 2017.
- [47] T Parr, G Pezzulo, and K. J Friston. *Active Inference: The Free Energy Principle in Mind, Brain, and Behavior*. MIT Press., 2022.
- [48] S Levine. Reinforcement learning and control as probabilistic inference. *arXiv:1805.00909*, 2018.
- [49] N Wiener. *Cybernetics*. MIT Press., 1948.
- [50] R Conant and W. R Ashby. Every good regulator of a system must be a model of that system. *Int. J. Systems Sci.*, 1970.
- [51] J Aczél and Z Daróczy. *On Measures of Information and Their Characterizations*. Academic Press., 1975.
- [52] P Auer, T Jaksch, and R Ortner. Near-optimal regret bounds for reinforcement learning (UCRL2). *J. Mach. Learn. Res.*, 11:1563–1600, 2010.

A Proof of the Composition Theorem

This appendix gives the full proof of Theorem 4.1. The proof composes the four key lemmas of Section 5 (full proofs of those in Appendix B) with a block decomposition at optimum-change events plus Cauchy–Schwarz across blocks.

Notation Throughout this section: $K_t(s) := D_{\text{KL}}(\delta_{a_t^*(s)} \parallel Q_t(\cdot \mid s))$ is the per-state reverse-KL coordinate; $d_h^{Q_t}(\cdot \mid s_0)$ is the round- t learner-induced state distribution at horizon step h starting from s_0 ; $\overline{\text{TV}}_t := \frac{1}{N_h} \sum_{h=0}^{N_h-1} \mathbb{E}_{s_h \sim d_h^{Q_t}} [\text{TV}(\pi_t^*(\cdot \mid s_h), Q_t(\cdot \mid s_h))]$ is the occupancy-weighted per-round coordinate of (A5); $B_T = |\{t : (P_t, r_t) \neq (P_{t-1}, r_{t-1})\}|$ counts piecewise-stationary segment boundaries (per Section 3); the partition is $0 = \tau_0 < \tau_1 < \dots < \tau_{B_T} \leq \tau_{B_T+1} = T$ with block i length $\Delta_i := \tau_{i+1} - \tau_i$.

A.1 Proof of (i) — Per-round identity-bounded regret

For deterministic $\pi_t^* = \delta_{a_t^*(s)}$ and any $Q_t(\cdot \mid s)$, Key Lemma 1 (Lemma 5.1, full proof in Appendix B.1) gives the exact identity $\text{TV}(\delta_{a_t^*(s)}, Q_t(\cdot \mid s)) = 1 - e^{-K_t(s)}$ in the extended real sense (when $Q_t(a_t^*(s) \mid s) = 0$, $K_t(s) = +\infty$ and both sides equal 1, recovering the trivial bound). Composed with the per-state TV-regret bounds (full proofs in Appendix B.1 and Appendix C.4): $\Delta_{\min} \cdot \text{TV}(\delta_{a_t^*(s)}, Q_t(\cdot \mid s)) \leq R(Q_t \mid s) \leq V_{\max} \cdot \text{TV}(\delta_{a_t^*(s)}, Q_t(\cdot \mid s))$, substituting the identity yields conclusion (i): $\Delta_{\min}(1 - e^{-K_t(s)}) \leq R(Q_t \mid s) \leq V_{\max}(1 - e^{-K_t(s)})$. $\square$

A.2 Proof of (ii) — Aggregate mismatch ultimate boundedness

Under (A2), Key Lemma 2 (Lemma 5.2, full proof in Appendix B.3) applied to the diagonal sector model with $\alpha_{ij}^{\text{ss}} = \nu_{ij} \cdot \iota_{ij} \cdot (1 - \lambda_{ij})$ gives ultimate boundedness of $\|\delta_\Sigma\|$ within $R_\Sigma^* = \rho_\Sigma / \mathcal{T}_\Sigma^{\text{bn,ss}}$. $\square$

A.3 Proof of (iii) — KL coordinate bias bound

Key Lemma 4 (Lemma 5.4, full proof in Appendix B.6) gives the per-state bias bound directly, under (A3)’s identifiability assumption (which gives $p_{\text{id}}(s)$ for each s) and the two-point support condition $Q_t(\cdot \mid s) \geq q_0$ at $a_{\text{ag},t}^*(s)$ and $\tilde{a}_t^*(s)$. $\square$

546 A.4 Proof of (iv) — 2×2 corrective-action routing

547 (A4) states the agent applies the satisfaction-gap / control-regret 2×2 (Section 4.1 Component 1).
 548 Conclusion (iv) is a direct restatement of the 2×2 cell partition. The decomposition is convention-
 549 dependent (the values of δ_{sat} and δ_{regret} depend on the continuation convention; see Appendix C.1 for
 550 monotonicity across C1/C2/C3); the 2×2 structure is preserved across all three conventions. $\square$

551 A.5 Proof of (v) — Trajectory-level cumulative regret

552 We chain Key Lemmas 1 and 4 with the simulation lemma and a Cauchy–Schwarz block argument.
 553 The proof is in four steps; the high-level sketch is in Section 5.5.

554 *Step 1: Simulation / performance-difference lemma.* For any two policies π
 555 and π' , any state s , and finite-horizon non-discounted MDP with horizon N_h ,
 556 the standard performance-difference lemma^[36–39] gives $V_t^{\pi'}(s) - V_t^\pi(s) =$
 557 $\mathbb{E}\left[\sum_{h=0}^{N_h-1} \left(\mathbb{E}_{a' \sim \pi'(\cdot | s_h)}[Q_h^{\pi'}(s_h, a')] - Q_h^\pi(s_h, a_h)\right) \mid s_0 = s, a_h \sim \pi\right]$. Each per-step bracket
 558 is bounded by value range times TV: $|\mathbb{E}_{a' \sim \pi'}[Q_h^{\pi'}] - \mathbb{E}_{a \sim \pi}[Q_h^\pi]| \leq V_{\max} \cdot \text{TV}(\pi'(\cdot | s_h), \pi(\cdot | s_h))$. Applied at round t with $\pi' = \pi_t^*$, $\pi = Q_t$: $V_t^{\pi_t^*}(s_0) - V_t^{Q_t}(s_0) \leq$
 559 $V_{\max} \sum_{h=0}^{N_h-1} \mathbb{E}_{s_h \sim d_h^{Q_t}(\cdot | s_0)}[\text{TV}(\pi_t^*(\cdot | s_h), Q_t(\cdot | s_h))] = V_{\max} \cdot N_h \cdot \overline{\text{TV}}_t$. The horizon factor N_h
 560 here is the *linear-in- N_h* simulation-lemma penalty — sharper than the N_h^2 of TRPO-style worst-state
 561 bounds^[40], which bound max-state TV rather than occupancy-weighted TV.

563 *Step 2: Per-state TV identity along the trajectory.* By the deterministic-per-state optimum scope of
 564 Theorem 4.1 ($\pi_t^*(\cdot | s) = \delta_{a_t^*(s)}$), at every visited state Key Lemma 1 gives $\text{TV}(\pi_t^*(\cdot | s), Q_t(\cdot | s)) = 1 - Q_t(a_t^*(s) | s) = 1 - e^{-K_t(s)}$. Summing over the trajectory and taking expectations:
 565 $\overline{\text{TV}}_t = \frac{1}{N_h} \sum_{h=0}^{N_h-1} \mathbb{E}_{s_h \sim d_h^{Q_t}(\cdot | s_0)}[1 - e^{-K_t(s_h)}]$. The per-state KL/TV identity is preserved underneath
 566 the occupancy aggregation; the per-round coordinate is sharper than Pinsker/BH at every visited state.

568 *Step 3: Block-sum + Cauchy–Schwarz across blocks.* Within each block $[\tau_i, \tau_{i+1})$ of length
 569 Δ_i , (A5)’s restart-on-change identity-tight base learner gives $\sum_{t \in \text{block}_i} \mathbb{E}[\overline{\text{TV}}_t] \leq 2c\sqrt{\Delta_i}$. Sum-
 570 ming over the $B_T + 1$ blocks: $\sum_{t=1}^T \mathbb{E}[\overline{\text{TV}}_t] = \sum_{i=0}^{B_T} \sum_{t \in \text{block}_i} \mathbb{E}[\overline{\text{TV}}_t] \leq 2c \sum_{i=0}^{B_T} \sqrt{\Delta_i}$.
 571 Cauchy–Schwarz gives $\sum_{i=0}^{B_T} \sqrt{\Delta_i} \leq \sqrt{(B_T + 1) \cdot \sum_i \Delta_i} = \sqrt{(B_T + 1)T}$. Multiplying by
 572 $V_{\max} N_h$: $\sum_{t=1}^T \mathbb{E}[V_{\max} N_h \overline{\text{TV}}_t] \leq 2c V_{\max} N_h \sqrt{(B_T + 1)T}$. This is the rate term.

573 *Step 4: Misidentification penalty via simulation lemma at the identified-vs-true-optimum gap.* Steps
 574 1–3 implicitly compared the agent’s policy Q_t to the *agent’s identified-optimum policy* $\pi_{\text{ag},t}^*(\cdot | s)$
 575 $:= \delta_{a_{\text{ag},t}^*(s)}$ (since (A5)’s base-learner guarantee is what the agent achieves on its identified
 576 objective). The trajectory-level dynamic regret in (v) compares to the *true* optimum $\tilde{\pi}_t^*(\cdot | s) := \delta_{\tilde{a}_t^*(s)}$.
 577 Decompose: $V_t^{\tilde{\pi}_t^*}(s_0) - V_t^{Q_t}(s_0) = [V_t^{\tilde{\pi}_t^*}(s_0) - V_t^{\pi_{\text{ag},t}^*}(s_0)] + [V_t^{\pi_{\text{ag},t}^*}(s_0) - V_t^{Q_t}(s_0)]$. Steps
 578 1–3 bound the second bracket. The first bracket is the misidentification penalty. Apply the simulation
 579 lemma to the two point-mass policies $(\tilde{\pi}_t^*, \pi_{\text{ag},t}^*)$: each per-step bracket reduces to $Q_h^{\tilde{\pi}_t^*}(s_h, \tilde{a}_t^*(s_h)) -$
 580 $Q_h^{\pi_{\text{ag},t}^*}(s_h, a_{\text{ag},t}^*(s_h)) \in [0, V_{\max}] \cdot \mathbf{1}[\tilde{a}_t^*(s_h) \neq a_{\text{ag},t}^*(s_h)]$. Lower bound by optimality of $\tilde{a}^*$; upper
 581 bound by the cumulative-Q range with the indicator on the misidentification event. Taking expectations,
 582 with the per-state floor $p_{\text{id}} := \min_s p_{\text{id}}(s)$: $\mathbb{E}[V_t^{\tilde{\pi}_t^*}(s_0) - V_t^{\pi_{\text{ag},t}^*}(s_0)] \leq V_{\max} \cdot N_h \cdot (1 - p_{\text{id}})$. Summing
 583 over T rounds: $V_{\max} \cdot N_h \cdot (1 - p_{\text{id}}) \cdot T$. This is the bias term, in *value coordinates* throughout — no
 584 detour through the KL coordinate. Key Lemma 4’s KL-readout bound is the right tool for conclusion
 585 (iii)’s diagnostic computability; for the value-coordinate misidentification penalty here, the indicator
 586 decomposition is direct and tighter (the support condition $Q_t \geq q_0$ at the two argmax candidates is
 587 *not* required for this term — only conclusion (iii) needs it).

588 *Combining.* Adding Steps 1–4: $\mathbb{E}[\text{DynReg}(T)] \leq 2c V_{\max} N_h \sqrt{(B_T + 1)T} + V_{\max} N_h (1 - p_{\text{id}}) T$,
 589 which is conclusion (v). $\square$

590 **Cesàro tracking corollary** Dividing through by T : $\frac{1}{T} \sum_t (V_t^{\tilde{\pi}_t^*} - V_t^{Q_t}) \leq$
 591 $2c V_{\max} N_h \sqrt{(B_T + 1)/T} + V_{\max} N_h (1 - p_{\text{id}})$. As $T \rightarrow \infty$ with $B_T = o(T)$ and $p_{\text{id}} \rightarrow 1$

(Regime A), the right side tends to zero — the agent’s average value over the trajectory tracks the moving optimum’s average value. In Regime B the second term is constant; tracking is only up to a $V_{\max}(1 - p_{\text{id}})$ residual. $\square$

Bandit case sharpening (Remark) In the bandit special case ($N_h = 1$, so $\overline{\text{TV}}_t = \mathbb{E}[1 - e^{-K_t}]$), Thompson sampling and UCB (Appendix D) satisfy (A5) with the *stronger logarithmic* per-block rate. By Lattimore and Szepesvári 2020^[32] Theorem 7.1 (or⁴⁴), $\mathbb{E}[N_a(L)] = O(\log L / \Delta_a^2)$ per suboptimal arm a ; aggregating over arms gives $\mathbb{E}[1 - e^{-K_t}] = \mathbb{E}[1 - Q_t(a^*)] = O(\log t / (t \Delta_{\min}^2))$ per stationary block. Substituting and summing gives cumulative regret $O(V_{\max}(B_T + 1) \log(T / (B_T + 1)) / \Delta_{\min}^2)$ — sharper than the worst-case $\sqrt{(B_T + 1)T}$ when $\Delta_{\min}$ is bounded away from zero. The framework’s $V_{\max} \cdot \text{TV}$ chain is structurally one factor of $\Delta_{\min}$ looser than direct gap-aware UCB analysis (which gives $O(\log T / \Delta_{\min})$ cumulative regret per block by weighting each suboptimal pull by $\Delta(a)$ rather than $V_{\max}$); this is a feature of the unification, not a bug. For $N_h > 1$ MDPs, UCRL2 and UCBVI provide the natural occupancy-weighted instantiation (Appendix D) with c of order $\tilde{O}(\sqrt{N_h SA})$, lifting cleanly to $\tilde{O}(N_h^2 \sqrt{SA(B_T + 1)T})$ — same piecewise-stationary B_T family as restart-on-change variants of^[1].

Best-of-Both-Worlds non-stationarity extension (Remark on (A5’)) Conclusion (v) above uses the restart-on-change form of (A5). The per-round identity route is *stationarity-agnostic* — Steps 1, 2, 4 above hold per round regardless of cross-round dynamics; only Step 3’s block-Cauchy–Schwarz aggregator is piecewise-stationary-specific. Replacing (A5) by

(A5’) *MASTER-wrapped base learner*. The base learner is the Wei and Luo 2021^[2] MASTER black-box wrapping of any base algorithm satisfying their Assumption 1 (a near-stationary regret guarantee with $C(t) = c_1 t^p + c_2$ for $p \in [1/2, 1)$, plus an auxiliary-quantity output condition; UCRL2, UCBVI, and Thompson sampling all qualify), giving stationary regret $\sum_{t \in W} \mathbb{E}[\overline{\text{TV}}_t] \leq 2c\sqrt{|W|}$ within each MASTER-selected window W .

substitutes Step 3’s block-Cauchy–Schwarz aggregator with MASTER’s adaptive-window aggregator (parallel base-learner instances at exponentially-spaced window sizes $W_k = 2^k$, online meta-learner over the $\log T$ instances). The wrapping is mechanical: the per-round identity (Steps 1–2) is preserved unchanged at each round; the misidentification penalty (Step 4) is round-additive and rides through; only the aggregator changes. The resulting rate, by Wei-Luo Theorem 2 applied with our framework’s $C(t) = 2cV_{\max}N_h\sqrt{t}$ at the boundary $p = 1/2$, is $\mathbb{E}[\text{DynReg}(T)] \leq \tilde{O}\left(V_{\max}N_h \cdot \min\left\{\sqrt{(B_T + 1)T}, (V_T + 1)^{1/3}T^{2/3}\right\}\right) + V_{\max}N_h(1 - p_{\text{id}})T$, automatically — without prior knowledge of B_T or V_T . The $V_T^{1/3}T^{2/3}$ exponent matches^[3]’s near-optimal RestartQ-UCB rate and is information-theoretically near-optimal: the^[27] lower bound $\Omega(K^{1/3}V_T^{1/3}T^{2/3})$ for non-stationary stochastic bandits applies *also at the deterministic- π corner** (their construction is a Bernoulli MAB with dynamic oracle playing the per-round argmax — deterministic-per-round optimum, identical to our canonical scope), so the per-round identity sharpens *constants and per-round form*, not the V_T exponent. The two regimes are *not* structurally distinct: they share the per-round coordinate $1 - e^{-K_t}$ and differ only in the aggregation strategy. $\square$

A.6 Component-by-component failure-mode analysis

This composition is a *bundle of compatible guarantees* (each conclusion supported by its own assumptions), not a single integrated theorem in which every component is necessary for the rate. The rate (v) goes through (A5) and (i) alone; the joint properties — rate plus persistence plus learnability plus corrective-action routing — require all four. The failure modes component-by-component:

Without Component 2 (point-mass identity) The per-round bound is Pinsker $V_{\max}\sqrt{D_{\text{KL}}/2}$ or BH $V_{\max}\sqrt{1 - e^{-D_{\text{KL}}}}$ — both strictly looser at this corner; Pinsker is vacuous for $D_{\text{KL}} > 2$. The metric layer loses *exactness*, and the “behavior-cloning loss against optimal trajectory” interpretation of the cumulative coordinate vanishes.

640 **Without Component 3 (strategic tempo)** The agent has no persistence guarantee. By the structural-
641 class theorem on $\mathcal{A}_{\text{decay}}$ (Appendix C.6), every count-accumulating update without forgetting even-
642 tually violates the threshold for any positive disturbance; the modeled mismatch $\|\delta_\Sigma\|$ grows un-
643 boundedly. The per-round bound (i) still holds, but $K_t(s)$ may itself drift — there is no guarantee the
644 per-round coordinate stays bounded.

645 **Without Component 4 (loop-Level-2)** The bound’s RHS is computable from Q_t , but its interpreta-
646 tion as regret-against-true-optimum requires identifying $a_t^*(s)$ — an interventional query under the
647 environment’s transition kernel. In Regime C ($\iota \approx 0$), the misidentification penalty saturates at $V_{\max}$
648 per state per step, contributing $V_{\max} N_h \cdot T$ to cumulative regret — the rate (v) becomes vacuous.

649 **Without Component 1 (two-gap diagnostic)** The regret signal δ_{regret} alone does not distinguish
650 “policy revision is the right move” from “the goal is unattainable from $\bar{M}_t$.” An agent might persist in
651 policy revision when the model, policy class, or horizon is the actual problem (Capability-limit cell),
652 or revise the goal when policy revision would have sufficed.

653 So while the rate (v) goes through (A5) and (i) alone, the *bundle* — rate plus persistence plus
654 learnability plus corrective-action routing — requires all four. Each drop is a concrete loss of an
655 epistemic property that the field treats as separate.

656 B Proofs of the Key Lemmas

657 This appendix gives full proofs of the four key lemmas surfaced in Section 5, plus the perturbative
658 extension and the ProST sector-level reduction.

659 B.1 Proof of Key Lemma 1 (Point-mass reverse-KL/TV identity)

660 For deterministic $\pi^* = \delta_{a^*}$ and any policy Q (extending $D_{\text{KL}} = +\infty$ when $Q(a^*) = 0$, which gives
661 $1 - e^{-D_{\text{KL}}} = 1 = \text{TV}$ trivially):

662 The reverse-KL collapses under the point-mass: $D_{\text{KL}}(\delta_{a^*} \| Q) = \sum_a \delta_{a^*}(a) \log \frac{\delta_{a^*}(a)}{Q(a)} =$
663 $\log \frac{1}{Q(a^*)} = -\log Q(a^*)$, where the convention $0 \log 0 = 0$ disposes of the $a \neq a^*$ terms.
664 Combined with the textbook total-variation calculation $\text{TV}(\delta_{a^*}, Q) = \frac{1}{2} \sum_a |\delta_{a^*}(a) - Q(a)| =$
665 $\frac{1}{2} (|1 - Q(a^*)| + \sum_{a \neq a^*} Q(a)) = \frac{1}{2} ((1 - Q(a^*)) + (1 - Q(a^*))) = 1 - Q(a^*)$, we have
666 $-\log Q(a^*) = -\log(1 - \text{TV})$. Solving for TV gives $\text{TV} = 1 - e^{-D_{\text{KL}}}$. $\square$

667 **Strict-improvement substitution** Substituting the identity into the Bretagnolle–Huber inequality
668 $\text{TV}(P, Q) \leq \sqrt{1 - e^{-D_{\text{KL}}(P \| Q)}}$ at $P = \delta_{a^*}$ yields $1 - e^{-D_{\text{KL}}} \leq \sqrt{1 - e^{-D_{\text{KL}}}}$, which holds *strictly*
669 on $D_{\text{KL}} > 0$ since $x < \sqrt{x}$ on $(0, 1)$. The identity therefore supplies the *exact* TV value at the
670 deterministic- π^* corner and lies strictly below the BH envelope there.

671 **Two-sided regret bound** (stated as Component 2’s two-sided characterization in Section 4.1 and as Key
672 Lemma 1’s *Per-round regret bound* paragraph in Section 5.1). From the regret expression $R(Q) =$
673 $\sum_{a \neq a^*} Q(a) \Delta(a)$ and $\Delta(a) \leq V_{\max}$: $R(Q) \leq V_{\max} \sum_{a \neq a^*} Q(a) = V_{\max} (1 - Q(a^*)) =$
674 $V_{\max} \text{TV}(\delta_{a^*}, Q) = V_{\max} (1 - e^{-D_{\text{KL}}})$. The matching action-gap lower bound (full statement in
675 Appendix C.4) gives $R(Q) = \sum_{a \neq a^*} Q(a) \Delta(a) \geq \Delta_{\min} \sum_{a \neq a^*} Q(a) = \Delta_{\min} \text{TV}(\delta_{a^*}, Q) =$
676 $\Delta_{\min} (1 - e^{-D_{\text{KL}}})$. Combining gives the two-sided characterization. The bound is *Lipschitz-equivalent*
677 with constants $\Delta_{\min}/V_{\max}$ and 1, both achieved on extremal value landscapes — therefore *coordinate-*
678 *optimal among bounds depending only on TV*. $\square$

679 B.2 Proof of the perturbative extension — ϵ -stochastic and softmax-regularized

680 We establish the perturbative identity via a uniform perturbation argument rather than an alignment-
681 specific calculation.

682 **ϵ -greedy stochastic optimum** Let $\pi_\epsilon^*(a^*) = 1 - \epsilon$ and $\pi_\epsilon^*(a) = \epsilon/(|\mathcal{A}| - 1)$ for $a \neq a^*$, and
683 assume the full-support lower bound $Q(a) \geq q_0 > 0$ for all $a \in \mathcal{A}$. (Full support is required since π_ϵ^*
684 has positive mass on every action; without $Q(a) > 0$ on the full support, $D_{\text{KL}}(\pi_\epsilon^* \| Q) = +\infty$.)

685 *Step 1 — TV is Lipschitz under the point-mass perturbation.* $\text{TV}(\pi_\epsilon^*, \delta_{a^*}) = \epsilon$. By the triangle
686 inequality for total variation, $|\text{TV}(\pi_\epsilon^*, Q) - \text{TV}(\delta_{a^*}, Q)| \leq \text{TV}(\pi_\epsilon^*, \delta_{a^*}) = \epsilon$, uniformly in Q —
687 no alignment hypothesis needed.

688 *Step 2 — Reverse-KL admits a uniform expansion.* Direct computation:

$$D_{\text{KL}}(\pi_\epsilon^* \| Q) - D_{\text{KL}}(\delta_{a^*} \| Q) = (1 - \epsilon) \log \frac{1 - \epsilon}{Q(a^*)} + \frac{\epsilon}{|\mathcal{A}| - 1} \sum_{a \neq a^*} \log \frac{\epsilon / (|\mathcal{A}| - 1)}{Q(a)} - (-\log Q(a^*))$$

689
$$= (1 - \epsilon) \log(1 - \epsilon) + \epsilon \log Q(a^*) + \epsilon \log \frac{\epsilon}{|\mathcal{A}| - 1} - \frac{\epsilon}{|\mathcal{A}| - 1} \sum_{a \neq a^*} \log Q(a),$$

690 where the second equality collects the $\log Q(a^*)$ contributions via $-(1 - \epsilon) \log Q(a^*) + \log Q(a^*) =$
691 $\epsilon \log Q(a^*)$ and splits the constant $\log$ term off the conditional sum. Expanding
692 $(1 - \epsilon) \log(1 - \epsilon) = -\epsilon + O(\epsilon^2)$ and $\epsilon \log(\epsilon / (|\mathcal{A}| - 1)) = \epsilon \log \epsilon - \epsilon \log(|\mathcal{A}| - 1)$,
693 the leading-order behavior is $D_{\text{KL}}(\pi_\epsilon^* \| Q) = D_{\text{KL}}(\delta_{a^*} \| Q) + \epsilon \log \epsilon - \epsilon (\log(|\mathcal{A}| -$
694 $1) + 1) + \epsilon \log Q(a^*) - \frac{\epsilon}{|\mathcal{A}| - 1} \sum_{a \neq a^*} \log Q(a) + O(\epsilon^2)$. Under $Q(a) \geq q_0$, each
695 $|\log Q(a)| \leq \log(1/q_0)$ is uniformly bounded; the dominant correction is the $\epsilon \log \epsilon$ leading
696 term, giving $|D_{\text{KL}}(\pi_\epsilon^* \| Q) - D_{\text{KL}}(\delta_{a^*} \| Q)| = O(\epsilon \log(1/\epsilon) + \epsilon \log(1/q_0)) = O(\epsilon \log(1/\epsilon))$
697 absorbing $\epsilon \log(1/q_0)$ into the leading $\epsilon \log(1/\epsilon)$ term as $\epsilon \rightarrow 0$ (constants depend on q_0 and $|\mathcal{A}|$).

698 *Step 3 — Map through $1 - e^{-x}$ via Lipschitzness.* The map $x \mapsto 1 - e^{-x}$ is 1-Lipschitz on $[0, \infty)$,
699 so $|(1 - e^{-D_{\text{KL}}(\pi_\epsilon^* \| Q)}) - (1 - e^{-D_{\text{KL}}(\delta_{a^*} \| Q)})| = O(\epsilon \log(1/\epsilon))$. Combining Steps 1 and 3 with
700 the unperturbed identity $\text{TV}(\delta_{a^*}, Q) = 1 - e^{-D_{\text{KL}}(\delta_{a^*} \| Q)}$: $\text{TV}(\pi_\epsilon^*, Q) = 1 - e^{-D_{\text{KL}}(\pi_\epsilon^* \| Q)} +$
701 $O(\epsilon \log(1/\epsilon))$ uniformly in Q over the class $\{Q : Q(a) \geq q_0\}$, with constants depending on q_0 and
702 $|\mathcal{A}|$. This establishes the ϵ -greedy form of the perturbative identity. $\square$

703 **Softmax-regularized optimum** Let $\pi_\tau^*(a) \propto \exp(Q_O(a)/\tau)$. Under isolated optimum with
704 action gap $\Delta_{\min}$, the off-optimum mass concentrates exponentially in $1/\tau$: $1 - \pi_\tau^*(a^*) =$
705 $\frac{\sum_{a \neq a^*} \exp(-(Q_O(a^*) - Q_O(a))/\tau)}{1 + \sum_{a \neq a^*} \exp(-(Q_O(a^*) - Q_O(a))/\tau)} \leq (|\mathcal{A}| - 1) \exp(-\Delta_{\min}/\tau)$. Treating π_τ^* as π_ϵ^* with ef-
706 fective $\epsilon = O(\exp(-\Delta_{\min}/\tau))$, the ϵ -greedy expansion above applies with $\epsilon \log(1/\epsilon) =$
707 $O((\Delta_{\min}/\tau) \exp(-\Delta_{\min}/\tau)) = O(\tau^{-1} \exp(-\Delta_{\min}/\tau))$, exponentially small with a polynomial
708 prefactor (equivalently, $O(\exp(-c\Delta_{\min}/\tau))$ for any fixed $c \in (0, 1)$ since the exponential dominates
709 the polynomial). $\square$

710 **Two-sided regret bound under perturbation** Composing with the TV-regret bounds: $\Delta_{\min}(1 -$
711 $e^{-D_{\text{KL}}(\pi_\epsilon^* \| Q)}) - O(\epsilon \log(1/\epsilon)) \leq R(Q) \leq V_{\max}(1 - e^{-D_{\text{KL}}(\pi_\epsilon^* \| Q)}) + O(\epsilon \log(1/\epsilon))$.

712 B.3 Proof of Key Lemma 2 (Multi-factor forgetting prerequisite)

713 We work in the continuous-time idealization of Model (Σ) : $\dot{\delta}_\Sigma = -\mathbf{F}(\delta_\Sigma) + \mathbf{w}$ with $\mathbf{F}(\delta_\Sigma) :=$
714 $(\alpha_{ij}^{\text{ss}} \delta_{ij})_{(i,j) \in E}$ the sector correction and $\mathbf{w}$ the deterministic disturbance with $\|\mathbf{w}\| \leq \rho_\Sigma$. The
715 discrete-time form recovers the same scaling up to absorbable constants via an AM-GM expansion of
716 $\|\delta_\Sigma - \mathbf{F} + \mathbf{w}\|^2$ (sketched below).

717 For the Lyapunov function $V(\delta_\Sigma) := \|\delta_\Sigma\|^2$, $\dot{V} = 2\delta_\Sigma^\top \dot{\delta}_\Sigma = -2\delta_\Sigma^\top \mathbf{F}(\delta_\Sigma) + 2\delta_\Sigma^\top \mathbf{w}$. The
718 first term is the *sector product*, bounded below by the bottleneck: $\delta_\Sigma^\top \mathbf{F}(\delta_\Sigma) = \sum_{(i,j)} \alpha_{ij}^{\text{ss}} \delta_{ij}^2 \geq$
719 $\min_{(i,j)} \alpha_{ij}^{\text{ss}} \cdot \|\delta_\Sigma\|^2 = \mathcal{T}_\Sigma^{\text{bn,ss}} \cdot \|\delta_\Sigma\|^2$, with the bound tight: setting δ_Σ proportional to $\mathbf{e}_{(i^*, j^*)}$ where
720 $(i^*, j^*) \in \arg \min_{(i,j)} \alpha_{ij}^{\text{ss}}$ saturates it. The cross-term is bounded by Cauchy-Schwarz: $|2\delta_\Sigma^\top \mathbf{w}| \leq$
721 $2\|\delta_\Sigma\| \cdot \|\mathbf{w}\| \leq 2\rho_\Sigma \|\delta_\Sigma\|$. Combining: $\dot{V} \leq -2\mathcal{T}_\Sigma^{\text{bn,ss}} \|\delta_\Sigma\|^2 + 2\rho_\Sigma \|\delta_\Sigma\| = -2\|\delta_\Sigma\| \cdot$
722 $(\mathcal{T}_\Sigma^{\text{bn,ss}} \|\delta_\Sigma\| - \rho_\Sigma)$. The drift is strictly negative whenever $\|\delta_\Sigma\| > \rho_\Sigma / \mathcal{T}_\Sigma^{\text{bn,ss}}$; the trajectory is
723 ultimately bounded by $R_\Sigma^* = \rho_\Sigma / \mathcal{T}_\Sigma^{\text{bn,ss}}$. So whenever $\mathcal{T}_\Sigma^{\text{bn,ss}} > \rho_\Sigma / R_\Sigma$, we have $R_\Sigma^* \leq R_\Sigma$ and
724 the modeled mismatch is ultimately bounded within the strategic reserve. This is the standard
725 Khalil sector-Lyapunov ultimate-boundedness setup with the cross-term retained^[45] (Section 9.2;
726 comparison-principle form). $\square$

727 *Discrete-time form.* For the discrete update $\delta_\Sigma^+ = \delta_\Sigma - \mathbf{F}(\delta_\Sigma) + \mathbf{w}$ with $\alpha_{ij}^{ss} \in [0, 1]$ and AM–GM
 728 with parameter $\eta = \mathcal{T}_\Sigma^{\text{bn,ss}} / (2(1 - \mathcal{T}_\Sigma^{\text{bn,ss}}))$, the steady-state bound is $\|\delta_\Sigma^*\| \leq 2\rho_\Sigma / \mathcal{T}_\Sigma^{\text{bn,ss}}$ — same
 729 scaling, factor-of-2 absorbable into the threshold constant.

730 *Mean-square corollary.* If instead $\mathbf{w}$ is interpreted as zero-mean stochastic disturbance with $\mathbb{E}[\|\mathbf{w}\|^2] \leq$
 731 ρ_Σ^2 , the cross-term $\mathbb{E}[2\delta_\Sigma^\top \mathbf{w} \mid \delta_\Sigma] = 0$ vanishes in expectation, and the drift inequality gives mean-
 732 square ultimate boundedness $\sqrt{\mathbb{E}[V^*]} \leq \rho_\Sigma / \sqrt{2\mathcal{T}_\Sigma^{\text{bn,ss}}}$ under threshold $\mathcal{T}_\Sigma^{\text{bn,ss}} > \rho_\Sigma^2 / (2R_\Sigma^2)$. The
 733 deterministic-disturbance form above is the operative one for Theorem 4.1(A2) since non-stationarity-
 734 induced drift is structurally adversarial-aligned (consistently moving mismatch in one direction) rather
 735 than zero-mean.

736 **Sharpness** When the inequality reverses ($\mathcal{T}_\Sigma^{\text{bn,ss}} \leq \rho_\Sigma / R_\Sigma$), the adversarial disturbance $\mathbf{w}^*(t) =$
 737 $\rho_\Sigma \delta_\Sigma(t) / \|\delta_\Sigma(t)\|$ saturates Cauchy–Schwarz and gives $d\|\delta_\Sigma\|/dt \geq -\mathcal{T}_\Sigma^{\text{bn,ss}} \|\delta_\Sigma\| + \rho_\Sigma$. By the
 738 comparison principle, $\|\delta_\Sigma(t)\|$ exceeds any $R < \rho_\Sigma / \mathcal{T}_\Sigma^{\text{bn,ss}}$ in finite time. The bottleneck-element
 739 refinement: pick $\mathbf{w}$ and δ_Σ both supported on $(i^*, j^*) \in \arg \min_{(i,j)} \alpha_{ij}^{ss}$ with $|w_{i^*j^*}| = \rho_\Sigma$ —
 740 saturates the sector inequality on the contracting side and aligns with $\mathbf{w}$ on the disturbance side
 741 simultaneously. The threshold is *sharp inside the diagonal sector model*. The theorem is silent about
 742 non-diagonal correction architectures or stabilization mechanisms outside Model (Σ). $\square$

743 B.4 Proof of the impulsive ProST sector-level reduction

744 Within Model (Σ), idealize ProST^[8] as an impulsive system: between scheduled update times
 745 $\{t_1, \dots, t_K\} \subset [0, T]$ the agent’s policy is held fixed and the strategic mismatch evolves under
 746 disturbance budget ρ_Σ alone (continuous-destabilizing in the Lyapunov $V = \|\delta_\Sigma\|^2$); at each update
 747 time t_i the policy is revised, contracting the modeled mismatch by per-update impulse gain $\gamma \in (0, 1]$
 748 so $\|\delta_\Sigma(t_i^+)\| \leq (1 - \gamma)\|\delta_\Sigma(t_i^-)\|$. Then $V(t_i^+) \leq (1 - \gamma)^2 V(t_i^-)$.

749 This fits the Hespanha–Liberzon–Teel^[42] (Theorem 1) framework for impulsive systems with desta-
 750 bilizing continuous evolution and stabilizing impulses (the *reverse-ADT* branch). The continuous-
 751 evolution rate is $\lambda_c = \rho_\Sigma / R_\Sigma$ (from linearizing the disturbance contribution to $\dot{V} \approx +\lambda_c V + \gamma_c(\rho_\Sigma)$).
 752 The impulse contraction in V is $\mu = (1 - \gamma)^2$. The reverse-ADT condition for ultimate boundedness —
 753 impulses cannot be too sparse — gives $\Delta_{\max} \cdot \lambda_c < -\ln \mu \iff \Delta_{\max} \cdot \rho_\Sigma / R_\Sigma < -\ln(1 - \gamma)^2$,
 754 where $\Delta_{\max} := \sup_i (t_{i+1} - t_i)$ is the longest inter-update gap. Under the condition the modeled
 755 mismatch is ultimately bounded within $R_\Sigma^* \approx \rho_\Sigma \Delta_{\max} / \gamma$ (leading order in γ).

756 Under uniform blocks $\Delta_i = T/K$ the condition reduces to $K/T > \rho_\Sigma / (2\gamma R_\Sigma)$ to leading order in γ
 757 — recovering ProST’s K/T form with the impulse gain made explicit. Under nonuniform schedules
 758 the impulsive form is *strictly stronger*: the longest block sets the threshold, not the average. The
 759 lemma is rigorous for the modeled impulsive sector dynamics; the per-update impulse gain γ is a
 760 domain quantity (depends on within-block sample size and update-time discount) and is left abstract.
 761 $\square$

762 B.5 Proof of Key Lemma 3 (Loop generates interventional samples)

763 By temporal ordering and the singular-trajectory commitment of Section 3, a_t causally precedes
 764 o_{t+1} . We invoke Pearl’s back-door criterion^[25] (Theorem 3.3.2): $P(o_{t+1} \mid \text{do}(a_t), H_t) = P(o_{t+1} \mid$
 765 $a_t, H_t)$ whenever $(o_{t+1} \perp\!\!\!\perp a_t \mid H_t)_{G_{\overline{a_t}}}$, where $G_{\overline{a_t}}$ denotes the graph G with all incoming arrows
 766 to a_t removed. In the singular-trajectory + agent-as-policy graph of Section 3, the only arrows into a_t
 767 come from H_t (the agent’s policy is $\pi_t(a_t \mid H_t)$); deleting them severs the only confounding path
 768 from a_t to o_{t+1} . Whether H_t d-separates a_t from o_{t+1} in $G_{\overline{a_t}}$ is precisely the content of (C2). The
 769 identity $P(o_{t+1} \mid \text{do}(a_t), H_t) = P(o_{t+1} \mid a_t, H_t)$ holds on the support where $\pi_t(a_t \mid H_t) > 0$,
 770 which is automatic since a_t was drawn from π_t . $\square$

771 *Remark on equivalent forms of (C2).* The d-separation condition $(o_{t+1} \perp\!\!\!\perp a_t \mid H_t)_{G_{\overline{a_t}}}$ is equivalent
 772 to the potential-outcome form $a_t \perp\!\!\!\perp Y^{(a_t)} \mid H_t$ (sequential ignorability in the Robins/Murphy sense),
 773 and to the truncated-factorization condition $P(\mathbf{U} \mid a_t, H_t) = P(\mathbf{U} \mid H_t)$ where $\mathbf{U}$ are the latent
 774 variables jointly affecting a_t and o_{t+1} — applied to the *prior* on $\mathbf{U}$ in the do-marginalization, not to
 775 the conditional on o_{t+1} . The three forms are interchangeable; the d-separation form is the most direct
 776 handle for the proof above.

On the role of (C1) and (C3). The identification proof above is load-bearing on (C2) alone: (C1)’s full-support positivity and (C3)’s known-action-mechanism are *automatic* under the paper’s architecture — (C1) reduces to realized-action positivity ($\pi_t(a_t \mid H_t) > 0$ tautologically since a_t was drawn from π_t); (C3) reduces to agent-policy-queryability (Q_t is the agent’s own policy, internally accessible). Both are stated as named conditions for forward-compatibility with off-policy / counterfactual extensions where they would do real work; for the on-policy identification claim of this lemma, (C2) is the substantive condition.

B.6 Proof of Key Lemma 4 (Bias bound)

On the matching event $\{a_{\text{ag},t}^*(s) = \tilde{a}_t^*(s)\}$, the difference $\hat{D}_t(s) - D_t^{\text{true}}(s)$ vanishes identically (both quantities equal $-\log Q_t$ at the same action). On the complementary misidentification event, the support condition $Q_t(\cdot \mid s) \geq q_0$ at *both* $a_{\text{ag},t}^*(s)$ and $\tilde{a}_t^*(s)$ implies $\log Q_t(a_{\text{ag},t}^*(s) \mid s) \in [\log q_0, 0]$ and $\log Q_t(\tilde{a}_t^*(s) \mid s) \in [\log q_0, 0]$, so their absolute difference is at most $\log(1/q_0)$. Therefore $|\hat{D}_t(s) - D_t^{\text{true}}(s)| \leq \mathbf{1}[a_{\text{ag},t}^*(s) \neq \tilde{a}_t^*(s)] \cdot \log(1/q_0)$. Taking expectations: $\mathbb{E}[|\hat{D}_t(s) - D_t^{\text{true}}(s)|] \leq (1 - p_{\text{id}}(s)) \log(1/q_0)$ by definition of $p_{\text{id}}(s)$. $\square$

High-probability form With probability $1 - \delta$ over a single round’s identification event, $|\hat{D}_t(s) - D_t^{\text{true}}(s)| \leq \log(1/q_0) \cdot \mathbf{1}[\text{misident}]$ where $\Pr[\text{misident}] = 1 - p_{\text{id}}(s)$. The bias decomposes cleanly: in Regime A ($p_{\text{id}}(s) \rightarrow 1$) it vanishes; in Regime B ($p_{\text{id}}(s) \in (0, 1)$) it is controlled by misidentification mass; in Regime C ($p_{\text{id}}(s) \rightarrow 0$) it saturates at $\log(1/q_0)$. This proves Key Lemma 4. $\square$

Per-state, per-horizon-step lift (used in Theorem 4.1(v)). The lemma applies pointwise at any visited state s_h along a horizon- N_h trajectory: $\mathbb{E}[|\hat{D}_t(s_h) - D_t^{\text{true}}(s_h)|] \leq (1 - p_{\text{id}}(s_h)) \log(1/q_0)$, with $p_{\text{id}}(s_h)$ the per-state argmax-correctness probability. Aggregating linearly over the horizon and assuming the per-state floor $p_{\text{id}} := \min_s p_{\text{id}}(s)$, the per-round bias contribution to the trajectory-level value gap is at most $N_h(1 - p_{\text{id}}) \log(1/q_0)$, contributing the second term of Theorem 4.1(v).

Remarks - Both $a_{\text{ag},t}^*(s)$ and $\tilde{a}_t^*(s)$ must satisfy $Q_t \geq q_0$ for the $\log(1/q_0)$ ceiling to apply; the support condition is on *both* points, not just the agent’s estimate. This is *strictly weaker* than the perturbative identity’s full-support condition $Q(a) \geq q_0$ for *all* a (Appendix B.2) — the perturbative identity needs full support because π_ϵ^* has positive mass on every action; the bias bound here needs only two-point support. - This is *not* an estimation-from-samples lemma. If Q_t is unavailable internally and must be estimated from on-policy rollouts, an additional concentration argument (Hoeffding under iid fixed-policy samples, or martingale concentration for adaptive policies) is needed; in the architectures this paper considers, Q_t is the agent’s policy and is internally available.

C Auxiliary Mathematical Material

This appendix collects supporting material used by the main results: convention hierarchy details, the admissible-divergence family, direction-forcing argument, action-gap matching lower bound, tied-optimum extensions, the structural-class theorem on gain-decay updates, strategic-tempo topologies, deployment modes, and the chain-rule uniqueness theorem.

C.1 Convention hierarchy: C1, C2, C3

Three named continuation conventions form a hierarchy of increasing diagnostic power and computational cost.

C1 — One-step improvement (canonical default) $\pi_{\text{cont}} = \pi_{\text{current}}$. Each action is evaluated assuming current behavior continues afterward. No fixed-point computation, no global-optimality assumption. Cheapest; weakest diagnostic.

C2 — Receding-horizon (N_r -step replanning) At each future step, re-optimize over a horizon of N_r steps using the model available at that step: $\pi_{\text{RH}}(M_\tau) =$

822 $\arg \max_{\pi} V_O(M_{\tau}, \pi; N_r)$ applied at each τ . Captures multi-step recovery: a goal that ap-
 823 pears unattainable under frozen continuation may be reachable with replanning. Moderate cost;
 824 moderate diagnostic.

825 **C3 — Bellman** $\pi_{\text{cont}} = \pi^*$ where $\pi^* = \arg \max_{\pi} V_O(M_t, \pi; N_h)$. The continuation is the optimal
 826 policy — a fixed-point equation. Strongest diagnostic; most expensive.

827 **Proposition C.1** (Monotonicity). *For any model M_t , horizon N_h , and policy class Π :*
 828 $A_O^{(1)}(M_t; \Pi, N_h) \leq A_O^{\text{RH}}(M_t; \Pi, N_r, N_h) \leq A_O^{\text{B}}(M_t; \Pi, N_h)$.

829 *Proof.* C1 freezes continuation at π_{current} (possibly suboptimal); C2 re-optimizes periodically, so
 830 $\pi_{\text{RH}} \succeq \pi_{\text{current}}$ at each future step; C3 uses the globally optimal π^* . A weakly better continuation
 831 yields a weakly higher expected trajectory value; taking the supremum over the first action preserves
 832 the ordering. $\square$

833 **Corollary** $\delta_{\text{sat}}^{\text{B}} \leq \delta_{\text{sat}}^{\text{RH}} \leq \delta_{\text{sat}}^{(1)}$ (since $\delta_{\text{sat}} = V_O^{\min} - A_O$, higher A_O means lower δ_{sat}). And
 834 $\delta_{\text{regret}}^{(1)} \leq \delta_{\text{regret}}^{\text{RH}} \leq \delta_{\text{regret}}^{\text{B}}$ (since $\delta_{\text{regret}} = A_O - V_O(M_t, \pi_{\text{current}}; N_h)$, higher A_O means higher δ_{regret}).

835 C1 is the most conservative diagnostic (most likely to diagnose “locally unattainable”); C3 is the most
 836 accurate (least false “unattainable” diagnoses). The 2×2 diagnostic structure is preserved under all
 837 three; only the inferential force varies.

838 C.2 Admissible-divergence family for the regret bound

839 The reverse-KL direction is forced by the deterministic- π^* regime (Appendix C.3). Within the
 840 direction-forced family, multiple smooth f -divergences yield valid regret bounds:

	Divergence	Bound on $R(Q)$	Tightness	Finite under det. π^* ?
	$\text{TV}(\pi^*, Q)$	$V_{\max} \cdot \text{TV}$	Tight (extremal V)	Yes
	$D_{\text{KL}}(\pi^* \parallel Q)$ via Pinsker	$V_{\max} \sqrt{D_{\text{KL}}/2}$	Loose by $\sqrt{\cdot}$	Yes
841	$D_{\text{KL}}(\pi^* \parallel Q)$ via point-mass identity	$V_{\max}(1 - e^{-D_{\text{KL}}})$	Tight (this paper)	Yes
	$\chi^2(\pi^* \parallel Q)$ (Le Cam)	$V_{\max} \cdot \frac{1}{2} \sqrt{\chi^2}$	Looser than Pinsker	$\chi^2 = 1/Q(a^*) - 1$
	$D_{\alpha}(\pi^* \parallel Q)$ (Rényi, $\alpha \geq 1$)	Various	Interpolates	Yes for $\alpha \geq 1$
	$D_{\text{KL}}(Q \parallel \pi^*)$ (forward-KL)	$+\infty$	Vacuous	No

842 Reverse-KL is selected uniquely within the direction-forced family by the chain-rule axiom (Ap-
 843 pendix C.10). The point-mass identity supplies the *exact* bound on reverse-KL under determinis-
 844 tic π^* ; the table reflects different bound shapes on the same divergence (with the BH inequality
 845 $V_{\max} \sqrt{1 - e^{-D_{\text{KL}}}}$ as the looser general form to which the identity reduces outside the deterministic- π^*
 846 corner).

847 C.3 Direction-forcing argument

848 For deterministic $\pi^* = \delta_{a^*}$ and any Q with $Q(a) > 0$ for some $a \neq a^*$: $D_{\text{KL}}(Q \parallel \pi^*) =$
 849 $\sum_a Q(a) \log \frac{Q(a)}{\pi^*(a)} = \sum_{a \neq a^*} Q(a) \log \frac{Q(a)}{0} = +\infty$. A bound “ $R \leq +\infty$ ” is vacuous. The
 850 reverse direction $D_{\text{KL}}(\pi^* \parallel Q)$ is finite (and equal to $-\log Q(a^*)$) whenever $Q(a^*) > 0$. The asym-
 851 metry is forced by the regime: regret is a one-sided quantity (contributes only from Q ’s off-optimum
 852 mass; π^* has no support off a^*); divergences whose role is to bound this one-sided quantity must
 853 themselves be one-sided. Symmetric divergences (squared Hellinger, Jensen-Shannon, symmetrized
 854 KL) introduce a vacuous symmetric term.

855 C.4 Action-gap matching lower bound

856 For any Q with $\Delta_{\min} = \min_{a \neq a^*} \Delta(a) > 0$: $R(Q) = \sum_{a \neq a^*} Q(a) \Delta(a) \geq \Delta_{\min} \sum_{a \neq a^*} Q(a) =$
 857 $\Delta_{\min} \cdot (1 - Q(a^*)) = \Delta_{\min} \cdot \text{TV}(\pi^*, Q)$. By the point-mass identity (Lemma 5.1), $\text{TV}(\pi^*, Q) =$
 858 $1 - e^{-D_{\text{KL}}(\pi^* \parallel Q)}$, giving the matching lower bound used in the two-sided characterization.

859 The lower bound is tight when sub-optimal actions are uniformly bad ($\Delta_{\min} = \max_{a \neq a^*} \Delta(a)$). For
 860 typical landscapes the gap between upper and lower bound is $V_{\max} - \Delta_{\min}$, controlled by the *spread*
 861 of action gaps.

862 C.5 Tied-optimum and softmax-smoothed extensions

863 **Tied-optimum** If π^* has support on a tied-optimum set $\mathcal{A}^* = \{a : Q_O(a) = Q_O(a^*)\}$
 864 with uniform mass, reverse-KL is finite whenever Q covers $\mathcal{A}^*$. The regret bound extends:
 865 $R(Q) = \sum_{a \notin \mathcal{A}^*} Q(a) \Delta(a) \leq V_{\max} \cdot \mathbb{P}_Q(a \notin \mathcal{A}^*)$. A multi-action analog of the point-mass
 866 identity holds with $\pi^*(a^*) = 1/|\mathcal{A}^*|$ replacing the point-mass form, yielding $D_{\text{KL}}(\pi^* \| Q) =$
 867 $-\sum_{a \in \mathcal{A}^*} (1/|\mathcal{A}^*|) \log(|\mathcal{A}^*| Q(a))$, a multi-action analog of $-\log Q(a^*)$. The two-sided bound be-
 868 comes one-sided in general; the upper bound holds with the modified KL form.

869 **Softmax-smoothed π^* with temperature $\tau \rightarrow 0$** Replace $\pi^* = \delta_{a^*}$ with $\pi_\tau^* \propto \exp(Q_O/\tau)$. As
 870 $\tau \rightarrow 0$, $\pi_\tau^* \rightarrow \delta_{a^*}$. For finite τ , the *perturbative* identity (Appendix B.2) applies: $\text{TV}(\pi_\tau^*, Q) =$
 871 $1 - e^{-D_{\text{KL}}(\pi_\tau^* \| Q)} + O(\tau^{-1} \exp(-\Delta_{\min}/\tau))$ — exponentially small with polynomial prefactor. The
 872 two-sided regret bound transfers with the same correction order. Outside the perturbative regime —
 873 for genuinely high-entropy optima — the BH inequality $\text{TV} \leq \sqrt{1 - e^{-D_{\text{KL}}}}$ becomes the relevant
 874 general bound; Pinsker also applies, with a one-sided regret form.

875 C.6 The gain-decay structural class $\mathcal{A}_{\text{decay}}$

876 The persistence inequality $\mathcal{T}_{\Sigma}^{\text{bn,ss}} > \rho_{\Sigma}/R_{\Sigma}$ (Section 5.2) is an *instantaneous* check at the current
 877 operating point. Define the structural class by gain-decay rather than sample-counting: $\mathcal{A}_{\text{decay}} :=$
 878 $\{\text{updates whose effective sector gain } \alpha_{ij}^{(t)} \rightarrow 0 \text{ as } t \rightarrow \infty \text{ on every revisable element}\}$.

879 **Theorem C.2** (Universal failure of $\mathcal{A}_{\text{decay}}$). *For any fixed $(\rho_{\Sigma}, R_{\Sigma})$ with $\rho_{\Sigma} > 0$, every agent in*
 880 $\mathcal{A}_{\text{decay}}$ *eventually violates the persistence threshold $\mathcal{T}_{\Sigma}^{\text{bn,ss}} > \rho_{\Sigma}/R_{\Sigma}$.*

881 *Proof.* At every element (i, j) , $\alpha_{ij}^{(t)} \rightarrow 0$ by definition of $\mathcal{A}_{\text{decay}}$, so $\min_{(i,j)} \alpha_{ij}^{(t)} \rightarrow 0$, so the
 882 bottleneck $\mathcal{T}_{\Sigma}^{\text{bn}} \rightarrow 0$ with experience. Once $\mathcal{T}_{\Sigma}^{\text{bn,ss}} < \rho_{\Sigma}/R_{\Sigma}$, the threshold is violated. This holds
 883 *regardless of prior calibration*: at any finite calibration level, the gain decays below threshold once
 884 enough experience accumulates. $\square$

885 *Scope clarification (per-element pull).* The class definition requires $\alpha_{ij}^{(t)} \rightarrow 0$ on *every* revisable
 886 element — i.e., the agent is being pulled at every element repeatedly. Beta-Bernoulli with $\eta_{\text{edge},ij} =$
 887 $1/(n_{ij} + 1)$ is in the class only for elements actually being updated; an agent that ignores a subset of
 888 elements is outside the class for those elements (their gain remains at initial calibration), but inside
 889 it for the elements it does touch. The theorem applies to the touched elements; persistence on the
 890 untouched ones reduces to whatever calibration was set initially.

891 This is a *structural* failure of the class, not a tuning problem. The class includes: - Count-accumulating
 892 Bayesian updates without forgetting (e.g., Beta-Bernoulli with $\eta_{\text{edge},ij} = 1/(n_{ij} + 1)$ where $n_{ij} \rightarrow \infty$).
 893 - Bounded-memory schemes with growing memory. - Observation-aggregating schemes without restart.
 894 - Gradient-based methods with vanishing step sizes ($\eta_t = c/t^{\beta}$ for $\beta \in (0, 1]$, the Robbins–Monro
 895 regime).

896 Mechanisms *outside* $\mathcal{A}_{\text{decay}}$ — constant-step-size stochastic approximation, sliding-window updates,
 897 bounded-memory learners, block-restart schemes, Kalman filters with bounded process noise —
 898 maintain a finite gain ceiling and escape asymptotic decay, but then face a *bidirectional* threshold:

899 For each row, the threshold is *bidirectional* and sharp within the sector-Lyapunov reduction.

900 C.7 Strategic-tempo across canonical topologies

901 The bottleneck strategic tempo $\mathcal{T}_{\Sigma}^{\text{bn,ss}}$ of Lemma 5.2 is verified — and the resulting per-topology
 902 threshold derived — across four canonical topologies under Beta-Bernoulli edge updates with per-
 903 element forgetting:

Table 1: Bidirectional thresholds for non-accumulating mechanisms outside $\mathcal{A}_{\text{decay}}$.

Mechanism	$\alpha_{\Sigma}^{\text{ss}}$	Prerequisite holds iff
Exponential forgetting with λ	$1 - \lambda$	$1 - \lambda > \rho_{\Sigma}/R_{\Sigma}$
Constant-step SA with rate η	η	$\eta > \rho_{\Sigma}/R_{\Sigma}$
Fixed-window mechanisms (sliding-window size W , bounded-memory size K , block-restart period T_R)	$1/W, 1/K, 1/T_R$	window/memory/period $< R_{\Sigma}/\rho_{\Sigma}$

- 904 • **(T1) Single edge.** $\mathcal{T}_{\Sigma}^{\text{bn,ss}} = \nu \cdot \iota \cdot (1 - \lambda)$. All three factors load-bearing; setting any one to
905 zero collapses the bottleneck.
- 906 • **(T2) Two-edge AND chain, observable intermediate.** $\mathcal{T}_{\Sigma}^{\text{bn,ss}} = \min(\nu_1(1 - \lambda_1), \nu_1\theta_1(1 -$
907 $\lambda_2))$. Edge 2’s effective observation rate is gated by edge 1’s success probability θ_1 —
908 *depth-gated attenuation*. For depth- d chains, the bottleneck is $\min_k \prod_{j < k} \theta_j(1 - \lambda_k)$.
- 909 • **(T3) Two-edge AND chain, unobservable intermediate.** Per-edge tempo ill-defined; plan-
910 level bottleneck $\mathcal{T}_{\Sigma, \text{plan}}^{\text{bn,ss}} = \nu(1 - \lambda_{\Phi})$ over a single tracked plan-level quantity $\hat{\Phi} = p_1p_2$.
- 911 • **(T4) Two-arm OR node, ε -greedy.** $\mathcal{T}_{\Sigma}^{\text{bn,ss}} = \min((1 - \varepsilon)(1 - \lambda_1), \varepsilon(1 - \lambda_2))$. Action
912 selection controls rate allocation; pure greedy ($\varepsilon = 0$) collapses the unexplored arm’s
913 bottleneck — *exploration-gated*.

914 The structural decomposition: AND-chains exhibit *depth-gated* geometric attenuation $\nu_k =$
915 $\nu \prod_{j < k} \theta_j$; OR-nodes exhibit *exploration-gated* allocation. Mixed AND/OR DAGs interleave both.
916 Each topology surfaces a different factor as the binding bottleneck.

917 **Regime-C contamination is structurally fatal** $\iota_{ij} = 0$ at any single element $\Rightarrow \mathcal{T}_{\Sigma}^{\text{bn,ss}} = 0$. An
918 agent with even one fully-confounded element cannot meet the threshold by forgetting alone.

919 C.8 Three deployment modes of loop-Level-2

920 The Pearl-do structure of closed-loop interventional access manifests at distinct layers with semanti-
921 cally distinct interventional mechanisms:

- 922 • **Mode 1 — agent-self-intervention.** The agent performs do-actions on its own action space
923 as part of its ordinary adaptive loop. Intervention is on the agent’s own action; target is the
924 environment’s response. This is Lemma 5.3’s primary content.
- 925 • **Mode 2 — observer-on-sub-agent.** An observer external to a composite performs do-
926 interventions on one sub-agent’s action space; target is another sub-agent’s response. Reveals
927 cross-coupling structure that component-marginal observation cannot identify.
- 928 • **Mode 3 — observer-on-agent-input** (sketched, future work). An observer intervenes on
929 the agent’s observation channel; target is the agent’s subsequent policy. Useful for offline
930 architectural-class diagnosis.

931 Modes share the Pearl-do structure but differ in who intervenes on what. The unification is at the
932 pattern level; the mechanism is semantically distinct per layer.

933 C.9 Distinction from active inference and causal-RL precursors

934 Action-perception-loop frameworks — active inference^[46,47], control-as-inference^[48], cybernet-
935 ics^[49,50] — implicitly use the action-causes-observation observation. Our distinctive moves:

- 936 • *Bareinboim-hierarchy connection.* Active inference / cybernetics rest on Bayesian-network
937 (Level 1) generative models; we invoke the causal-hierarchy theorem^[26] to position the
938 policy DAG as causal with do-conditioning in Q_O .
- 939 • *Regime-indexed identifiability (A/B/C).* Active inference literature treats causal identifiability
940 uniformly within modeling assumptions; we surface the regime split at framework level.

941 • *Scope honesty.* We distinguish “data generated under intervention” from “cleanly identified
942 do-estimates”;^[43] documents that the active-inference literature sometimes elides this.

943 The causal-RL line^[13–17] is the direct ancestor for regime-indexed identifiability and on-policy inter-
944 ventional access; all are stationary-MDP. Composition with non-stationarity is, to our knowledge,
945 novel.

946 C.10 Chain-rule uniqueness of reverse-KL

947 **Theorem C.3** (Chain-rule uniqueness (Hobson 1969; Csiszár 1991)). *Let $D_f(P \| Q) =$
948 $\sum_x Q(x) f(P(x)/Q(x))$ be a smooth f -divergence with f convex and $f(1) = 0$. The chain rule
949 $D_f(P_{XY} \| Q_{XY}) = D_f(P_X \| Q_X) + \mathbb{E}_{P_X}[D_f(P_{Y|X} \| Q_{Y|X})]$ holds for all joint distributions
950 if and only if $f(t) = c \cdot t \log t$ for some $c > 0$ — i.e., D_f is reverse-KL up to positive scaling.*

951 *Proof.* Take any joint P_{XY}, Q_{XY} on a finite product space. Writing $r_x := P(x)/Q(x)$ and $s_{y|x} :=$
952 $P(y | x)/Q(y | x)$, the joint ratio factorizes as $P_{XY}(x, y)/Q_{XY}(x, y) = r_x \cdot s_{y|x}$. Expanding both
953 sides of the chain rule: $D_f(P_{XY} \| Q_{XY}) = \sum_{x,y} Q_{XY}(x, y) f(r_x s_{y|x}) = \sum_x Q(x) \sum_y Q(y | x)$
954 $f(r_x s_{y|x})$, $D_f(P_X \| Q_X) + \mathbb{E}_{P_X}[D_f(P_{Y|X} \| Q_{Y|X})] = \sum_x Q(x) f(r_x) + \sum_x P(x) \sum_y Q(y | x)$
955 $f(s_{y|x})$, the second term using $\mathbb{E}_{P_X} = \sum_x P(x) = \sum_x Q(x) r_x$. Equating and collecting:
956 $\sum_x Q(x) \left[\sum_y Q(y | x) f(r_x s_{y|x}) - f(r_x) - r_x \sum_y Q(y | x) f(s_{y|x}) \right] = 0$. The chain rule
957 must hold for *all* joints, including those where r_x takes arbitrary positive value at one x and $s_{y|x}$
958 varies independently — which forces the bracketed expression to vanish pointwise. A two-point
959 reduction (take Y binary with $s_{y|x}$ taking values s and $1/s$ on equal-weight $Q(y | x)$) collapses
960 the inner sums and yields the functional equation $f(rs) = f(r) + r f(s)$ for all $r, s > 0$.
961 With f convex and $f(1) = 0$, the unique solution is $f(t) = c \cdot t \log t$ for $c > 0$ ^[51] (§4) — this
962 gives $D_f(P \| Q) = c \sum_x Q(x) (P(x)/Q(x)) \log(P(x)/Q(x)) = c \sum_x P(x) \log(P(x)/Q(x)) =$
963 $c \cdot D_{\text{KL}}(P \| Q)$, reverse-KL up to positive scaling. $\square$

964 **References** Hobson 1969 (“A new theorem of information theory,” *J. Stat. Phys.*); Csiszár 1991
965 (“Why least squares and maximum entropy?” *Annals of Statistics*; Theorem 3 corollary, Theorem 5);
966 Shore-Johnson 1980 (“Axiomatic derivation of the principle of maximum entropy,” *IEEE Trans. Info.*
967 *Theory*, system-independence axiom); Sanov 1957 (large-deviation rate function); Aczél-Daróczy
968 1975 (functional-equation machinery).

969 These references give *structurally equivalent reformulations* of the same axiom. The Cauchy functional
970 equation each reduces to is the common content. No known uniqueness route outside the independence-
971 on-sub-problems family exists.

972 **Why other family members fail the chain rule** Concrete counterexample for χ^2 : take Q_X uniform
973 on $\{x_1, x_2\}$, $P_X = (3/4, 1/4)$, $Q(y|x)$ uniform, $P(y|x) = (3/4, 1/4)$. Direct calculation gives
974 $\chi^2(P_{XY} \| Q_{XY}) = 9/16$ while $\chi^2(P_X \| Q_X) + \mathbb{E}_{P_X}[\chi^2(P_{Y|X} \| Q_{Y|X})] = 1/4 + 1/4 = 8/16$.
975 Non-additive. Rényi- α for $\alpha \neq 1$ fails analogously; squared Hellinger likewise fails.

976 D Algorithm Sketch and Base-Learner Instantiation

977 A practical algorithm achieving the joint guarantees of Theorem 4.1 requires four components, each
978 translating one of the assumptions (A1)–(A5) into operational form.

979 **(a) Identity-tight base learner satisfying (A5)** In the stationary deterministic- π^* bandit case ($N_h =$
980 1), Thompson sampling^[18] (posterior-randomized; pointwise $Q_t(a^*) > 0$) and vanilla deterministic
981 UCB^[32] (Dirac-deployed; $Q_t(a^*) \in \{0, 1\}$, with the (A1) extended-real reading of Theorem 4.1) both
982 achieve the per-round coordinate in expectation: $\mathbb{E}[1 - e^{-K_t}] = \mathbb{E}[1 - Q_t(a^*)] = O\left(\frac{\log t}{t \Delta_{\min}^2}\right)$,
983 following from Lattimore and Szepesvári 2020^[32] Theorem 7.1 ($\mathbb{E}[N_a(t)] \leq 16 \log(t)/\Delta_a^2 + O(1)$)
984 per suboptimal arm; aggregating over arms with the per-state floor gives the per-round form). This
985 satisfies (A5) of Theorem 4.1 with the *stronger logarithmic* rate, yielding cumulative dynamic

regret $O(V_{\max}(B_T + 1) \log(T/(B_T + 1))/\Delta_{\min}^2)$. The $V_{\max} \cdot \text{TV}$ chain is structurally one factor of $\Delta_{\min}$ looser than direct gap-aware UCB analysis (which gives $O(\log T/\Delta_{\min})$ cumulative regret per block by weighting each suboptimal pull by its own $\Delta(a)$ rather than $V_{\max}$); this is a feature of the unification — the framework’s contribution is composition + non-stationarity-handling + identity-coordinate exactness, not a sharper bandit constant. For $N_h > 1$ MDPs, UCRL2^[52] and UCBVI^[39] give per-block cumulative trajectory-level regret $\tilde{O}(N_h^{3/2} \sqrt{SAL})$ with S, A the state/action sizes; rewriting in the occupancy-weighted coordinate satisfies (A5) with c of order $\tilde{O}(\sqrt{N_h SA})$, lifting cleanly to the $\tilde{O}(N_h^2 \sqrt{SA(B_T + 1)T})$ cumulative rate — same piecewise-stationary B_T family as restart-on-change variants of^[1] (the continuous-variation regime of^[3] is structurally distinct; see Section 4 *Comparator regime*). The occupancy-weighted form is the natural object for trajectory-level base learners; the per-state KL/TV identity is preserved underneath. Information-directed sampling^[19] requires a different analysis here because $H(\pi^*) = 0$ collapses its information ratio at the deterministic- π^* corner; we leave the IDS analysis as future work.

(b) Multi-factor forgetting schedule satisfying (A2) Choose per-element discount rates $\{\lambda_{ij}\}$ such that the bottleneck $\mathcal{T}_{\Sigma}^{\text{bn,ss}} = \min_{(i,j)} \nu_{ij} \ell_{ij} (1 - \lambda_{ij})$ exceeds ρ_{Σ}/R_{Σ} . When ρ_{Σ} is unknown, the^[2] black-box reduction gives an adaptive choice (applied per element if heterogeneous); when ρ_{Σ} is known via a variation-budget^[1] or change-point detection^[35], the choice is direct. As a fail-fast pre-check, $\mathcal{T}_{\Sigma}^{\text{agg,ss}} > |E| \cdot \rho_{\Sigma}/R_{\Sigma}$ is necessary.

(c) Identifiability check satisfying (A3) Test whether the loop data is in Regime A (full identifiability), Regime B (partial), or Regime C (none). In Regime A, no adjustment is needed. In Regime B, apply^[16] (DOVI)-style confounding-bias adjustment. In Regime C, the bound is provable but not realizable; algorithm flags the regime as out-of-scope.

(d) Two-gap diagnostic readout satisfying (A4) Compute δ_{sat} (requires A_O , intractable in general; estimable via the policy-class supremum over recent rounds) and δ_{regret} (requires V_O at current policy; tractable in simulation). Route to corrective action class.

The algorithm is sketch-grade; full instantiation and empirical evaluation are deferred to a follow-up paper.

E Pinsker vs. Point-Mass Identity Numerical Comparison

We compare $V_{\max}(1 - e^{-D_{\text{KL}}})$ (point-mass identity bound, exact under deterministic π^* per Lemma 5.1 and the two-sided regret theorem) with $V_{\max} \sqrt{D_{\text{KL}}/2}$ (Pinsker bound, also valid under deterministic π^* but loose), $V_{\max} \sqrt{1 - e^{-D_{\text{KL}}}}$ (Bretagnolle–Huber inequality, the general envelope below which the identity sits at this corner), and the trivial envelope $V_{\max}$. Set $V_{\max} = 1$ for normalization.

Table 2: Identity vs. Pinsker vs. BH at $V_{\max} = 1$.

D_{KL}	$1 - e^{-D_{\text{KL}}}$ (identity)	$\sqrt{1 - e^{-D_{\text{KL}}}}$ (BH)	$\sqrt{D_{\text{KL}}/2}$ (Pinsker)	ratio
0.01	0.00995	0.0998	0.0707	7.10×
0.1	0.0952	0.308	0.224	2.35×
0.5	0.393	0.627	0.500	1.27×
1.0	0.632	0.795	0.707	1.12×
2.0	0.865	0.930	1.000	1.16×
4.0	0.982	0.991	1.414	1.02×
10.0	0.99995	0.99998	2.236	1.00×

The point-mass identity bound is uniformly tighter than Pinsker, by a factor of $7\times$ at $D_{\text{KL}} = 0.01$ and converging to 1 as D_{KL} grows large (where both saturate at $V_{\max}$). It is also uniformly tighter than the BH inequality $\sqrt{1 - e^{-D_{\text{KL}}}}$ — the relation $x < \sqrt{x}$ on $(0, 1)$ — although both saturate together at $V_{\max}$ for large D_{KL} . Pinsker, by contrast, hits the trivial envelope $V_{\max} = 1$ at $D_{\text{KL}} = 2$ and exceeds it beyond — vacuous from $D_{\text{KL}} = 4$ onward (where the unclipped form returns 1.414), fully vacuous by $D_{\text{KL}} = 10$.

1024 The matching lower bound $\Delta_{\min}(1 - e^{-D_{\text{KL}}})$ has the same shape, scaled by $\Delta_{\min}/V_{\max}$. In the extremal
 1025 value landscape ($\Delta_{\min} = V_{\max}$), the upper and lower bounds coincide and the regret is *identified*
 1026 *exactly* by the KL coordinate. For typical landscapes, regret and KL are Lipschitz-equivalent with
 1027 constants $\Delta_{\min}/V_{\max}$ and 1.

1028 F Prior-Art Retrieval Summary

1029 We conducted a structured prior-art retrieval (via Undermind) with the goal of identifying any frame-
 1030 work composing all four target elements: (1) goal-feasibility-vs-policy-quality two-gap diagnostic;
 1031 (2) an exact reverse-KL/TV regret identity (point-mass) under deterministic- π^* , or — more loosely
 1032 — Bretagnolle–Huber-style use in regret analysis; (3) strategic-tempo / forgetting prerequisite tying
 1033 convergence to rate of useful policy revision; (4) closed-loop interventional access making regret
 1034 bounds learnable from on-policy data.

1035 **Result: 63 papers retrieved, abstract-level coverage estimated 75%, no direct anticipation in**
 1036 **the retrieved set** The landscape splits into four largely separate strands corresponding to our four
 1037 components; we did not find a published framework composing all four in this retrieval. Retrieval
 1038 queries used the structured prior-art tool Undermind; the four-strand cite-and-distinguish in Section 2
 1039 surfaces the closest neighbor in each strand.

1040 **Strongest negative signal (in the retrieved corpus): the Bretagnolle–Huber inequality is absent**
 1041 **as an upper-bound coordinate for regret analysis in the retrieved RL / non-stationary literature**
 1042 The information-theoretic regret line —^[18–23] — uses entropy / mutual information / information
 1043 ratio / Pinsker / Hellinger; we did not find the BH inequality deployed *as an upper-bound regret*
 1044 *coordinate*. (BH is well-known in bandit *lower-bound* proofs via Le Cam’s method and Tsybakov’s
 1045 framework^[31], including in standard RL textbook treatments^[32]; that use is not contested here.) The
 1046 *exact point-mass identity at the deterministic- π corner** (Lemma 5.1) — which strictly improves both
 1047 Pinsker and the BH envelope at this corner — was likewise not found in the retrieved corpus, in either
 1048 upper-bound or lower-bound use. We frame this as a literature-search observation against an explicitly
 1049 bounded retrieval (63 papers, ~75% abstract-level coverage), not a formal absence claim against all
 1050 published RL work; the claim is best read as “to our knowledge / in the retrieved set” rather than
 1051 “absolutely novel.”

1052 **Closest neighbors per strand:**^[1] (Strand 1 — variation-budget dynamic regret);^[5] (Strand 2 —
 1053 closest two-term decomposition, different axis);^[8] (Strand 3 — closest tempo result, different form);^[13]
 1054 (Strand 4 — closest causal-RL, stationary only; the foundational MDPUC paper introducing the
 1055 closed-loop interventional-vs-observational structural argument that Component 4 generalizes).

1056 NeurIPS Paper Checklist

1057 1. Claims

1058 Question: Do the main claims made in the abstract and introduction accurately reflect the
1059 paper’s contributions and scope?

1060 Answer: [Yes]

1061 Justification: The abstract and the contribution subsection of the introduction state the
1062 four-component composition (two-gap diagnostic, point-mass reverse-KL/TV identity, multi-
1063 factor strategic tempo with forgetting prerequisite, closed-loop interventional access) and
1064 the cumulative dynamic regret rate. Each component is developed in the main-result and
1065 mechanism sections and proved in full in the appendices.

1066 Guidelines:

- 1067 • The answer [N/A] means that the abstract and introduction do not include the claims
1068 made in the paper.
- 1069 • The abstract and/or introduction should clearly state the claims made, including the
1070 contributions made in the paper and important assumptions and limitations. A [No] or
1071 [N/A] answer to this question will not be perceived well by the reviewers.
- 1072 • The claims made should match theoretical and experimental results, and reflect how
1073 much the results can be expected to generalize to other settings.
- 1074 • It is fine to include aspirational goals as motivation as long as it is clear that these goals
1075 are not attained by the paper.

1076 2. Limitations

1077 Question: Does the paper discuss the limitations of the work performed by the authors?

1078 Answer: [Yes]

1079 Justification: The scope-and-limitations subsection of the introduction (§1.2) and the conclu-
1080 sion (§6) together discuss theory-only scope, canonical-scope assumptions (deterministic
1081 π^* , bounded value range, isolated optimum, singular trajectory), per-state vs. cumulative
1082 regret scope (Cesàro tracking under $B_T = o(T)$), partial uniqueness at the metric layer
1083 (chain-rule uniqueness over reverse-KL only; joint uniqueness across all three properties
1084 left open), and the dependence on architectural separation between M_t and goal state for
1085 coupled-goal architectures.

1086 Guidelines:

- 1087 • The answer [N/A] means that the paper has no limitation while the answer [No] means
1088 that the paper has limitations, but those are not discussed in the paper.
- 1089 • The authors are encouraged to create a separate “Limitations” section in their paper.
- 1090 • The paper should point out any strong assumptions and how robust the results are to
1091 violations of these assumptions (e.g., independence assumptions, noiseless settings,
1092 model well-specification, asymptotic approximations only holding locally). The authors
1093 should reflect on how these assumptions might be violated in practice and what the
1094 implications would be.
- 1095 • The authors should reflect on the scope of the claims made, e.g., if the approach was
1096 only tested on a few datasets or with a few runs. In general, empirical results often
1097 depend on implicit assumptions, which should be articulated.
- 1098 • The authors should reflect on the factors that influence the performance of the approach.
1099 For example, a facial recognition algorithm may perform poorly when image resolution
1100 is low or images are taken in low lighting. Or a speech-to-text system might not be
1101 used reliably to provide closed captions for online lectures because it fails to handle
1102 technical jargon.
- 1103 • The authors should discuss the computational efficiency of the proposed algorithms
1104 and how they scale with dataset size.
- 1105 • If applicable, the authors should discuss possible limitations of their approach to address
1106 problems of privacy and fairness.

- While the authors might fear that complete honesty about limitations might be used by reviewers as grounds for rejection, a worse outcome might be that reviewers discover limitations that aren't acknowledged in the paper. The authors should use their best judgment and recognize that individual actions in favor of transparency play an important role in developing norms that preserve the integrity of the community. Reviewers will be specifically instructed to not penalize honesty concerning limitations.

3. Theory assumptions and proofs

Question: For each theoretical result, does the paper provide the full set of assumptions and a complete (and correct) proof?

Answer: [Yes]

Justification: All theorems and lemmas are numbered, stated with their assumptions in-place, and cross-referenced. Theorem 4.1 (composition: joint guarantees with conditional cumulative reduction) is stated in §4 and proved in full in Appendix A. Lemmas 5.1–5.4 (point-mass reverse-KL/TV identity; multi-factor forgetting prerequisite; closed-loop interventional access; bias bound under partial identifiability) are stated with proof sketches in §5 and proved in full in Appendix B, including the perturbative extension to ϵ -stochastic and softmax-regularized optima and the impulsive-ProST sector-level reduction. Proposition C.1 (monotonicity of the conventions hierarchy), Theorem C.2 (universal failure of $\mathcal{A}_{\text{decay}}$), and Theorem C.3 (chain-rule uniqueness of reverse-KL) are stated and proved in Appendix C.

Guidelines:

- The answer [N/A] means that the paper does not include theoretical results.
- All the theorems, formulas, and proofs in the paper should be numbered and cross-referenced.
- All assumptions should be clearly stated or referenced in the statement of any theorems.
- The proofs can either appear in the main paper or the supplemental material, but if they appear in the supplemental material, the authors are encouraged to provide a short proof sketch to provide intuition.
- Inversely, any informal proof provided in the core of the paper should be complemented by formal proofs provided in appendix or supplemental material.
- Theorems and Lemmas that the proof relies upon should be properly referenced.

4. Experimental result reproducibility

Question: Does the paper fully disclose all the information needed to reproduce the main experimental results of the paper to the extent that it affects the main claims and/or conclusions of the paper (regardless of whether the code and data are provided or not)?

Answer: [N/A]

Justification: Theory-only paper; no original experiments. The worked-example reduction to ProST (Lemma 5.2) draws on Lee et al. 2023^[8]'s published experiments; reproducibility there is theirs.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- If the paper includes experiments, a [No] answer to this question will not be perceived well by the reviewers: Making the paper reproducible is important, regardless of whether the code and data are provided or not.
- If the contribution is a dataset and/or model, the authors should describe the steps taken to make their results reproducible or verifiable.
- Depending on the contribution, reproducibility can be accomplished in various ways. For example, if the contribution is a novel architecture, describing the architecture fully might suffice, or if the contribution is a specific model and empirical evaluation, it may be necessary to either make it possible for others to replicate the model with the same dataset, or provide access to the model. In general, releasing code and data is often one good way to accomplish this, but reproducibility can also be provided via detailed instructions for how to replicate the results, access to a hosted model (e.g., in the case of a large language model), releasing of a model checkpoint, or other means that are appropriate to the research performed.

- While NeurIPS does not require releasing code, the conference does require all submissions to provide some reasonable avenue for reproducibility, which may depend on the nature of the contribution. For example
 - (a) If the contribution is primarily a new algorithm, the paper should make it clear how to reproduce that algorithm.
 - (b) If the contribution is primarily a new model architecture, the paper should describe the architecture clearly and fully.
 - (c) If the contribution is a new model (e.g., a large language model), then there should either be a way to access this model for reproducing the results or a way to reproduce the model (e.g., with an open-source dataset or instructions for how to construct the dataset).
 - (d) We recognize that reproducibility may be tricky in some cases, in which case authors are welcome to describe the particular way they provide for reproducibility. In the case of closed-source models, it may be that access to the model is limited in some way (e.g., to registered users), but it should be possible for other researchers to have some path to reproducing or verifying the results.

5. Open access to data and code

Question: Does the paper provide open access to the data and code, with sufficient instructions to faithfully reproduce the main experimental results, as described in supplemental material?

Answer: [N/A]

Justification: Theory-only paper; no code or data accompanies the contribution.

Guidelines:

- The answer [N/A] means that paper does not include experiments requiring code.
- Please see the NeurIPS code and data submission guidelines (<https://neurips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- While we encourage the release of code and data, we understand that this might not be possible, so [No] is an acceptable answer. Papers cannot be rejected simply for not including code, unless this is central to the contribution (e.g., for a new open-source benchmark).
- The instructions should contain the exact command and environment needed to run to reproduce the results. See the NeurIPS code and data submission guidelines (<https://neurips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- The authors should provide instructions on data access and preparation, including how to access the raw data, preprocessed data, intermediate data, and generated data, etc.
- The authors should provide scripts to reproduce all experimental results for the new proposed method and baselines. If only a subset of experiments are reproducible, they should state which ones are omitted from the script and why.
- At submission time, to preserve anonymity, the authors should release anonymized versions (if applicable).
- Providing as much information as possible in supplemental material (appended to the paper) is recommended, but including URLs to data and code is permitted.

6. Experimental setting/details

Question: Does the paper specify all the training and test details (e.g., data splits, hyperparameters, how they were chosen, type of optimizer) necessary to understand the results?

Answer: [N/A]

Justification: Theory-only paper; no experiments.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- The experimental setting should be presented in the core of the paper to a level of detail that is necessary to appreciate the results and make sense of them.
- The full details can be provided either with the code, in appendix, or as supplemental material.

7. Experiment statistical significance

Question: Does the paper report error bars suitably and correctly defined or other appropriate information about the statistical significance of the experiments?

Answer: [N/A]

Justification: Theory-only paper; no experiments.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- The authors should answer [Yes] if the results are accompanied by error bars, confidence intervals, or statistical significance tests, at least for the experiments that support the main claims of the paper.
- The factors of variability that the error bars are capturing should be clearly stated (for example, train/test split, initialization, random drawing of some parameter, or overall run with given experimental conditions).
- The method for calculating the error bars should be explained (closed form formula, call to a library function, bootstrap, etc.)
- The assumptions made should be given (e.g., Normally distributed errors).
- It should be clear whether the error bar is the standard deviation or the standard error of the mean.
- It is OK to report 1-sigma error bars, but one should state it. The authors should preferably report a 2-sigma error bar than state that they have a 96% CI, if the hypothesis of Normality of errors is not verified.
- For asymmetric distributions, the authors should be careful not to show in tables or figures symmetric error bars that would yield results that are out of range (e.g., negative error rates).
- If error bars are reported in tables or plots, the authors should explain in the text how they were calculated and reference the corresponding figures or tables in the text.

8. Experiments compute resources

Question: For each experiment, does the paper provide sufficient information on the computer resources (type of compute workers, memory, time of execution) needed to reproduce the experiments?

Answer: [N/A]

Justification: Theory-only paper; no experiments.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- The paper should indicate the type of compute workers CPU or GPU, internal cluster, or cloud provider, including relevant memory and storage.
- The paper should provide the amount of compute required for each of the individual experimental runs as well as estimate the total compute.
- The paper should disclose whether the full research project required more compute than the experiments reported in the paper (e.g., preliminary or failed experiments that didn't make it into the paper).

9. Code of ethics

Question: Does the research conducted in the paper conform, in every respect, with the NeurIPS Code of Ethics <https://neurips.cc/public/EthicsGuidelines>?

Answer: [Yes]

Justification: Theory paper without human subjects, datasets, or model release; conforms with the NeurIPS Code of Ethics throughout.

Guidelines:

- The answer [N/A] means that the authors have not reviewed the NeurIPS Code of Ethics.
- If the authors answer [No], they should explain the special circumstances that require a deviation from the Code of Ethics.
- The authors should make sure to preserve anonymity (e.g., if there is a special consideration due to laws or regulations in their jurisdiction).

10. Broader impacts

Question: Does the paper discuss both potential positive societal impacts and negative societal impacts of the work performed?

Answer: [N/A]

Justification: This is foundational theory work on convergence guarantees for non-stationary reinforcement learning; the contribution is a regret-bound analysis tool, not a deployable algorithm or model. We see no direct path from sharper non-stationary-RL convergence theory to either positive or negative societal applications that would not equally apply to the broader RL literature, and we do not release any models, datasets, or code that would carry deployment risks.

Guidelines:

- The answer [N/A] means that there is no societal impact of the work performed.
- If the authors answer [N/A] or [No], they should explain why their work has no societal impact or why the paper does not address societal impact.
- Examples of negative societal impacts include potential malicious or unintended uses (e.g., disinformation, generating fake profiles, surveillance), fairness considerations (e.g., deployment of technologies that could make decisions that unfairly impact specific groups), privacy considerations, and security considerations.
- The conference expects that many papers will be foundational research and not tied to particular applications, let alone deployments. However, if there is a direct path to any negative applications, the authors should point it out. For example, it is legitimate to point out that an improvement in the quality of generative models could be used to generate Deepfakes for disinformation. On the other hand, it is not needed to point out that a generic algorithm for optimizing neural networks could enable people to train models that generate Deepfakes faster.
- The authors should consider possible harms that could arise when the technology is being used as intended and functioning correctly, harms that could arise when the technology is being used as intended but gives incorrect results, and harms following from (intentional or unintentional) misuse of the technology.
- If there are negative societal impacts, the authors could also discuss possible mitigation strategies (e.g., gated release of models, providing defenses in addition to attacks, mechanisms for monitoring misuse, mechanisms to monitor how a system learns from feedback over time, improving the efficiency and accessibility of ML).

11. Safeguards

Question: Does the paper describe safeguards that have been put in place for responsible release of data or models that have a high risk for misuse (e.g., pre-trained language models, image generators, or scraped datasets)?

Answer: [N/A]

Justification: No data or models are released with this paper.

Guidelines:

- The answer [N/A] means that the paper poses no such risks.
- Released models that have a high risk for misuse or dual-use should be released with necessary safeguards to allow for controlled use of the model, for example by requiring that users adhere to usage guidelines or restrictions to access the model or implementing safety filters.
- Datasets that have been scraped from the Internet could pose safety risks. The authors should describe how they avoided releasing unsafe images.
- We recognize that providing effective safeguards is challenging, and many papers do not require this, but we encourage authors to take this into account and make a best faith effort.

12. Licenses for existing assets

Question: Are the creators or original owners of assets (e.g., code, data, models), used in the paper, properly credited and are the license and terms of use explicitly mentioned and properly respected?

Answer: [N/A]

Justification: No external code, datasets, or models are used; all cited prior work is credited via the references list.

Guidelines:

- The answer [N/A] means that the paper does not use existing assets.
- The authors should cite the original paper that produced the code package or dataset.
- The authors should state which version of the asset is used and, if possible, include a URL.
- The name of the license (e.g., CC-BY 4.0) should be included for each asset.
- For scraped data from a particular source (e.g., website), the copyright and terms of service of that source should be provided.
- If assets are released, the license, copyright information, and terms of use in the package should be provided. For popular datasets, paperswithcode.com/datasets has curated licenses for some datasets. Their licensing guide can help determine the license of a dataset.
- For existing datasets that are re-packaged, both the original license and the license of the derived asset (if it has changed) should be provided.
- If this information is not available online, the authors are encouraged to reach out to the asset's creators.

13. New assets

Question: Are new assets introduced in the paper well documented and is the documentation provided alongside the assets?

Answer: [N/A]

Justification: No new assets (datasets, code, models) are released with this paper.

Guidelines:

- The answer [N/A] means that the paper does not release new assets.
- Researchers should communicate the details of the dataset/code/model as part of their submissions via structured templates. This includes details about training, license, limitations, etc.
- The paper should discuss whether and how consent was obtained from people whose asset is used.
- At submission time, remember to anonymize your assets (if applicable). You can either create an anonymized URL or include an anonymized zip file.

14. Crowdsourcing and research with human subjects

Question: For crowdsourcing experiments and research with human subjects, does the paper include the full text of instructions given to participants and screenshots, if applicable, as well as details about compensation (if any)?

Answer: [N/A]

Justification: No crowdsourcing or research with human subjects.

Guidelines:

- The answer [N/A] means that the paper does not involve crowdsourcing nor research with human subjects.
- Including this information in the supplemental material is fine, but if the main contribution of the paper involves human subjects, then as much detail as possible should be included in the main paper.
- According to the NeurIPS Code of Ethics, workers involved in data collection, curation, or other labor should be paid at least the minimum wage in the country of the data collector.

15. Institutional review board (IRB) approvals or equivalent for research with human subjects

1370 Question: Does the paper describe potential risks incurred by study participants, whether
1371 such risks were disclosed to the subjects, and whether Institutional Review Board (IRB)
1372 approvals (or an equivalent approval/review based on the requirements of your country or
1373 institution) were obtained?

1374 Answer: [N/A]

1375 Justification: No human subjects.

1376 Guidelines:

- 1377 • The answer [N/A] means that the paper does not involve crowdsourcing nor research
1378 with human subjects.
- 1379 • Depending on the country in which research is conducted, IRB approval (or equivalent)
1380 may be required for any human subjects research. If you obtained IRB approval, you
1381 should clearly state this in the paper.
- 1382 • We recognize that the procedures for this may vary significantly between institutions
1383 and locations, and we expect authors to adhere to the NeurIPS Code of Ethics and the
1384 guidelines for their institution.
- 1385 • For initial submissions, do not include any information that would break anonymity (if
1386 applicable), such as the institution conducting the review.

1387 16. Declaration of LLM usage

1388 Question: Does the paper describe the usage of LLMs if it is an important, original, or
1389 non-standard component of the core methods in this research? Note that if the LLM is used
1390 only for writing, editing, or formatting purposes and does *not* impact the core methodology,
1391 scientific rigor, or originality of the research, declaration is not required.

1392 Answer: [N/A]

1393 Justification: LLMs were used only for writing assistance, formatting, and citation-
1394 verification cross-checks against published sources. The core mathematical content (theorem
1395 statements, proofs, the four-component composition, the point-mass identity, the sector-
1396 Lyapunov reduction) is author-developed and author-verified; the LLM did not originate
1397 methodological or theoretical contributions.

1398 Guidelines:

- 1399 • The answer [N/A] means that the core method development in this research does not
1400 involve LLMs as any important, original, or non-standard components.
- 1401 • Please refer to our LLM policy in the NeurIPS handbook for what should or should not
1402 be described.